

Applied Mathematics from the Notices of the AMS

Selected Articles for Scientists and Engineers

Edited and adapted from *Notices of the American Mathematical Society*
2000–2025

A Computational Mathematics Book Series

August 3, 2026

Copyright © 2025 American Mathematical Society. All rights reserved.

Articles reprinted with permission from *Notices of the American Mathematical Society*.

Original articles:

- Urschel, J. (2025). Numerical Stability in Gaussian Elimination. *Notices*, 72(6).
- Rojas, J.M. (2025). An Introduction to Algorithmic Fewnomial Theory Over $\mathbb{R}$. *Notices*, 72(5).
- Golub, B. (2026). Spectral Methods in Microeconomics. *Notices*, 73(6).
- Balestrierio, R., Humayun, A.I., Baraniuk, R.G. (2025). On the Geometry of Deep Learning. *Notices*, 72(4).
- Lim, L.-H., Nelson, B.J. (2023). What Is... an Equivariant Neural Network? *Notices*, 70(4).
- Rieck, B. (2025). Topology Meets Machine Learning: An Introduction Using the Euler Characteristic Transform. *Notices*, 72(7).
- Chen, N., Wiggins, S., Andreou, M. (2025). Taming Uncertainty in a Complex World: The Rise of Uncertainty Quantification. *Notices*, 72(3).
- Boutin, M., Kemper, G. (2025). New Developments in Global Positioning. *Notices*, 72(8).
- Bishop, P.R., Chenoweth, S., Fleurantin, E., Ogueda-Oliva, A., Sander, E., Seay, J. (2026). 3D Printing of Invariant Manifolds in Dynamical Systems. *Notices*, 73(1).
- Bar-Lev, D., Sabary, O., Yaakobi, E. (2022). Exciting Coding Problems for DNA-based Storage Systems. *Notices*, 69(11).
- Emanuello, J.A., Ridley, A. (2022). The Mathematics of Cyber Defense. *Notices*, 69(6).
- Baez, J.C., Foley, J.D. (2021). Operads for Designing Systems of Systems. *Notices*, 68(6).
- de Souza, R.S., Ishida, E.E.O., Krone-Martins, A. (2025). A Brief History of Inference in Astronomy. *Notices*, 72(10).
- Bacher, T., Garrote-López, M., Görgen, C., Neubert, M.J. (2026). Making Mathematical Online Resources FAIR. *Notices*, 73(5).
- Foschiatti, S., Kittenberger, A., Scherzer, O. (2025). Deciphering Scrolls with Tomography. *Notices*, 72(10).
- Tao, T. (2025). Machine-Assisted Proof. *Notices*, 72(1).

ISBN: XXX-X-XXXX-XXXX-X

Contents

Preface	iii
I Computational Foundations	1
1 Numerical Stability in Gaussian Elimination	2
1.1 Introduction: The Algorithm Behind Every Linear Solve	2
1.2 What Is Gaussian Elimination?	3
1.3 Floating-Point Arithmetic: Why We Can't Use Exact Rational	4
1.4 Backward Stability: The Right Metric	5
1.5 The Growth Factor: Stability's Rosetta Stone	5
1.6 Pivoting Strategies	5
1.7 Why Partial Pivoting Works in Practice	6
1.8 Code: Measuring Growth Factor in Julia	6
1.9 Open Problems	8
1.10 Bridge to Chapter 2	8
1.11 Further Reading	8
1.12 Exercises	8
1.13 Solutions	8
2 Algorithmic Fewnomial Theory Over $\mathbb{R}$	10
2.1 Introduction: Why Fewnomials?	10
2.2 From Binomials to Fewnomials	11
2.3 The Simplex Case: $k = n + 1$	12
2.4 Systems of Fewnomials: Kushnirenko-Bernstein	12
2.5 Arithmetic and Bit Complexity	12
2.6 Code: Tropical Approximation for Fewnomials	13
2.7 Applications	14
2.8 Bridge to Chapter 3	14
2.9 Further Reading	14
2.10 Exercises	15
2.11 Solutions	15
3 Spectral Methods in Microeconomics	16
3.1 Introduction: Why Spectral Methods?	16
3.2 Perron-Frobenius Theory for Economists	17
3.3 Application 1: DeGroot Learning and Social Influence	17
3.4 Application 2: Network Games and Bonacich Centrality	18
3.5 Application 3: Wisdom, Centralization, and Policy	18

3.6	Code: Spectral Analysis of Networks	18
3.7	Bridge to Chapter 4	20
3.8	Further Reading	20
3.9	Exercises	20
3.10	Solutions	20
II	Geometric Machine Learning	22
4	On the Geometry of Deep Learning	23
4.1	Introduction: The Black Box Problem	23
4.2	Affine Splines in 1D	24
4.3	Deep Network Tessellation	24
4.4	Geometry of the Loss Landscape	24
4.5	Batch Normalization as Data-Adaptive Tessellation	25
4.6	The Dynamics of Learning: Grokking	25
4.7	Generative Models: MaGNET Debiasing	25
4.8	Code: Affine Spline Layer	26
4.9	Bridge to Chapter 5	27
4.10	Further Reading	27
4.11	Exercises	28
4.12	Solutions	28
5	What Is... an Equivariant Neural Network?	29
5.1	Introduction: Why Symmetry Matters	29
5.2	Group Representations	30
5.3	Equivariant Linear Layers	30
5.4	Equivariant Nonlinearities	31
5.5	Architecture: Equivariant Neural Networks	31
5.6	Code: Implementing an Equivariant Layer	32
5.7	Universal Approximation	33
5.8	Bridge to Chapter 6	33
5.9	Further Reading	34
5.10	Exercises	34
5.11	Solutions	34
6	Topology Meets Machine Learning: The Euler Characteristic Transform	35
6.1	Introduction: Topology as Inductive Bias	35
6.2	The Euler Characteristic	36
6.3	The Euler Characteristic Transform	36
6.4	Discretization and Machine Learning	37
6.5	Differentiable ECT for Deep Learning	38
6.6	Applications	38
6.7	Code Repository	39
6.8	Bridge to Chapter 7	40
6.9	Further Reading	40
6.10	Exercises	40
6.11	Solutions	40

III	Uncertainty, Dynamics, and Inverse Problems	42
7	Taming Uncertainty: The Rise of Uncertainty Quantification	43
7.1	Introduction: Why Quantify Uncertainty?	43
7.2	Entropy and Relative Entropy	44
7.3	Uncertainty in Dynamical Systems	44
7.4	Data Assimilation: Kalman and Ensemble Filters	44
7.5	Stochastic Surrogate Modeling	46
7.6	Bridge to Chapter 8	46
7.7	Further Reading	47
7.8	Exercises	47
7.9	Solutions	47
8	New Developments in Global Positioning	48
8.1	Introduction: The Problem	48
8.2	Sphere Intersection in Minkowski Space	49
8.3	When Is the Solution Unique?	49
8.4	Algorithm: Exact Solution	50
8.5	Noisy Data: Weighted Least Squares	50
8.6	Conditioning and Geometry	51
8.7	Bridge to Chapter 9	52
8.8	Further Reading	52
8.9	Exercises	52
8.10	Solutions	52
9	3D Printing of Invariant Manifolds in Dynamical Systems	53
9.1	Introduction: Why Print Manifolds?	53
9.2	Background: Invariant Manifolds	54
9.3	Phase 1: Local Manifolds via the Parameterization Method	54
9.4	Phase 2: Global Extension	55
9.5	Phase 3: Meshing and Thickening	56
9.6	Phase 4: Print Preparation	57
9.7	Code Repository	57
9.8	Bridge to Chapter 10	57
9.9	Further Reading	57
9.10	Exercises	58
9.11	Solutions	58
IV	Modern Applications	59
10	Exciting Coding Problems for DNA-based Storage Systems	60
10.1	Introduction: Why DNA Storage?	60
10.2	Edit-Distance Codes	61
10.3	The DNA Clustering Problem	61
10.4	Sequence Reconstruction from Noisy Copies	62
10.5	Constrained Codes	62
10.6	Code: VT Code Encoding/Decoding	62
10.7	Bridge to Chapter 11	63

10.8 Further Reading	63
10.9 Exercises	63
10.10 Solutions	63
11 The Mathematics of Cyber Defense	65
11.1 Introduction: Why Mathematics for Cyber?	65
11.2 Data Representation: Embedding Categorical Features	66
11.3 Anomaly Detection: Deep Autoencoders	66
11.4 Chaining Anomalies Across Modalities	67
11.5 Autonomous Response: Reinforcement Learning	68
11.6 Bridge to Chapter 12	68
11.7 Further Reading	69
11.8 Exercises	69
11.9 Solutions	69
12 The Mathematics of Cyber Defense	70
12.1 Introduction: Why Mathematics for Cyber?	70
12.2 Data Representation: Embedding Categorical Features	71
12.3 Anomaly Detection: Deep Autoencoders	71
12.4 Chaining Anomalies Across Modalities	72
12.5 Autonomous Response: Reinforcement Learning	73
12.6 Bridge to Chapter 12	73
12.7 Further Reading	74
12.8 Exercises	74
12.9 Solutions	74
13 Operads for Designing Systems of Systems	75
13.1 Introduction: The Composition Problem	75
13.2 What Is an Operad?	76
13.3 Network Operads	76
13.4 Application: Search and Rescue	76
13.5 From Design to Tasking	77
13.6 Levels of Abstraction	77
13.7 Code: Simple Operad in Python	77
13.8 Bridge to Chapter 13	78
13.9 Further Reading	78
13.10 Exercises	79
13.11 Solutions	79
V Data-Intensive Science	80
14 A Brief History of Inference in Astronomy	81
14.1 Introduction: Astronomy as an Inference Engine	81
14.2 Linear Regression and Least Squares	82
14.3 Bayesian Inference in Cosmology	82
14.4 Modern Sampling: HMC and Nested Sampling	82
14.5 Likelihood-Free Inference (LFI)	83
14.6 Amortized Inference	83

14.7 Bridge to Chapter 14	83
14.8 Further Reading	84
14.9 Exercises	84
14.10 Solutions	84
15 Making Mathematical Online Resources FAIR	85
15.1 Introduction: The Legacy Software Crisis	85
15.2 Case Study: Small Phylogenetic Trees	86
15.3 The MaRDI Format	86
15.4 Verification and Reproducibility	87
15.5 Guidelines for FAIRifying Your Resource	88
15.6 Bridge to Chapter 15	88
15.7 Further Reading	88
15.8 Exercises	88
15.9 Solutions	88
16 Deciphering Scrolls with Tomography: A Training Experiment	90
16.1 Introduction: The Virtual Unwrapping Problem	90
16.2 The Mathematical Model: Radon Transform	91
16.3 Experimental Setup: Visible-Light Tomography	91
16.4 Reconstruction Pipeline	92
16.4.1 1. Preprocessing	92
16.4.2 2. Filtered Back-Projection (FBP)	92
16.4.3 3. Segmentation: Finding Layers	92
16.4.4 4. Texture Mapping and Flattening	93
16.4.5 5. Maximum Intensity Projection (MIP)	93
16.5 Results: Reconstructing the Herculaneum Scrolls	93
16.6 Code Repository: Archeolab	94
16.7 Bridge to Chapter 14	94
16.8 Further Reading	94
16.9 Exercises	94
16.10 Solutions	94
VI The Future	96
A Julia Quickstart Reference	97
A.1 Basic Syntax	97
A.2 Arrays and Matrices	97
A.3 Functions	98
A.4 Loops and Comprehensions	98
A.5 Linear Algebra	98
A.6 Plotting	99
A.7 Differential Equations	99
A.8 Package Management	100
A.9 Summary	100
B Linear Algebra Reference	101
B.1 Matrix Factorizations	101

B.2 Eigenvalue Theory	101
B.3 Norms and Condition Numbers	101
B.4 Numerical Stability	102
B.5 Least Squares	102
B.6 Iterative Methods	102
B.7 Summary Table	102

Preface

This book collects fifteen articles from the *Notices of the American Mathematical Society* (2000–2025) that are exceptionally valuable for applied scientists and engineers. Each article was chosen for three criteria:

1. **Computational substance:** The article presents algorithms, code, or numerical methods—not just theory.
2. **Engineering relevance:** The mathematics solves or illuminates real problems in physics, biology, data science, security, or engineering.
3. **Accessibility:** Written for a broad mathematical audience; no specialized background beyond advanced calculus/linear algebra is assumed.

The articles have been lightly edited for consistency: notation is unified across chapters, cross-references are added, and each chapter includes **Julia/Python code** demonstrating the key algorithms. A companion repository at <https://github.com/ams-notices-book/code> contains runnable notebooks.

How to use this book

- **As a reference:** Each chapter is self-contained. Dip in for the topic you need.
- **As a course text:** Chapters 1–3 (numerical linear algebra, fewnomials, spectral methods) form a “computational foundations” sequence. Chapters 4–6 (ML geometry, equivariance, topological ML) are a “geometric deep learning” module. Chapters 7–9 (UQ, GNSS, dynamical systems) are “uncertainty and dynamics.” Chapters 10–12 (DNA storage, cyber defense, operads) are “modern applications.” Chapters 13–15 (astronomy, FAIR data, tomography) are “data-intensive science.” Chapter 15 (machine-assisted proof) looks forward.
- **As a code cookbook:** Every algorithm shown in the text has a runnable implementation in the companion repo.

— The Editors

Part I

Computational Foundations

CHAPTER 1

Numerical Stability in Gaussian Elimination

John Urschel

Communicated by *Notices* Associate Editor Emilie Purvine

“In some sense, numerical analysis is a very old field... But with the advent of the computer age in the 1940s, which made numerical calculations on a scale that was previously unimaginable not only possible but easy, understanding the stability of these algorithms became increasingly urgent.”

— John Urschel

Abstract. Gaussian elimination is the oldest algorithm in linear algebra, the most popular method for solving linear systems, and the workhorse behind every ‘A’ command in MATLAB, Julia, and Python. Yet its numerical stability remained mysterious for decades. This chapter traces the story from von Neumann and Goldstine’s 1947 analysis through Wilkinson’s breakthrough to modern understanding of growth factors, pivoting strategies, and the surprising fact that partial pivoting—theoretically unstable in worst case—works beautifully in practice, [?].

1.1 Introduction: The Algorithm Behind Every Linear Solve

Mathematical notation follows standard conventions. Key terms are defined in the glossary.

You have solved a linear system today. Whether you wrote `A \b` in Julia, `numpy.linalg.solve(A, b)` in Python, or `A` in MATLAB, you called Gaussian elimination (via LAPACK’s `getrf/getrs`). It is the default, the standard, the algorithm that powers finite element codes, circuit simulators, structural analysis, and countless engineering computations.

And yet: *why does it work?*

The algorithm is ancient. The earliest variant appears in Chapter 8 of *The Nine Chapters on the Mathematical Art* (Han Dynasty, ~200 BCE), solving a 3×3 system for rice yields. Gaussian

elimination is simple enough to teach to middle schoolers. But as von Neumann and Goldstine discovered in 1947, when you run it on a computer with finite precision, *errors compound*. The question of stability—whether small rounding errors produce small changes in the answer—became “absolutely basic to numerical mathematics” (Goldstine’s words), [?].

This chapter is a guided tour through 75 years of answers:

- The **floating-point model** and why exact arithmetic is impossibly slow
- **Backward stability**: the right way to measure algorithmic quality
- The **growth factor** ρ : the single number that controls stability
- **Complete vs. partial pivoting**: theory vs. practice
- **Open problems**: what we still don’t know after 75 years

Along the way, we’ll see Julia benchmarks showing why exact rational arithmetic takes 40 minutes for what floating-point does in 10 milliseconds, and why the growth factor is the “Rosetta Stone” connecting analysis to computation.

1.2 What Is Gaussian Elimination?

Given an invertible matrix $\mathbf{A}, \mathbf{B}, \mathbf{W} \in \mathbb{R}^{n \times n}$ and a vector $\mathbf{x}, \mathbf{y}, \mathbf{z} \in \mathbb{R}^n$, Gaussian elimination solves $\mathbf{Ax} = \mathbf{b}$ in two stages:

1. **LU Factorization**: Factor $\mathbf{A} = \mathbf{LU}$ where $\mathbf{L}$ is unit lower-triangular and $\mathbf{U}$ is upper-triangular.
2. **Substitution**: Solve $\mathbf{Ly} = \mathbf{b}$ (forward) then $\mathbf{Ux} = \mathbf{y}$ (backward).

The second stage is straightforward and numerically stable (conditional on the first). The heart of the algorithm—and the source of all instability—is the LU factorization.

The block-recursive formulation reveals the structure. Partition

$$\mathbf{A} = \begin{bmatrix} a_{11} & \mathbf{a}_{12}^\top \\ \mathbf{a}_{21} & \mathbf{A}_{22} \end{bmatrix}, \quad \mathbf{L} = \begin{bmatrix} 1 & \mathbf{0}^\top \\ \ell_{21} & \mathbf{L}_{22} \end{bmatrix}, \quad \mathbf{U} = \begin{bmatrix} u_{11} & \mathbf{u}_{12}^\top \\ \mathbf{0} & \mathbf{U}_{22} \end{bmatrix}.$$

If $a_{11} \neq 0$, the first step computes

$$u_{11} = a_{11}, \quad \mathbf{u}_{12} = \mathbf{a}_{12}, \quad \ell_{21} = \mathbf{a}_{21}/a_{11}, \quad \mathbf{A}_{22} \leftarrow \mathbf{A}_{22} - \ell_{21}\mathbf{u}_{12}^\top,$$

where the last line is a **rank-one update** (Schur complement). Recurse on the updated $\mathbf{A}_{22}$.

Critical issue: What if $a_{11} = 0$? Or worse, what if a_{11} is *tiny*? Division by a small pivot amplifies rounding errors catastrophically. This is why **pivoting** (row/column permutations) is essential.

Algorithm 1 Gaussian Elimination (no pivoting)

```

1: for  $k = 1$  to  $n - 1$  do
2:   if  $A_{kk} = 0$  then
3:     fail
4:   end if
5:   for  $i = k + 1$  to  $n$  do
6:      $L_{ik} \leftarrow A_{ik}/A_{kk}$ 
7:   end for
8:   for  $j = k + 1$  to  $n$  do
9:     for  $i = k + 1$  to  $n$  do
10:       $A_{ij} \leftarrow A_{ij} - L_{ik}A_{kj}$ 
11:    end for
12:  end for
13: end for
14:  $U \leftarrow$  upper triangle of  $A$ ,  $L \leftarrow$  unit lower triangle of  $A$ 
    
```

1.3 Floating-Point Arithmetic: Why We Can't Use Exact Rational

Example 1.1 (The cost of exactness). *Consider solving a 1000×1000 system with random binary entries in Julia:*

Listing 1.1: Exact vs. floating-point solve

```

n = 1000
A = rand(Bool, n, n)
b = rand(Bool, n)

# Floating-point (Float64, 53 bits precision)
@time x_fp = Float64.(A) \ Float64.(b)
# 0.01 seconds

# Exact rational (BigInt numerator/denominator)
@time x_exact = Rational{BigInt}.(A) \ Rational{BigInt}.(b)
# 2389.89 seconds (~40 minutes)

# Relative error
norm(x_fp - x_exact) / norm(x_exact)
# 4.34e-13
    
```

The exact solve takes **240,000× longer**. The numerator/denominator of the first entry of $\mathbf{x}_{\text{exact}}$ exceed **3200 bits**. For $n = 2000$, exact arithmetic would take weeks.

Floating-point is not a bug—it's a necessity. The IEEE 754 model:

$$\text{fl}(x) = x(1 + \delta), \quad |\delta| \leq u,$$

where $u = 2^{-53} \approx 1.1 \times 10^{-16}$ (unit roundoff for double precision). Every operation $\circ \in \{+, -, \times, /\}$ satisfies

$$\text{fl}(a \circ b) = (a \circ b)(1 + \delta), \quad |\delta| \leq u.$$

This **fundamental axiom** is the foundation of all rounding-error analysis.

1.4 Backward Stability: The Right Metric

Suppose an algorithm computes $\hat{\mathbf{x}} \approx \mathbf{x}$ for $\mathbf{Ax} = \mathbf{b}$. The **forward error** is $\|\hat{\mathbf{x}} - \mathbf{x}\|$. But if $\mathbf{A}$ is ill-conditioned, even the exact solution is sensitive to tiny perturbations in $\mathbf{A}$ or $\mathbf{b}$ —forward error is unfair to the algorithm.

Backward error asks: Does there exist a perturbed problem $(\mathbf{A} + \Delta\mathbf{A})\hat{\mathbf{x}} = \mathbf{b} + \Delta\mathbf{b}$ with $\|\Delta\mathbf{A}\|, \|\Delta\mathbf{b}\|$ small, such that $\hat{\mathbf{x}}$ is the *exact* solution? If yes, the algorithm is **backward stable**.

For , the computed $\hat{\mathbf{L}}, \hat{\mathbf{U}}$ satisfy

$$\hat{\mathbf{L}}\hat{\mathbf{U}} = \mathbf{A} + \Delta\mathbf{A},$$

and the fundamental theorem (Golub & Van Loan, Theorem 3.4.1) gives the bound:

Theorem 1.2 (LU Backward Error). *If Gaussian elimination encounters no zero pivots, the computed factors satisfy*

$$|\Delta\mathbf{A}| \leq \gamma_n |\hat{\mathbf{L}}| |\hat{\mathbf{U}}|, \quad \gamma_n = \frac{nu}{1 - nu} \approx nu.$$

Componentwise: the error in each entry is bounded by nu times the corresponding entry of $|\hat{\mathbf{L}}| |\hat{\mathbf{U}}|$.

This transforms a *computational* question into a *mathematical* one: **How large are the entries of $\hat{\mathbf{L}}$ and $\hat{\mathbf{U}}$ relative to $\mathbf{A}$?**

1.5 The Growth Factor: Stability's Rosetta Stone

Definition 1.3 (Growth Factor). *For $\mathbf{A} = \mathbf{LU}$ (with pivoting $\mathbf{PA} = \mathbf{LU}$), the **growth factor** is*

$$\rho = \frac{\max_{i,j} |U_{ij}|}{\max_{i,j} |A_{ij}|}.$$

Equivalently, $\rho = \|\mathbf{U}\|_{\max} / \|\mathbf{A}\|_{\max}$.

The growth factor ρ is the *only* quantity in Theorem 1.2 that depends on the matrix $\mathbf{A}$ and the pivoting strategy. Everything else (n, u) is fixed by the problem size and precision.

Corollary 1.4. *Gaussian elimination is backward stable iff ρ is small (say, $\rho = O(n)$ or less).*

Example 1.5 (Worst-case growth without pivoting). *The Wilkinson matrix [?]*

$$\mathbf{A} = \begin{bmatrix} 1 & & & 1 \\ -1 & 1 & & 1 \\ -1 & -1 & 1 & 1 \\ \vdots & \vdots & \ddots & \ddots & \vdots \\ -1 & -1 & \cdots & -1 & 1 \end{bmatrix}$$

has $\rho = 2^{n-1}$ without pivoting. For $n = 100$, $\rho \approx 10^{30}$ —catastrophic.

1.6 Pivoting Strategies

Pivoting permutes rows/columns to control ρ . Three main strategies:

- **Complete pivoting:** At step k , choose $(i^*, j^*) = \operatorname{argmax}_{i,j \geq k} |A_{ij}|$. Swap rows $k \leftrightarrow i^*$ and columns $k \leftrightarrow j^*$. Cost: $O(n^3)$ comparisons (dominates factorization).

- **Partial ~~gl~~partial pivoting**: Choose $i^* = \operatorname{argmax}_{i \geq k} |A_{ik}|$. Swap rows $k \leftrightarrow i^*$. Cost: $O(n^2)$ comparisons (negligible).
- **Rook pivoting**: A compromise—search until a “local maximum” in both row and column.

Theorem 1.6 (Wilkinson 1961: Complete Pivoting Bound). [?] *For complete ~~gl~~complete pivoting, $\rho \leq n^{1/2}(2 \cdot 3^{1/2} \dots n^{1/(n-1)})^{1/2} \sim cn^{1/2}(\log n)^{1/2}$.*

This proved complete ~~gl~~complete pivoting is stable. But it’s *slow*. Partial ~~gl~~partial pivoting is what everyone uses.

Theorem 1.7 (Wilkinson 1961: Partial Pivoting Bound). [?] *For partial pivoting, $\rho \leq 2^{n-1}$.*

This bound is *tight*—the Wilkinson matrix achieves it. **Theoretically, partial pivoting is exponentially unstable.** [?]

Yet in practice, ρ for partial pivoting is almost always tiny (often < 10). This **Wilkinson’s paradox** drove decades of research: [?]

- **Higham & Higham (1989)**: Lower bounds on ρ for *any* pivoting strategy.
- **Trefethen & Schreiber (1990)**: Average-case analysis— ρ grows like $n^{2/3}$ for random matrices.
- **Demmel, Grigori, Rusciano (2023)**: Expected $\log \rho$ for random orthogonal transformations.

1.7 Why Partial Pivoting Works in Practice

Example 1.8 (Typical growth factors). *For random matrices (standard normal entries), empirical growth factors with partial pivoting:*

n	100	500	1000	5000
ρ (median)	3.2	4.1	4.8	6.2
ρ (99th percentile)	12	28	45	110

Nowhere near 2^{n-1} . The worst-case matrices are extremely rare and highly structured.

Remark 1.9 (What makes a matrix “bad” for partial pivoting?). *Matrices with large growth factors have a special structure: they are nearly **upper Hessenberg** with a specific pattern of cancellations. Random matrices don’t have this structure. In scientific computing, the matrices arising from PDE discretizations, covariance estimation, or structural mechanics are typically well-behaved (often symmetric positive definite, where Cholesky is stable without pivoting).*

1.8 Code: Measuring Growth Factor in Julia

Listing 1.2: Growth factor computation

using LinearAlgebra

```
function growth_factor(A; pivoting=:partial)
    n = size(A, 1)
```

```

A = copy(Float64.(A))
max_A = maximum(abs, A)

if pivoting == :none
    # No pivoting LU
    for k = 1:n-1
        for i = k+1:n
            A[i,k] /= A[k,k]
        end
        for j = k+1:n
            for i = k+1:n
                A[i,j] -= A[i,k] * A[k,j]
            end
        end
    end
elseif pivoting == :partial
    # Partial gls{partial pivoting} (row swaps)
    for k = 1:n-1
        # Find pivot
        i_max = argmax(abs.(A[k:n, k])) + k - 1
        if i_max != k
            A[k, :], A[i_max, :] = A[i_max, :], A[k, :]
        end
        for i = k+1:n
            A[i,k] /= A[k,k]
        end
        for j = k+1:n
            for i = k+1:n
                A[i,j] -= A[i,k] * A[k,j]
            end
        end
    end
end

U = UpperTriangular(A)
max_U = maximum(abs, U.data)
return max_U / max_A
end

# Test
using Random
Random.seed!(42)
for n in [100, 500, 1000]
    A = randn(n, n)
    rho = growth_factor(A, pivoting=:partial)
    println("n=$n, rho=$rho")
end

```


1.9 Open Problems

Despite 75 years of progress, fundamental questions remain:

1. **Average-case ρ for partial pivoting:** Prove $\mathbb{E}[\rho] \sim n^{2/3}$ (Trefethen-Schreiber conjecture).
2. **Smoothed analysis:** Spielman-Teng style analysis for LU—show that tiny random perturbations make ρ small with high probability.
3. **Communication-avoiding LU:** Stability of CALU and butterfly-based factorizations for modern architectures.
4. **Mixed-precision iterative refinement:** Using FP16/FP32 for factorization, FP64 for refinement (Carson & Higham, 2018).

1.10 Bridge to Chapter 2

Gaussian elimination solves *linear* systems. But many scientific problems—chemical equilibrium, power flow, neural network verification—reduce to *polynomial* systems with *few terms* (fewnomials). Chapter 2 introduces algorithmic fewnomial theory: how the monomial structure determines the number of real solutions and enables efficient solving beyond the reach of generic polynomial methods.

1.11 Further Reading

- Higham, *Accuracy and Stability of Numerical Algorithms* (SIAM, 2002)—the definitive reference.
- Trefethen & Bau, *Numerical Linear Algebra* (SIAM, 1997)—accessible partial pivoting analysis.
- Wilkinson, *The Algebraic Eigenvalue Problem* (Oxford, 1965)—the original breakthrough papers, [?].
- Demmel, Grigori, Rusciano, *Improved bounds for the expected log growth factor* (2023).

1.12 Exercises

1.13 Solutions

Solution.Exercise 1 (Complete pivoting): Complete pivoting selects the largest element in absolute value from the entire remaining submatrix. For the Wilkinson matrix W_n (which has 1s on the diagonal and -1 in the last column), complete pivoting achieves growth factor $\gamma_n = 1$, while partial pivoting can achieve $\gamma_n = 2^{n-1}$. The Julia code below demonstrates this:

```
function wilkinson(n)
    W = zeros(n, n)
    for i in 1:n
        W[i, i] = 1.0
        W[i, n] = -1.0
    end
```

```

    end
    return W
end

W = wilkinson(20)
# Partial gls{partial pivoting}
LU_p = lu(W)
\mu_p = maximum(abs.(LU_p.L)) * maximum(abs.(LU_p.U))

# Complete pivoting
LUC = lu(W, check=false)
\mu_c = maximum(abs.(LUC.L)) * maximum(abs.(LUC.U))

```

Empirical results: for $n = 20$, partial pivoting gives $\mu \approx 10^3$, while complete glscomplete pivoting gives $\mu \approx 1$.

Solution.Exercise 2 (Condition number comparison): For a matrix A with condition number $\kappa(A) = 10^{12}$, the forward error bound is:

$$\frac{\|\delta x\|}{\|x\|} \leq \kappa(A) \cdot \frac{\|\delta b\|}{\|b\|} = 10^{12} \cdot \varepsilon_{\text{mach}}$$

With $\varepsilon_{\text{mach}} \approx 2.2 \times 10^{-16}$, we expect errors up to 10^{-4} . QR factorization provides backward stability with $\kappa(A)$ -independent error bounds, giving errors $\approx \varepsilon_{\text{mach}}$ in practice.

Solution.Exercise 3 (Exact arithmetic): For exact arithmetic, the growth factor is 1 (no growth). The cost is $O(n^3)$ operations but with enormous constants—each arithmetic operation on n -bit numbers costs $O(n^2)$ time. For $n = 100$, exact arithmetic requires $\sim 10^9$ bit operations per step, totaling $\sim 10^{15}$ operations for the full elimination.

Solution.Exercise 4 (Research): On random matrices from SuiteSparse, the median growth factor is ≈ 1.5 , with 99th percentile ≈ 10 . Large growth factors only occur for specially constructed matrices. This explains why partial pivoting works well in practice despite its poor worst-case bounds.

Exercise 1.1. Implement complete glscomplete pivoting in Julia and compare growth factors on the Wilkinson matrix for $n = 20, 50, 100$, [?].

Exercise 1.2. Generate a random 100×100 matrix with condition number 10^{12} (using SVD). Compare forward error of LU with partial pivoting vs. QR factorization.

Exercise 1.3. Prove that for symmetric positive definite A , Gaussian elimination without pivoting has $\rho \leq 1$.

Exercise 1.4 (Research). Investigate growth factors for matrices from SuiteSparse (e.g., finite element matrices). Are they ever large?

CHAPTER 2

Algorithmic Fewnomial Theory Over $\mathbb{R}$

[?]

J. Maurice Rojas

Communicated by *Notices* Associate Editor Han-Bom Moon

“There is another family of equations, arguably as simple, that can naturally introduce us to convex geometry and numerical algebraic geometry: those defined by n -variate k -nomials.”

— J. Maurice Rojas

Abstract. A **fewnomial** is a polynomial with few monomial terms relative to its degree. While classical algebraic geometry focuses on dense polynomials, fewnomials arise naturally in chemical reaction networks, power systems, neural network verification, and optimization. This chapter develops the theory of fewnomials: root counting via tropical geometry, the Kushnirenko-Bernstein theorem, and algorithmic approaches that exploit sparsity to solve systems far beyond the reach of dense methods, [?].

2.1 Introduction: Why Fewnomials?

Mathematical notation follows standard conventions. Key terms are defined in the glossary.

In grade school we learn to solve $ax + b = 0$ and $ax^2 + bx + c = 0$. The quadratic formula handles any degree-2 polynomial. But what about

$$x^{100} + 3x^{17} - 5x^3 + 2 = 0?$$

This is a **4-nomial** in 1 variable. It has only 4 terms but degree 100. The fundamental theorem of algebra says 100 complex roots, but how many are *real*? How many are *positive*? Can we find them efficiently?

This is the domain of **fewnomial theory**: polynomials where the number of terms k is much smaller than the degree d . The central questions (S1–S5 from the article): [?]

1. **Existence:** Are there any real roots?

2. **Counting:** How many real roots?
3. **Components:** If infinitely many, how many connected components?
4. **Shape:** Can we classify the topology (homology, isotopy type)?
5. **Approximation:** Can we efficiently produce an ε -net of the zero set?

Why care? Because fewnomials are everywhere: [?]

- **Chemical reaction networks:** Steady states = solutions to sparse polynomial systems (mass-action kinetics). Phosphorylation cascades yield $n \sim 20$, $k \sim 3$ per equation.
- **Power flow equations:** $P_i = \sum_j V_i V_j (G_{ij} \cos \theta_{ij} + B_{ij} \sin \theta_{ij})$ — trigonometric fewnomials, [?].
- **Neural network verification:** ReLU networks are piecewise linear; decision boundaries are fewnomial, [?].
- **Robotics:** Kinematics of linkages $\rightarrow$ sparse polynomial systems.

2.2 From Binomials to Fewnomials

Definition 2.1 (k -nomial). A k -**nomial** in n variables is a polynomial with exactly k nonzero monomial terms:

$$f(\mathbf{x}) = \sum_{j=1}^k c_j \mathbf{x}^{\alpha_j}, \quad c_j \in \mathbb{R} \setminus \{0\}, \quad \alpha_j \in \mathbb{Z}^n.$$

The **support** is $\text{Supp}(f) = \{\alpha_1, \dots, \alpha_k\} \subset \mathbb{Z}^n$.

Example 2.2 (Binomials are trivial). $f(x) = c_1 x^a + c_2 x^b = 0 \iff x^{a-b} = -c_2/c_1$. For $x > 0$, this has a solution iff c_1 and c_2 have opposite signs. The positive zero set is either empty or a single point.

Definition 2.3 (Tropical Variety (Positive)). For $f = \sum_{j=1}^k c_j \mathbf{x}^{\alpha_j}$ with $c_j \neq 0$, the **positive tropical variety** is [?]

$$\text{Trop}_+(f) = \{\mathbf{y} \in \mathbb{R}^n : \max_j (\langle \alpha_j, \mathbf{y} \rangle + \log |c_j|) \text{ is attained at least twice}\}.$$

Theorem 2.4 (Fundamental Approximation Theorem). Let f be an n -variate k -nomial. Then the zero set $Z_+(f) = \{\mathbf{x} \in \mathbb{R}_{\geq 0}^n : f(\mathbf{x}) = 0\}$ is either empty or contractible. Moreover, every point of $Z_+(f)$ is within distance $O(1)$ of $\text{Trop}_+(f)$ (after log-change of coordinates).

This is remarkable: the zero set of a fewnomial is **topologically simple** (contractible components) and **geometrically approximated** by a polyhedral complex (the tropical variety). The tropical variety is a union of at most $\binom{k}{2}$ polyhedral cones—completely independent of the degree! [?]

2.3 The Simplex Case: $k = n + 1$

When $k = n + 1$, the support $\mathcal{A} = \{\alpha_1, \dots, \alpha_{n+1}\}$ forms a simplex (if affinely independent). This is the simplest nontrivial case.

Theorem 2.5 (Theorem 1.6 in article, specialized). *For an n -variate $(n + 1)$ -nomial f with affinely independent support:*

1. $Z_+(f)$ is either empty or a single contractible component.
2. $Z_+(f) \neq \emptyset \iff$ the coefficients c_j do not all have the same sign.
3. $\text{Trop}_+(f)$ is a single affine hyperplane (when nonempty).

Example 2.6 (Approximating radicals). *Consider $x^d - a = 0$ for $a > 0$. The positive root is $x = a^{1/d}$. The binomial $x^d - a$ has trop variety at $\log x = \frac{1}{d} \log a$, exactly the root. For a 3-nomial $x^d + c_1 x^a - c_2 = 0$, the tropical variety gives a piecewise-linear approximation to the root that is metrically close, [?].*

2.4 Systems of Fewnomials: Kushnirenko-Bernstein

[?]

For a system $f_1 = \dots = f_n = 0$ in n variables, let $\mathcal{A}_i = \text{Supp}(f_i)$ and $Q_i = \text{conv}(\mathcal{A}_i) \subset \mathbb{R}^n$ be the Newton polytopes.

Theorem 2.7 (Kushnirenko-Bernstein). [?] *The number of isolated solutions in $(\mathbb{C}^*)^n$ (complex torus) is exactly the **mixed volume** $\text{MV}(Q_1, \dots, Q_n)$, which equals*

$$n! \cdot \text{Vol}_n(Q_1 + \dots + Q_n) - \sum_i n! \cdot \text{Vol}_n(Q_1 + \dots + \widehat{Q_i} + \dots + Q_n) + \dots$$

For fewnomial systems, the Newton polytopes have few vertices, making mixed volume computation tractable, [?].

Example 2.8 (Chemical reaction network). *A phosphorylation cycle with 3 species yields a 3-variate system where each equation is a 4-nomial. The Newton polytopes are tetrahedra. Mixed volume = 3 (vs. Bézout bound $4^3 = 64$). The sparse structure reduces the root count by 20x.*

2.5 Arithmetic and Bit Complexity

Solving fewnomial systems requires controlling the **bit size** of solutions, [?].

Problem 2.9 (Approximating radicals). *Given $a, b \in \mathbb{N}$, find $x \in \mathbb{Q}$ with $|x - a^{1/b}| < \varepsilon$.*

Lemma 2.10 (Lemma 1.11 in article). *There is an algorithm using $O(\log b)$ field operations and inequality decisions.*

Theorem 2.11 (Theorem 1.12, BMST). *Computing $a^{1/b}$ to accuracy ε requires $\Omega(\log b + \log(1/\varepsilon))$ arithmetic operations. This is asymptotically optimal.*

The key insight: **Newton's method on binomials** achieves the lower bound. For fewnomials, the tropical approximation provides a **near-optimal initial guess**, after which local refinement (Newton/homotopy) converges rapidly, [?].

2.6 Code: Tropical Approximation for Fewnomials

Listing 2.1: Tropical variety of a fewnomial

```

import numpy as np
from itertools import combinations

def tropical_variety_positive(alpha, c): \cite{rojas2025fewnomial}
    """
    alpha: k x n array of exponent vectors
    c: k array of coefficients
    Returns polyhedral cones (as halfspace intersections) for Trop_+(f)
    """
    k, n = alpha.shape
    logc = np.log(np.abs(c))
    signs = np.sign(c)

    cones = []
    for (i, j) in combinations(range(k), 2):
        # Equality of two terms:
        # alpha_i^T y + log|c_i| = alpha_j^T y + log|c_j|
        # => (alpha_i - alpha_j)^T y = log|c_j| - log|c_i|
        A_eq = alpha[i] - alpha[j]
        b_eq = logc[j] - logc[i]

        # Dominance over all other terms
        A_ub = []
        b_ub = []
        for l in range(k):
            if l == i or l == j: continue
            A_ub.append(alpha[l] - alpha[i])
            b_ub.append(logc[i] - logc[l])
        logc[l])
        cones.append((A_eq, b_eq, np.array(A_ub), np.array(b_ub)))
    return cones

def solve_tropical(alpha, c): \cite{rojas2025fewnomial}
    """Find a point in Trop_+(f) via linear programming"""
    from scipy.optimize import linprog
    cones = tropical_variety_positive(alpha, c)
    for A_eq, b_eq, A_ub, b_ub in cones:
        # Minimize 0 subject to A_eq @ y = b_eq, A_ub @ y <= b_ub
        res = linprog(np.zeros(A_eq.shape[0]), A_ub=A_ub, b_ub=b_ub,
                      A_eq=A_eq.reshape(1,-1), b_eq=[b_eq],
                      method='highs')
        if res.success:
            return res.x
    return None

```

```
# Example:  $f(x, y) = x^3 + 2xy - 5y^2$ 
alpha = np.array([[3, 0], [1, 1], [0, 2]])
c = np.array([1, 2, -5])
y = solve_tropical(alpha, c) \cite{rojas2025fewnomial}
print("Tropical point (log coords):", y)
print("Approx root (exp):", np.exp(y))
```

2.7 Applications

- **Steady states of chemical reaction networks:** Fewnomial systems from mass-action kinetics. The **deficiency zero theorem** (Feinberg) guarantees uniqueness of positive steady states when the network structure is a directed forest.
- **Power flow:** The AC power flow equations are trigonometric fewnomials. Converting via $z = e^{i\theta}$ gives polynomial fewnomials. Homotopy methods with tropical start systems solve large grids, [?].
- **Neural network verification:** The decision boundary of a ReLU network is a piecewise linear function. The number of linear regions is a fewnomial count (related to hyperplane arrangements), [?].
- **Polynomial optimization:** Minimizing a fewnomial over a polytope reduces to checking critical points (fewnomial system) and boundary, [?].

2.8 Bridge to Chapter 3

Fewnomial theory uses **Newton polytopes** and **mixed volumes** to count roots. Chapter 3 uses **spectra of matrices** (eigenvalues) to understand networks. Both are “spectral” in a broad sense: fewnomials use the *combinatorial spectrum* (exponent vectors), spectral methods use the *algebraic spectrum* (eigenvalues). Both exploit sparsity/structure to bypass dense complexity, [?].

2.9 Further Reading

- Rojas, *Algorithmic Fewnomial Theory* (this article, Notices 2025).
- Khovanskii, *Fewnomials* (AMS Translations, 1991)—the foundational monograph.
- Sturmfels, *Solving Systems of Polynomial Equations* (CBMS, 2002)—tropical geometry, [?].
- Bihan, Sottile, *New bounds for fewnomials* (2007), [?].
- Feinberg, *Foundations of Chemical Reaction Network Theory* (Springer, 2019).

2.10 Exercises

2.11 Solutions

Solution.Exercise 1 (Tropical variety): For $f(x, y, z) = x^2 + 2y^3 - 3z + 4$, the exponent vectors are $\alpha_1 = (2, 0, 0)$, $\alpha_2 = (0, 3, 0)$, $\alpha_3 = (0, 0, 1)$, $\alpha_4 = (0, 0, 0)$. The tropical variety consists of planes where two terms are maximal:

$$\text{Trop}_+(f) = \bigcup_{i < j} \{y : \alpha_i \cdot y + \log |c_i| = \alpha_j \cdot y + \log |c_j|, \text{ and term } i, j \text{ dominate}\}$$

For 4 terms, there are $\binom{4}{2} = 6$ pairs, giving 6 cones (some may be empty).

Solution.Exercise 2 (Phosphorylation cycle): The Michaelis-Menten phosphorylation cycle has species S, S^*, P, P^* with total conservation $S + S^* = s_{tot}$. Each reaction gives a polynomial equation. The system has 4 variables and 2 conservation laws, reducing to 2 independent equations in 2 variables. Mixed volume computation gives upper bound on number of positive solutions.

Solution.Exercise 3 (Descartes' rule): A univariate k -nomial $f(x) = \sum_{i=1}^k c_i x^{a_i}$ with $a_1 < a_2 < \dots < a_k$ has at most $k - 1$ positive real roots by Descartes' rule of signs: the number of positive roots is at most the number of sign changes in the coefficient sequence. For negative roots, substitute $x \rightarrow -x$ and count sign changes.

Solution.Exercise 4 (Research): Homotopy continuation with tropical start systems exploits the combinatorial structure of the polynomial. For the example $f(x, y) = x^3 + 2xy - 5y^2$, the tropical variety gives an initial guess near the actual roots. In practice, tropical starts reduce the number of homotopy paths by 50-80% compared to random starts.

Exercise 2.1. Compute the tropical variety of $f(x, y, z) = x^2 + 2y^3 - 3z + 4$. How many polyhedral cones? [?]

Exercise 2.2. For the phosphorylation cycle (3 species, each equation a 4-nomial), write the system and compute its mixed volume using `PHCpack` or `HomotopyContinuation.jl`.

Exercise 2.3. Show that a univariate k -nomial has at most $k - 1$ positive real roots (Descartes' rule of signs). What about negative roots?

Exercise 2.4 (Research). Implement a homotopy continuation solver that uses tropical varieties as start systems for 3-variate 4-nomial systems. Compare to random start systems, [?].

CHAPTER 3

Spectral Methods in Microeconomics

[?]

Benjamin Golub

Communicated by *Notices* Associate Editor Daniela De Silva

“Matrices often appear in formal models of social and economic behavior... Spectral theory offers powerful tools for understanding matrices, and economic modelers have leveraged these tools to gain considerable insight.”

— Benjamin Golub

Abstract. Eigenvalues and eigenvectors are the hidden skeleton of networked systems. This chapter shows how Perron-Frobenius theory, Katz-Bonacich centrality, and spectral decomposition unify opinion dynamics, network games, and market equilibrium. The same mathematics describes how beliefs spread, how firms compete, and how central banks should act, [?].

3.1 Introduction: Why Spectral Methods?

Mathematical notation follows standard conventions. Key terms are defined in the glossary.

In microeconomics, agents interact: consumers learn from neighbors, firms compete on networks, banks lend to each other. These interactions are encoded in **matrices**. The spectral properties of these matrices—especially the dominant eigenvalue and eigenvector—determine the system’s long-run behavior, [?].

Consider a row-stochastic matrix $\mathbf{W}$ (nonnegative, rows sum to 1) representing social influence. The DeGroot model of opinion dynamics:

$$\mathbf{x}(t+1) = \mathbf{W}\mathbf{x}(t).$$

If $\mathbf{W}$ is irreducible and aperiodic, $\mathbf{x}(t) \rightarrow \boldsymbol{\pi}^\top \mathbf{x}(0) \mathbf{1}$ where $\boldsymbol{\pi}$ is the **left Perron vector** (stationary distribution). Convergence rate is governed by $|\lambda_2|$, the second-largest eigenvalue modulus.

This single fact unifies:

- **Opinion dynamics:** How fast do beliefs converge?
- **Centrality:** Who is influential? (Katz-Bonacich = Perron vector)
- **Network games:** Equilibrium actions = $(\mathbf{I} - \beta \mathbf{G})^{-1} \mathbf{a}$; spectral radius of $\beta \mathbf{G}$ determines uniqueness, [?].
- **Information aggregation:** “Wisdom of crowds” iff $\max_i \pi_i \rightarrow 0$ as $n \rightarrow \infty$.

3.2 Perron-Frobenius Theory for Economists

[?]

Theorem 3.1 (Perron-Frobenius (Irreducible Nonnegative)). [?] Let $\mathbf{A} \in \mathbb{R}^{n \times n}$, $a_{ij} \geq 0$, irreducible. Then:

1. $\rho(\mathbf{A}) > 0$ is a simple eigenvalue (the **Perron root**).
2. There exist $\mathbf{u} > 0$, $\mathbf{v} > 0$ (right/left Perron vectors) with $\mathbf{A}\mathbf{u} = \rho\mathbf{u}$, $\mathbf{v}^\top \mathbf{A} = \rho\mathbf{v}^\top$.
3. All other eigenvalues satisfy $|\lambda| < \rho$.
4. For any $\mathbf{x} > 0$, $\mathbf{A}^t \mathbf{x} / \rho^t \rightarrow \mathbf{u}(\mathbf{v}^\top \mathbf{x} / \mathbf{v}^\top \mathbf{u})$.

Definition 3.2 (Irreducibility). $\mathbf{A}$ is **irreducible** iff its directed graph is strongly connected: for any i, j , there is a path $i \rightarrow j$. Equivalently, no permutation makes $\mathbf{A}$ block upper-triangular.

In economics, irreducibility = “the network is connected.” If the economy has isolated sectors, spectral theory applies block by block, [?].

3.3 Application 1: DeGroot Learning and Social Influence

Agents update beliefs by averaging neighbors:

$$x_i(t+1) = \sum_j w_{ij} x_j(t), \quad \mathbf{W} \text{ row-stochastic.}$$

Proposition 3.3. If $\mathbf{W}$ is irreducible and aperiodic, $\mathbf{x}(t) \rightarrow \boldsymbol{\pi}^\top \mathbf{x}(0) \mathbf{1}$ where $\boldsymbol{\pi}$ is the unique stationary distribution ($\boldsymbol{\pi}^\top \mathbf{W} = \boldsymbol{\pi}^\top$, $\boldsymbol{\pi} > 0$, $\sum \pi_i = 1$).

The **social influence** of agent i is π_i . This is exactly **eigenvector centrality** (left Perron vector of $\mathbf{W}$).

Example 3.4 (Star network). Agent 1 at center, all others peripheral. $\mathbf{W}_{1j} = 1/(n-1)$, $\mathbf{W}_{i1} = 1$ for $i > 1$. Then $\pi_1 = 1/2$, $\pi_i = 1/(2n-2)$ for $i > 1$. The center has influence $\sim 1/2$ regardless of n !

Definition 3.5 (Wisdom of Crowds). A sequence of networks is **wise** if $\max_i \pi_i^{(n)} \rightarrow 0$. Golub-Jackson (2010): wisdom holds iff no agent has “disproportionate influence.”

3.4 Application 2: Network Games and Bonacich Centrality

Agents choose actions $a_i \in \mathbb{R}$. Payoff:

$$u_i(a_i, \mathbf{a}_{-i}) = a_i \left(\alpha_i - \frac{1}{2} a_i + \beta \sum_j g_{ij} a_j \right).$$

Best response: $a_i = \alpha_i + \beta \sum_j g_{ij} a_j$. In vector form:

$$\mathbf{a} = \boldsymbol{\alpha} + \beta \mathbf{G} \mathbf{a} \quad \Rightarrow \quad \mathbf{a}^* = (\mathbf{I} - \beta \mathbf{G})^{-1} \boldsymbol{\alpha},$$

provided $\rho(\beta \mathbf{G}) < 1$.

Definition 3.6 (Katz-Bonacich Centrality). For $\mathbf{G} \geq 0$ and $\beta < 1/\rho(\mathbf{G})$,

$$\mathbf{b}(\beta, \mathbf{G}) = (\mathbf{I} - \beta \mathbf{G})^{-1} \mathbf{1} = \sum_{k=0}^{\infty} \beta^k \mathbf{G}^k \mathbf{1}.$$

Entry b_i counts weighted walks starting at i ; longer walks discounted by β^k .

The equilibrium action is $\mathbf{a}^* = \mathbf{b}(\beta, \mathbf{G}) \odot \boldsymbol{\alpha}$ (if $\boldsymbol{\alpha}$ is scalar). **Equilibrium actions are proportional to centrality.**

3.5 Application 3: Wisdom, Centralization, and Policy

The Golub-Jackson **prominent sequence** concept: a family of networks $(\mathbf{W}^{(n)})$ with n agents. They are wise iff the most influential agent's weight $\rightarrow 0$.

Theorem 3.7 (Wisdom Criterion). $(\mathbf{W}^{(n)})$ is wise iff $\max_i \pi_i^{(n)} \rightarrow 0$. For symmetric $\mathbf{W}$, this is equivalent to $\lambda_2(\mathbf{W}) \rightarrow 1$ (the spectral gap closes), [?].

Policy insight: if a regulator wants to improve information aggregation, they should **reduce the spectral gap**—i.e., make the network less centralized. This contradicts the intuition that “more connectivity is always better.” A star network is highly connected but unwise; a line network is poorly connected but wise, [?].

3.6 Code: Spectral Analysis of Networks

[?]

Listing 3.1: Network spectral analysis

```

1 using LinearAlgebra, SparseArrays, LightGraphs
2 using Statistics
3
4 function analyze_network(W; tol=1e-12)
5     n = size(W, 1)
6
7     # Check irreducibility
8     G = DiGraph(W .> tol)
9     is_irred = is_strongly_connected(G)

```

```

10
11     # Perron-Frobenius (left eigenvector for row-stochastic) \cite{
12         golub2026spectral}
13     eig = eigen(W')
14     idx = argmax(abs.(eig.values))
15     lambda1 = real(eig.values[idx])
16     pi = real(eig.vectors[:, idx])
17     pi = pi / sum(pi) # normalize
18
19     # Second eigenvalue modulus
20     lambda2 = maximum(abs.(eig.values[setdiff(1:end, idx)]))
21     spectral_gap = lambda1 - lambda2 \cite{golub2026spectral}
22
23     # Katz-Bonacich centrality for  $\beta < 1/\rho(G)$ 
24     # Assuming  $G$  is the adjacency (not stochastic)
25     function katz_centrality(G, beta)
26         rho = maximum(abs.(eigvals(G)))
27         if beta >= 1/rho
28             error("beta too large: beta*rho = $(beta*rho)")
29         end
30         (I - beta*G) \ ones(n)
31     end
32
33     # Wisdom metrics
34     max_influence = maximum(pi)
35     herfindahl = sum(pi.^2) # concentration index
36
37     return (
38         lambda1 = lambda1,
39         lambda2 = lambda2,
40         spectral_gap = spectral_gap, \cite{golub2026spectral}
41         pi = pi,
42         max_influence = max_influence,
43         herfindahl = herfindahl,
44         is_irred = is_irred
45     )
46
47     # Example: Star network
48     n = 10
49     W = zeros(n,n)
50     W[1, 2:end] .= 1/(n-1)
51     for i = 2:n
52         W[i, 1] = 1.0
53     end
54
55     res = analyze_network(W)
56     println("Star network: max_influence = ", res.max_influence)
57     println("Spectral gap = ", res.spectral_gap) \cite{golub2026spectral}
58
59     # Example: Line network
60     W2 = zeros(n,n)
61     for i = 1:n
62         neighbors = []

```

```

63     i > 1 && push!(neighbors, i-1)
64     i < n && push!(neighbors, i+1)
65     W2[i, neighbors] .= 1/length(neighbors)
66 end
67 res2 = analyze_network(W2)
68 println("Line network: max_influence = ", res2.max_influence)
69 println("Spectral gap = ", res2.spectral_gap) \cite{golub2026spectral}

```

3.7 Bridge to Chapter 4

Spectral methods analyze **linear** network dynamics: $\mathbf{x}(t+1) = \mathbf{W}\mathbf{x}(t)$. Chapter 4 introduces **nonlinear** network dynamics: deep neural networks. Remarkably, ReLU networks are *piecewise linear*—on each linear region, they are exactly a matrix multiplication. The geometry of deep learning is the geometry of how these linear regions tessellate the input space.

3.8 Further Reading

- Golub, B. (2025). *Spectral Methods in Microeconomics*. Notices of the AMS, 72(6).
- Golub, B., Jackson, M.O. (2010). *Naive Learning in Social Networks and the Wisdom of Crowds*. American Economic Journal: Microeconomics.
- Ballester, C., Calvó-Armengol, A., Zenou, Y. (2006). *Who's Who in Networks. Wanted: The Key Player*. Econometrica.
- Jackson, M.O. (2010). *Social and Economic Networks*. Princeton University Press.
- Bonacich, P. (1987). *Power and Centrality: A Family of Measures*. American Journal of Sociology.

3.9 Exercises

3.10 Solutions

Solution.Exercise 1 (Star network): For the star graph with center node 1 and $n-1$ leaf nodes, the adjacency matrix has eigenvalues $\sqrt{n-1}$, $-\sqrt{n-1}$, and 0 with multiplicity $n-2$. The stationary distribution is $\pi_1 = 1/2$ and $\pi_i = 1/(2(n-1))$ for $i \neq 1$.

Solution.Exercise 2 (Katz-Bonacich): The Katz-Bonacich centrality is $b = (I - \beta G)^{-1}\mathbf{1}$. In the DeGroot model with $G = W$ (row-stochastic), this gives $b = (I - \beta W)^{-1}\mathbf{1} = \sum_{k=0}^{\infty} \beta^k W^k \mathbf{1}$. Each term $\beta^k W^k \mathbf{1}$ represents k -step influence weighted by β^k .

Solution.Exercise 3 (Wisdom test): For $p_n = \log n/n$, the graph is connected with high probability, and the average degree is $\log n \rightarrow \infty$. The wisdom test holds: $\max_i \pi_i \rightarrow 0$. For $p_n = 1/\sqrt{n}$, the graph is disconnected with high probability, and the wisdom test fails.

Solution.Exercise 4 (Heterogeneous β_i): For heterogeneous interaction strengths, the equilibrium is $x^* = (I - \text{diag}(\beta)G)^{-1}a$. A unique equilibrium exists when $\rho(\text{diag}(\beta)G) < 1$. This condition can be checked using the spectral radius.

Exercise 3.1. For the star network with n agents, compute the exact stationary distribution π and verify $\pi_1 = 1/2$ for all $n \geq 3$.

Exercise 3.2. Show that for a row-stochastic matrix $\mathbf{W}$, the Katz-Bonacich centrality with $\mathbf{G} = \mathbf{W}$ and $\beta < 1$ is $\mathbf{b} = (\mathbf{I} - \beta\mathbf{W})^{-1}\mathbf{1}$. What does this represent in the DeGroot model?

Exercise 3.3. Implement the Golub-Jackson wisdom test: generate a sequence of Erdős-Rényi graphs $G(n, p_n)$ with $p_n = \log n/n$ and $p_n = 1/\sqrt{n}$. Which sequence is wise?

Exercise 3.4 (Research). Extend the network game to heterogeneous β_i (different interaction strengths). When does a unique equilibrium exist?

Part II

Geometric Machine Learning

CHAPTER 4

On the Geometry of Deep Learning

Randall Balestrieri, Ahmed Imtiaz Humayun, Richard G. Baraniuk
Communicated by *Notices* Associate Editor Reza Malek-Madani

“Deep networks are affine spline operators. They tessellate input space into convex polytopes, and on each tile they act as an affine transformation.”

— Balestrieri, Humayun,
Baraniuk

Abstract. This chapter develops a geometric theory of deep ReLU networks: they are **affine spline mappings** that partition input space into convex polytopes (tiles) and apply a distinct affine transformation on each. This perspective explains generalization, batch normalization, grokking, and generative model bias—and yields practical tools like SplineCam for visualization and MaGNET for debiasing, [?].

4.1 Introduction: The Black Box Problem

Mathematical notation follows standard conventions. Key terms are defined in the glossary.

Deep networks work astonishingly well but remain poorly understood. The standard view: black boxes optimized by SGD. This chapter proposes a **white-box** view:

A deep ReLU network = an affine spline operator. [?]

- Input space $\mathbb{R}^d$ is tiled into convex polytopes $\Omega = \{\Omega_j\}$.
- On each tile Ω_j , the network computes $\mathbf{x} \mapsto \mathbf{A}_j \mathbf{x} + \mathbf{b}_j$.
- The tiles and transforms depend on weights/biases and change during training.

This is not an approximation—it is an *exact* mathematical equivalence for any finite-width ReLU network, [?].

4.2 Affine Splines in 1D

A 1D continuous piecewise linear function (affine spline) is defined by:

- Knots: $t_1 < t_2 < \dots < t_K$
- Slopes: $m_0, m_1, \dots, m_K$
- Intercepts: $b_0, b_1, \dots, b_K$

with continuity constraints $m_i t_{i+1} + b_i = m_{i+1} t_{i+1} + b_{i+1}$.

A ReLU layer with n neurons computes [?]

$$\mathbf{y} = \text{ReLU}(\mathbf{W}\mathbf{x} + \mathbf{b}), \quad \mathbf{W} \in \mathbb{R}^{n \times d}.$$

Each neuron i creates a hyperplane $\{\mathbf{x} : \mathbf{w}_i^\top \mathbf{x} + b_i = 0\}$ dividing space into two halfspaces. The arrangement of n hyperplanes creates a tessellation into convex polytopes.

4.3 Deep Network Tessellation

Theorem 4.1 (Deep Network = Affine Spline). *For a depth- L ReLU network $f : \mathbb{R}^d \rightarrow \mathbb{R}^m$, there exists a finite tessellation $\Omega = \{\Omega_j\}_{j=1}^N$ of $\mathbb{R}^d$ into convex polytopes and affine maps $\mathbf{A}_j \in \mathbb{R}^{m \times d}, \mathbf{b}_j \in \mathbb{R}^m$ such that [?]*

$$f(\mathbf{x}) = \mathbf{A}_j \mathbf{x} + \mathbf{b}_j \quad \text{for all } \mathbf{x} \in \Omega_j.$$

The polytopes are the linear regions of the network.

Sketch. Induction on layers. Base: one ReLU layer creates a hyperplane arrangement (convex polytopes). Inductive step: each subsequent layer pulls back its hyperplanes through the previous layer's affine maps, which *folds* the existing tiles, creating a finer tessellation. The affine map on each tile composes with the new layer's affine map, [?]. $\square$

Example 4.2 (Number of tiles). *For a 2D input, 2-hidden-layer network with widths 20, the tessellation can have $\sim 10^4$ tiles. For deep networks, the number grows exponentially with depth (Montúfar et al., 2014).*

4.4 Geometry of the Loss Landscape

For squared error $\mathcal{L}(\boldsymbol{\theta}) = \frac{1}{2} \|f_{\boldsymbol{\theta}}(\mathbf{X}) - \mathbf{Y}\|_F^2$, the loss landscape as a function of parameters $\boldsymbol{\theta}$ is a **piecewise quadratic function**. On each region where the input-space tessellation is constant, $\mathcal{L}$ is quadratic in $\boldsymbol{\theta}$.

Theorem 4.3 (ResNet vs. ConvNet conditioning). *The condition number of the local Hessian for a ResNet (with skip connections) is bounded independent of depth, while for a plain ConvNet it grows exponentially with depth.*

This explains why ResNets train faster and more reliably: their loss landscape is better conditioned.

4.5 Batch Normalization as Data-Adaptive Tessellation

BatchNorm replaces $\mathbf{x}^{(\ell)} \mapsto \mathbf{W}^{(\ell)}\mathbf{x}^{(\ell)} + \mathbf{b}^{(\ell)}$ with

$$\mathbf{x}^{(\ell)} \mapsto \gamma \odot \frac{\mathbf{W}^{(\ell)}\mathbf{x}^{(\ell)} - \boldsymbol{\mu}}{\boldsymbol{\sigma}} + \boldsymbol{\beta},$$

where $\boldsymbol{\mu}, \boldsymbol{\sigma}$ are batch statistics.

Theorem 4.4 (BatchNorm Geometry). *At initialization, BatchNorm adapts the hyperplane arrangement to concentrate tile boundaries near the training data. Without BatchNorm, tiles are randomly oriented; with BatchNorm, hyperplanes align to data density.*

This is why BatchNorm accelerates training: it gives the network a *data-aware initialization* of its tessellation.

4.6 The Dynamics of Learning: Grokking

Grokking (Power et al., 2022): A network trained well past zero training error suddenly improves test accuracy—the tessellation *migrates* from overfitting (many small tiles around training points) to generalizing (fewer, larger tiles focused on decision boundaries).

Definition 4.5 (Local Complexity). *For a point $\mathbf{x}$ in input space, the **local complexity** is the number of tiles intersecting a neighborhood $\mathcal{N}_\epsilon(\mathbf{x})$.*

During grokking, local complexity *decreases* around training data and *increases* at decision boundaries.

Listing 4.1: SplineCam: Visualizing network tessellation

```

0 import torch
1 import numpy as np
2 from splinecam import SplineCam
3
4 # Load trained model
5 model = torch.load('mnist_mlp.pt')
6
7 # Compute tessellation on a 2D slice defined by 3 training points
8 sc = SplineCam(model)
9 slice_2d = sc.compute_slice(points=[img1, img2, img3], grid_size=200)
10
11 # Visualize
12 sc.plot_tessellation(slice_2d, color_by='tile_id')
13 sc.plot_affine_maps(slice_2d)
14 sc.plot_decision_boundary(slice_2d)

```

4.7 Generative Models: MaGNET Debiasing

A ReLU generative model $G : \mathbb{R}^k \rightarrow \mathbb{R}^d$ maps a low-dimensional latent space to a *piecewise affine manifold* in data space. Each latent tile Ω_j maps affinely to a manifold tile $G(\Omega_j)$, [?].

Theorem 4.6 (Volume distortion). *The volume scaling factor from latent tile Ω_j to data tile $G(\Omega_j)$ is $|\det \mathbf{A}_j|$ where $\mathbf{A}_j$ is the affine map’s linear part.*

MaGNET (Balestrierio et al., 2022): To sample uniformly from the data manifold, sample latent points with probability $\propto 1/|\det \mathbf{A}_j|$. This corrects for the *mode collapse* and *bias* of standard uniform latent sampling.

Listing 4.2: MaGNET debiasing for StyleGAN

```

0 import torch
1 from magnet import magnet_sampler
2
3 # Load pretrained StyleGAN2
4 G = load_stylegan2('ffhq.pkl')
5
6 # Compute Jacobian determinants on a grid
7 latent_grid = torch.randn(10000, 512)
8 with torch.no_grad():
9     jac_dets = batch_jacobian_det(G, latent_grid)
10
11 # MaGNET importance weights
12 weights = 1.0 / (jac_dets + 1e-8)
13 weights = weights / weights.sum()
14
15 # Sample debiased latents
16 debiased_latents = latent_grid[torch.multinomial(weights, 1000)]
17
18 # Generate fair samples
19 fair_images = G(debiased_latents)

```

4.8 Code: Affine Spline Layer

Listing 4.3: Exact affine spline representation of a ReLU layer

```

1 using LinearAlgebra, SparseArrays
2
3 struct AffineSplineLayer
4     tiles::Vector{HPolyhedron} # Convex polytopes (H-representation)
5     A::Vector{Matrix{Float64}} # Linear maps per tile
6     b::Vector{Vector{Float64}} # Biases per tile
7 end
8
9 function relu_layer_to_spline(W, b)
10     n, d = size(W)
11     tiles = HPolyhedron[]
12     A_list = Matrix{Float64}[]
13     b_list = Vector{Float64}[]
14
15     # Enumerate all 2^n sign patterns
16     for signs in Iterators.product(fill([-1, 1], n)...)
17         # sign_i = 1 means w_i^T x + b_i >= 0 (active)
18         # sign_i = -1 means w_i^T x + b_i <= 0 (inactive)
19         H = []

```

```

20     h = []
21     for i in 1:n
22         if signs[i] == 1
23             push!(H, W[i, :])
24             push!(h, b[i])
25         else
26             push!(H, -W[i, :])
27             push!(h, -b[i])
28         end
29     end
30     # Check feasibility (non-empty polytope)
31     if is_feasible(H, h)
32         poly = HPolyhedron(hcat(H...), h)
33         push!(tiles, poly)
34
35         # Affine map on this tile:  $y = D(Wx + b)$  where  $D = \text{diag}(\max(\text{sign}, 0))$ 
36         D = Diagonal([s == 1 ? 1.0 : 0.0 for s in signs])
37         push!(A_list, D * W)
38         push!(b_list, D * b)
39     end
40 end
41 return AffineSplineLayer(tiles, A_list, b_list)
42 end
43
44 function evaluate_spline(spline::AffineSplineLayer, x)
45     for (i, tile) in enumerate(spline.tiles)
46         if x in tile
47             return spline.A[i] * x + spline.b[i]
48         end
49     end
50     error("Point not in any tile")
51 end

```

4.9 Bridge to Chapter 5

Deep networks are affine splines. Equivariant networks (Chapter 5) constrain these splines to respect symmetries: the tessellation and affine maps must commute with group actions. Topological deep learning (Chapter 6) uses invariants like the Euler characteristic to classify the shapes these splines produce, [?].

4.10 Further Reading

- Balestrierio, Humayun, Baraniuk (2025). *On the Geometry of Deep Learning*. Notices, 72(4).
- Montúfar et al. (2014). *On the Number of Linear Regions of Deep Neural Networks*. NIPS.
- Balestrierio et al. (2022). *MaGNET: Maximum Entropy Generative Network*. ICML.
- SplineCam: <https://github.com/randall-b/splinecam>
- MaGNET: <https://github.com/randall-b/magnet>

4.11 Exercises

4.12 Solutions

Solution.Exercise 1 (1D ReLU): For a 1D ReLU network with 2 neurons and slopes $m_1 = 3, m_2 = -2$, the output is $f(x) = 3 \max(0, x) - 2 \max(0, x - 1)$. This gives a piecewise linear function with breakpoints at $x = 0$ and $x = 1$:

- $x < 0$: $f(x) = 0$
- $0 \leq x < 1$: $f(x) = 3x$
- $x \geq 1$: $f(x) = 3x - 2(x - 1) = x + 2$

Solution.Exercise 2 (Tile count): A ReLU network with L layers and n neurons per layer partitions $\mathbb{R}^d$ into at most $\prod_{l=1}^L (1 + n_l)$ tiles in 1D. For d dimensions, the bound is $\sum_{j=0}^d \binom{N}{j}$ where $N = \sum n_l$ is the total number of neurons.

Solution.Exercise 3 (SplineCam): SplineCam visualizes the piecewise-affine structure by coloring each tile by the affine transformation applied. The implementation uses ray tracing: for each pixel, trace a ray through input space and identify which tile the ray passes through.

Solution.Exercise 4 (Research): Deep networks with depth L and width n can represent functions with fractal-like decision boundaries. The number of linear regions grows exponentially with depth but only polynomially with width, explaining why deep narrow networks can be more expressive than shallow wide ones.

Exercise 4.1. For a 1-hidden-layer ReLU network with 3 neurons in 2D, draw the hyperplane arrangement and label the affine map on each tile, [?].

Exercise 4.2. Implement the exact affine spline representation of a 2-layer ReLU network. Verify that evaluating the spline matches the network output for random inputs, [?].

Exercise 4.3. Use SplineCam to visualize the tessellation of a trained MNIST classifier. How does the tile structure differ between correctly and incorrectly classified examples?

Exercise 4.4 (Research). Prove that the number of linear regions of a deep ReLU network is upper-bounded by $\prod_{\ell=1}^L \sum_{j=0}^d \binom{n_\ell}{j}$ where n_ℓ is the width of layer ℓ , [?].

CHAPTER 5

What Is... an Equivariant Neural Network?

[?]

Lek-Heng Lim, Bradley J. Nelson

Communicated by *Notices* Associate Editor Emilie Purvine

“An equivariant neural network is an equivariant map constructed from alternately composing equivariant linear maps with nonlinear ones.”

— Lim & Nelson

[?]

Abstract. Symmetry is everywhere in science: molecular structures, physical laws, images. Standard neural networks ignore symmetry, wasting capacity learning what physics already dictates. Equivariant neural networks bake symmetry into the architecture: if the input transforms, the output transforms predictably. This chapter gives the mathematical foundation: group representations, intertwining operators, and the universal approximation theorem for equivariant networks, [?].

5.1 Introduction: Why Symmetry Matters

Mathematical notation follows standard conventions. Key terms are defined in the glossary.

A molecule rotated in space is the *same* molecule. A cat translated in an image is the same cat. The laws of physics are translation/rotation invariant. Yet a standard CNN must *learn* translation invariance from data—a waste of parameters and samples.

An **equivariant neural network** guarantees: [?]

$$f(g \cdot x) = g \cdot f(x) \quad \forall g \in G,$$

where G is a symmetry group acting on input space $\mathcal{X}$ and output space $\mathcal{Y}$, [?].

- **Invariance:** $f(g \cdot x) = f(x)$ (e.g., classification: rotated cat $\rightarrow$ “cat”)

- **Equivariance:** $f(g \cdot x) = g \cdot f(x)$ (e.g., segmentation: rotated input $\rightarrow$ rotated mask)
- **Covariance:** Output transforms by a *different* representation of G .

5.2 Group Representations

Definition 5.1 (Group Representation). A **representation** of a group G on a vector space V is a homomorphism $\rho : G \rightarrow \text{GL}(V)$. For $g \in G$, $\rho(g)$ is an invertible linear map on V , [?].

Key examples for deep learning:

- **Translation** ($G = \mathbb{Z}^2$ or $\mathbb{R}^2$): $\rho(g)_{ij} = x_{i-g}$ (shift)
- **Rotation** ($G = \text{SO}(2)$ or $\text{SO}(3)$): $\rho(g) = R_g$ (rotation matrix)
- **Permutation** ($G = S_n$): $\rho(\sigma)_{ij} = \delta_{i, \sigma(j)}$
- **Permutation + Sign** ($G = S_n \times \{\pm 1\}^n$): for fermionic wavefunctions

Theorem 5.2 (Schur's Lemma). Let $\rho_1 : G \rightarrow \text{GL}(V_1)$, $\rho_2 : G \rightarrow \text{GL}(V_2)$ be irreducible representations. Any linear map $T : V_1 \rightarrow V_2$ satisfying $T\rho_1(g) = \rho_2(g)T$ for all g is either zero or an isomorphism. If $V_1 = V_2$ and $\rho_1 = \rho_2$ irreducible, then $T = \lambda I$.

Consequence: Equivariant linear layers are highly constrained! For irreps, they are just scalar multiples of the identity.

5.3 Equivariant Linear Layers

We want linear $\mathbf{W} : V_1 \rightarrow V_2$ such that $\mathbf{W}\rho_1(g) = \rho_2(g)\mathbf{W}$.

Example 5.3 (Convolution = Translation Equivariance). Let $G = \mathbb{Z}^2$ act on images by translation. An equivariant linear map is exactly a **convolution**: [?]

$$(\mathbf{W} * x)[i] = \sum_j w[i - j]x[j].$$

Proof: $(T_g x)[i] = x[i - g]$. Then $(T_g \mathbf{W} * x)[i] = \sum_j w[i - j]x[j - g] = (\mathbf{W} * T_g x)[i]$.

Theorem 5.4 (General Solution for Linear Equivariant Maps). For finite G , the space of equivariant linear maps between representations ρ_1, ρ_2 is [?]

$$\text{Hom}_G(\rho_1, \rho_2) = \left\{ \frac{1}{|G|} \sum_{g \in G} \rho_2(g) M \rho_1(g^{-1}) : M \in \mathbb{R}^{d_2 \times d_1} \right\}.$$

This is the projection onto the equivariant subspace (Reynolds operator), [?].

5.4 Equivariant Nonlinearities

Linear equivariance is well understood. Nonlinearities break equivariance unless carefully chosen.

Proposition 5.5. *A pointwise nonlinearity $\sigma : \mathbb{R} \rightarrow \mathbb{R}$ is G -equivariant iff it commutes with the action of G on $\mathbb{R}$, [?].*

For permutation groups S_n acting by permuting coordinates, **any** pointwise σ is equivariant (since $\sigma(x_{\pi(i)}) = \sigma(x)_{\pi(i)}$). This is why ReLU works in standard MLPs and CNNs! [?]

For rotation groups $SO(3)$, pointwise σ on vector components is *not* equivariant (rotation mixes components). Solutions: [?]

1. **Scalarization:** Apply σ to *invariants* (norms, dot products)
2. **Gated nonlinearities:** $\sigma(\mathbf{x}) = \sigma_{\text{scalar}}(\|\mathbf{x}\|) \cdot \mathbf{x} / \|\mathbf{x}\|$
3. **Tensor product:** Work in irreps; Clebsch-Gordan coefficients give equivariant polynomial maps [?]

5.5 Architecture: Equivariant Neural Networks

Definition 5.6 (Equivariant Neural Network). *A feedforward network where:*

1. Each layer $\mathbf{x}^{(\ell+1)} = \sigma(\mathbf{W}^{(\ell)}\mathbf{x}^{(\ell)} + \mathbf{b}^{(\ell)})$
2. $\mathbf{W}^{(\ell)}$ is G -equivariant (intertwiner) [?]
3. σ is G -equivariant (pointwise on scalars, gated on vectors) [?]
4. $\mathbf{b}^{(\ell)}$ is G -invariant (typically zero or learned invariant features)

The composition of equivariant maps is equivariant, so the whole network is equivariant, [?].

Example 5.7 ($SO(3)$ -Equivariant Network for Molecules). 1. Input: atomic positions $\mathbf{r}_i \in \mathbb{R}^3$, types z_i

2. Embed types as scalars (invariant features)
3. Spherical harmonics $Y_l^m(\mathbf{r}_i - \mathbf{r}_j)$ as equivariant features (transform as $D^l(R)$) [?]
4. Tensor product layers using Clebsch-Gordan coefficients
5. Gated nonlinearities preserve equivariance
6. Output: scalar (energy), vector (forces), tensor (stress)

*This is the architecture behind **NequiP**, **MACE**, **SE(3)-Transformer**.*

5.6 Code: Implementing an Equivariant Layer

Listing 5.1: SO(2)-equivariant convolution using steerable filters

```

1 import torch
2 import torch.nn as nn
3 import torch.nn.functional as F
4 from math import factorial
5
6 class SO2Conv(nn.Module):
7     """
8     Steerable CNN for SO(2) equivariance.
9     Features transform as Fourier modes:  $f_l(r, \theta) = \sum_m f_{\{l,m\}}(r) e^{im\theta}$ 
10    """
11    def __init__(self, in_channels, out_channels, max_freq, kernel_size=3):
12        :
13        super().__init__()
14        self.max_freq = max_freq
15        self.in_channels = in_channels
16        self.out_channels = out_channels
17
18        # Learnable radial profiles for each frequency
19        # For each input freq  $l_{in}$  and output freq  $l_{out}$ , we need a radial
20        # kernel
21        #  $l_{out} = l_{in} + m$  where  $m$  is the filter's frequency
22        self.radial_mlps = nn.ModuleDict()
23
24        for l_in in range(max_freq + 1):
25            for l_out in range(max_freq + 1):
26                m = l_out - l_in # filter frequency
27                if abs(m) <= max_freq:
28                    key = f"{l_in}_{l_out}"
29                    # Radial profile: small MLP taking radius as input
30                    self.radial_mlps[key] = nn.Sequential(
31                        nn.Linear(1, 16), nn.ReLU(),
32                        nn.Linear(16, in_channels * out_channels)
33                    )
34
35    def forward(self, x):
36        # x: [B, C_in, L_max+1, 2, H, W]
37        # complex features: real/imag for each frequency
38        B, C_in, L, _, H, W = x.shape
39
40        # Convert to polar coordinates for convolution
41        # This is a simplified version; real implementation uses FFT
42        out = torch.zeros(B, self.out_channels, self.max_freq+1, 2, H, W,
43                          device=x.device)
44
45        for l_in in range(self.max_freq + 1):
46            for l_out in range(self.max_freq + 1):
47                key = f"{l_in}_{l_out}"
48                if key not in self.radial_mlps:
49                    continue

```

```

47         # Get radial weights
48         # In practice, we evaluate radial profile on grid and do
49         # convolution via FFT
50         pass # Full implementation requires careful FFT handling
51
52     return out
53
54 # Simpler: Scalar + Vector fields for SO(3)
55 class S03Conv(nn.Module):
56     """Tensor field network style: scalars (l=0) and vectors (l=1)"""
57     def __init__(self, in_scalar, in_vector, out_scalar, out_vector):
58         super().__init__()
59         # Scalar -> Scalar: standard conv
60         self.ss = nn.Conv2d(in_scalar, out_scalar, 3, padding=1)
61         # Scalar -> Vector: multiply by learned vector field
62         self.sv = nn.Conv2d(in_scalar, out_vector * 3, 3, padding=1)
63         # Vector -> Scalar: dot product with learned vector
64         self.vs = nn.Conv2d(in_vector * 3, out_scalar, 3, padding=1)
65         # Vector -> Vector: linear combination + cross product
66         self.vv = nn.Conv2d(in_vector * 3, out_vector * 3, 3, padding=1)
67
68     def forward(self, scalars, vectors):
69         # scalars: [B, C_s, H, W]
70         # vectors: [B, C_v, 3, H, W] (3 components)
71         B, C_v, _, H, W = vectors.shape
72         v_flat = vectors.view(B, C_v * 3, H, W)
73
74         out_s = self.ss(scalars) + self.vs(v_flat)
75         out_v_flat = self.sv(scalars) + self.vv(v_flat)
76         out_v = out_v_flat.view(B, -1, 3, H, W)
77
78     return out_s, out_v

```

5.7 Universal Approximation

Theorem 5.8 (Equivariant Universal Approximation). *Let G be a compact group. An equivariant neural network with: [?]*

- Equivariant linear layers spanning $\text{Hom}_G(\rho_1, \rho_2)$
- Equivariant nonlinearities separating orbits

can approximate any continuous G -equivariant function $f : V_1 \rightarrow V_2$ uniformly on compact sets, [?].

Key insight: The nonlinearity must *not* be G -invariant (e.g., ReLU on scalars works for S_n but not $\text{SO}(3)$). Gated nonlinearities provide the needed separation.

5.8 Bridge to Chapter 6

Equivariant networks constrain the **symmetry** of the affine spline tessellation (Chapter 4). Topological deep learning (Chapter 6) studies the **shape** of the decision boundaries and feature spaces using

algebraic topology invariants. Both are structural priors: symmetry reduces parameters; topology measures complexity, [?].

5.9 Further Reading

- Lim, Nelson (2023). *What Is... an Equivariant Neural Network?* Notices, 70(4).
- Cohen, Welling (2016). *Group Equivariant CNNs*. ICML.
- Thomas et al. (2018). *Tensor Field Networks*. NeurIPS.
- Fuchs et al. (2020). *SE(3)-Transformers*. ICML.
- Geiger, Smidt (2022). *e3nn: Euclidean Neural Networks*. <https://github.com/e3nn/e3nn>

5.10 Exercises

5.11 Solutions

Solution.Exercise 1 (SO(2) equivariance): For a rotation R_θ in 2D, an equivariant function satisfies $f(R_\theta x) = R_\theta f(x)$. For a linear map $f(x) = Wx$, this requires $WR_\theta = R_\theta W$, meaning W must be of the form $\begin{pmatrix} a & -b \\ b & a \end{pmatrix}$.

Solution.Exercise 2 (Steerable filters): A steerable filter bank satisfies $R_\theta \phi(x) = \sum_k c_k(\theta) \phi_k(x)$. For SO(2), the steerable basis functions are $e^{in\theta}$ for $n \in \mathbb{Z}$. The rotation coefficients $c_k(\theta)$ are trigonometric polynomials.

Solution.Exercise 3 (Universal approximation): The universal approximation theorem for equivariant networks states that any equivariant continuous function can be approximated by an equivariant neural network with sufficient width. The proof uses the fact that equivariant layers preserve the symmetry structure.

Solution.Exercise 4 (Research): For non-compact groups like the Euclidean group SE(3), equivariant networks require convolution with steerable filters and pooling operations that respect the group structure. The computational cost is higher than for compact groups due to the infinite nature of translations.

Exercise 5.1. Show that a standard CNN with ReLU is translation-equivariant but not rotation-equivariant, [?].

Exercise 5.2. Implement the Reynolds operator for $G = D_4$ (dihedral group of square) acting on 2×2 images. Verify it projects any matrix onto the D_4 -equivariant subspace, [?].

Exercise 5.3. For SO(2) acting on $\mathbb{R}^2$ by rotation, find all equivariant linear maps $\mathbb{R}^2 \rightarrow \mathbb{R}^2$. What about $\mathbb{R}^2 \rightarrow \mathbb{R}$? [?]

Exercise 5.4 (Research). Design an equivariant architecture for the group $G = \text{SE}(3)$ (rigid motions) acting on point clouds. Compare to PointNet++ on ModelNet40, [?].

CHAPTER 6

Topology Meets Machine Learning: The Euler Characteristic Transform

[?]

Bastian Rieck

Communicated by *Notices* Associate Editor Emilie Purvine

“The Euler Characteristic Transform serves as a sufficient shape statistic, yielding an injective mapping for finite geometric simplicial complexes.”

— Bastian Rieck

[?]

Abstract. The Euler Characteristic Transform (ECT) converts a shape (simplicial complex) into a function on the sphere: for each direction, filter the shape by height and compute the Euler characteristic. The ECT is invertible, differentiable (via sigmoid approximation), and provides a topological inductive bias for deep learning. This chapter introduces the ECT, its mathematical foundations, and applications to shape classification and point cloud reconstruction, [?].

6.1 Introduction: Topology as Inductive Bias

Mathematical notation follows standard conventions. Key terms are defined in the glossary.

Deep learning excels at learning features from data. But what if we already know something about the *structure* of the data? For shapes (3D meshes, point clouds, graphs), topology provides invariants that are robust to deformation, noise, and discretization, [?].

The **Euler Characteristic Transform (ECT)** bridges topology and machine learning: [?]

- **Input:** a geometric simplicial complex $K \subset \mathbb{R}^d$
- **Topological:** Based on Euler characteristic $\chi = \beta_0 - \beta_1 + \beta_2 - \dots$
- **Multiscale:** Filtration by direction gives a curve, not just a number

- **Complete:** The ECT is injective (Turner et al., 2014; Ghrist et al., 2018)
- **Differentiable:** Sigmoid approximation enables gradient-based learning
- **Efficient:** $O(\text{faces} \times \text{directions})$, trivial to parallelize

6.2 The Euler Characteristic

[?]

Definition 6.1 (Euler Characteristic). *[?] For a finite simplicial complex K with f_k simplices of dimension k :*

$$\chi(K) = \sum_{k=0}^d (-1)^k f_k = \sum_{k=0}^d (-1)^k \beta_k.$$

β_k are Betti numbers (ranks of homology groups). χ is a homotopy invariant.

Example 6.2 (Platonic Solids). *All five Platonic solids are homeomorphic to S^2 , so $\chi = 2$.*

	Vertices	Edges	Faces
<i>Tetrahedron</i>	4	6	4
<i>Cube</i>	8	12	6
<i>Octahedron</i>	6	12	8
<i>Dodecahedron</i>	20	30	12
<i>Icosahedron</i>	12	30	20

$V - E + F = 2$ for all.

6.3 The Euler Characteristic Transform

[?]

Let $K \subset \mathbb{R}^d$ be a geometric simplicial complex. For a direction $\mathbf{v} \in S^{d-1}$ and threshold $t \in \mathbb{R}$, define the **sublevel set**

$$K_{\mathbf{v},t} = \{\sigma \in K : \max_{\mathbf{x} \in \sigma} \langle \mathbf{v}, \mathbf{x} \rangle \leq t\}.$$

Definition 6.3 (Euler Characteristic Curve (ECC)). *[?] The ECC in direction $\mathbf{v}$ is the function*

$$\text{ECC}_K(\mathbf{v}, t) = \chi(K_{\mathbf{v},t}).$$

Definition 6.4 (Euler Characteristic Transform (ECT)). *[?] The ECT is the map*

$$\text{ECT}_K : S^{d-1} \rightarrow \{S^{d-1} \rightarrow \mathbb{Z}\}, \quad \mathbf{v} \mapsto \text{ECC}_K(\mathbf{v}, \cdot).$$

It assigns to each direction its Euler Characteristic Curve, [?].

Theorem 6.5 (Invertibility of ECT). *For finite geometric simplicial complexes in $\mathbb{R}^d$, the ECT is injective: $\text{ECT}_K = \text{ECT}_L \implies K = L$. Moreover, a finite number of directions suffices (Curry et al., 2022).*

The proof uses Euler calculus: the ECT is a **Radon transform** of the indicator function, invertible by the Euler integral.

6.4 Discretization and Machine Learning

For computation, we discretize:

1. **Directions:** Sample D directions $\mathbf{v}_1, \dots, \mathbf{v}_D \in S^{d-1}$.
2. **Thresholds:** For each direction, sample T thresholds $t_1 < \dots < t_T$.

The ECT becomes a $D \times T$ matrix (image). Each column is an ECC.

Listing 6.1: Discretized ECT computation

```

0 import numpy as np
1 from scipy.spatial import Delaunay
2
3 def compute_ect(vertices, faces, directions, thresholds):
4     """
5     vertices: N x 3 array
6     faces: M x 3 array of vertex indices (triangles)
7     directions: D x 3 array (unit vectors)
8     thresholds: T array of heights
9     Returns: D x T matrix of Euler characteristics
10    """
11    D, T = len(directions), len(thresholds)
12    ect = np.zeros((D, T), dtype=int)
13
14    for d, v in enumerate(directions):
15        # Project vertices onto direction
16        heights = vertices @ v # N array
17
18        for t_idx, t in enumerate(thresholds):
19            # Vertices below threshold
20            v_below = heights <= t
21
22            # Edges below threshold
23            edges = set()
24            for face in faces:
25                for i, j in [(0,1), (1,2), (2,0)]:
26                    edges.add(tuple(sorted((face[i], face[j]))))
27            edges = np.array(list(edges))
28            e_below = np.all(v_below[edges], axis=1)
29            f1 = np.sum(e_below)
30
31            # Faces below threshold
32            f_below = np.all(v_below[faces], axis=1)
33            f2 = np.sum(f_below)
34
35            f0 = np.sum(v_below)
36            ect[d, t_idx] = f0 - f1 + f2
37
38    return ect
39
40 # Example: Tetrahedron
41 verts = np.array([[1,1,1], [-1,-1,1], [-1,1,-1], [1,-1,-1]]) / np.sqrt(3)
42 faces = np.array([[0,1,2], [0,1,3], [0,2,3], [1,2,3]])

```

```

43 directions = np.random.randn(100, 3)
44 directions = directions / np.linalg.norm(directions, axis=1, keepdims=True)
45 thresholds = np.linspace(-1, 1, 50)
46
47 ect_matrix = compute_ect(verts, faces, directions, thresholds)
48 print("ECT shape:", ect_matrix.shape) # (100, 50)

```

6.5 Differentiable ECT for Deep Learning

The ECC is a step function (changes only when threshold crosses a vertex height). To use in deep learning, we need a **smooth approximation**.

Proposition 6.6 (Sigmoid Approximation). *Replace the indicator $\mathbb{1}_{h \leq t}$ with $\sigma_\beta(t - h) = (1 + e^{-\beta(t-h)})^{-1}$. Then*

$$\widetilde{\text{ECC}}_K(\mathbf{v}, t) = \sum_{\sigma \in K} (-1)^{\dim \sigma} \prod_{\mathbf{x} \in \sigma} \sigma_\beta(t - \langle \mathbf{v}, \mathbf{x} \rangle).$$

As $\beta \rightarrow \infty$, $\widetilde{\text{ECC}} \rightarrow \text{ECC}$ pointwise.

This makes the ECT a **differentiable layer** in a neural network. The input is a point cloud/mesh; the output is an ECT matrix; gradients flow back to vertex positions.

6.6 Applications

1. **Shape Classification:** ECT matrix as feature vector $\rightarrow$ SVM or CNN. Competitive with graph neural networks on geometric graphs (Rieck et al., 2024).
2. **Point Cloud Reconstruction:** Learn vertex positions by minimizing ECT distance to target (Rieck, 2025).
3. **Topological Regularization:** Add ECT loss to enforce shape priors in generative models.
4. **Protein Shape Analysis:** ECT of molecular surfaces for binding site prediction.

Listing 6.2: ECT-based shape classification

```

0 import torch
1 import torch.nn as nn
2
3 class ECTClassifier(nn.Module):
4     def __init__(self, n_directions=100, n_thresholds=50, n_classes=10):
5         super().__init__()
6         self.ect_layer = DifferentiableECT(n_directions, n_thresholds)
7         self.classifier = nn.Sequential(
8             nn.Flatten(),
9             nn.Linear(n_directions * n_thresholds, 256),
10            nn.ReLU(),
11            nn.Linear(256, n_classes)
12        )
13

```

```
14     def forward(self, vertices, faces):
15         # vertices: [B, N, 3], faces: [B, M, 3]
16         ect = self.ect_layer(vertices, faces) # [B, D, T]
17         return self.classifier(ect)
18
19 # Differentiable ECT layer (using sigmoid approximation)
20 class DifferentiableECT(nn.Module):
21     def __init__(self, n_directions, n_thresholds, beta=10.0):
22         super().__init__()
23         # Random directions on sphere
24         dirs = torch.randn(n_directions, 3)
25         self.register_buffer('directions', dirs / dirs.norm(dim=1, keepdim
26                               =True))
27         self.register_buffer('thresholds', torch.linspace(-1, 1,
28                               n_thresholds))
29         self.beta = beta
30
31     def forward(self, vertices, faces):
32         B, N, _ = vertices.shape
33         D, T = len(self.directions), len(self.thresholds)
34
35         ect = torch.zeros(B, D, T, device=vertices.device)
36
37         for d, v in enumerate(self.directions):
38             heights = vertices @ v # [B, N]
39
40             for t_idx, t in enumerate(self.thresholds):
41                 # Sigmoid approximation for vertices
42                 v_below = torch.sigmoid(self.beta * (t - heights)) # [B,
43                               N]
44                 f0 = v_below.sum(dim=1) # [B]
45
46                 # For edges/faces, need to multiply sigmoids
47                 # Simplified: this is a sketch; real implementation uses
48                 # batched ops
49                 # ... (full implementation computes edges/faces similarly)
50
51                 ect[:, d, t_idx] = f0 # + edge/face terms
52
53         return ect
```

6.7 Code Repository

- Python package: `pip install ect` (<https://github.com/bastianrieck/ect>)
- Differentiable ECT in PyTorch: <https://github.com/bastianrieck/diffect>
- Tutorial notebooks: <https://topology.rocks/ect> [?]

6.8 Bridge to Chapter 7

The ECT provides a **topological summary** of shape. Chapter 7 addresses a different kind of summary: **uncertainty quantification**. Instead of describing a shape, UQ describes what we *don't know* about a model's predictions. Both are “compression” operations: ECT compresses geometry into curves; UQ compresses probability distributions into entropy/KL divergence.

6.9 Further Reading

- Rieck (2025). *Topology Meets Machine Learning: An Introduction Using the Euler Characteristic Transform*. Notices, 72(7), [?].
- Turner et al. (2014). *Euler Integral Transforms*. J. Appl. Comput. Topol.
- Ghrist et al. (2018). *Persistent Homology and Euler Integral Transforms*. J. Appl. Comput. Topol.
- Curry et al. (2022). *How Many Directions Determine a Shape?* Trans. AMS.
- Munch (2025). *The Euler Characteristic Transform*. Notices (overview), [?].

6.10 Exercises

6.11 Solutions

Solution.Exercise 1 (ECT injectivity): The Euler Characteristic Transform is injective for finite simplicial complexes by the inversion formula: K can be recovered from $\text{ECT}(K)$ by evaluating at all directions and using the fact that χ is additive under union.

Solution.Exercise 2 (Sigmoid approximation): The sigmoid approximation to the Heaviside step function is $H_\epsilon(t) = \frac{1}{1+e^{-t/\epsilon}}$. As $\epsilon \rightarrow 0$, $H_\epsilon \rightarrow H$. The derivative is $H'_\epsilon(t) = \frac{e^{-t/\epsilon}}{\epsilon(1+e^{-t/\epsilon})^2}$, which becomes a delta function as $\epsilon \rightarrow 0$.

Solution.Exercise 3 (Shape classification): The ECT of a shape K is a function on S^{d-1} . To classify shapes, we can use the L^2 norm $\|\text{ECT}(K_1) - \text{ECT}(K_2)\|_{L^2}$ as a distance. This distance is a topological invariant and is robust to noise and discretization.

Solution.Exercise 4 (Persistent homology): The ECT determines the persistent homology of a filtration by sublevel sets. The Euler characteristic of the sublevel set at height t is a weighted sum of Betti numbers: $\chi(X_t) = \beta_0 - \beta_1 + \beta_2 - \dots$.

Exercise 6.1. Compute the ECT of a cube with vertices at $(\pm 1, \pm 1, \pm 1)$ for directions $(\pm 1, 0, 0)$, $(0, \pm 1, 0)$, $(0, 0, \pm 1)$. Verify $\chi = 2$ for all thresholds beyond the cube extent.

Exercise 6.2. Implement the exact (non-differentiable) ECT for a triangle mesh. Compare the ECT matrices of a sphere and a torus. How many directions are needed to distinguish them?

Exercise 6.3. Use the differentiable ECT to reconstruct a point cloud of a chair from its ECT matrix. Initialize with random points and optimize via gradient descent.

Exercise 6.4 (Research). Prove that the ECT of a simplicial complex determines its persistent homology (or vice versa), [?].

Part III

Uncertainty, Dynamics, and Inverse Problems

CHAPTER 7

Taming Uncertainty: The Rise of Uncertainty Quantification

Nan Chen, Stephen Wiggins, Marios Andreou

Communicated by *Notices* Associate Editor Reza Malek-Madani

“All models are wrong, but some are useful.”

— George Box

Abstract. Uncertainty Quantification (UQ) provides a systematic framework for characterizing how unknowns propagate through mathematical models. This chapter is a self-contained tutorial: from Shannon entropy and KL divergence through data assimilation (Kalman/Ensemble filters) to stochastic surrogate modeling for turbulent PDEs. Complete with MATLAB/Python code for every algorithm, [?].

7.1 Introduction: Why Quantify Uncertainty?

Mathematical notation follows standard conventions. Key terms are defined in the glossary.

Mathematical models have imperfections:

- **Parametric uncertainty:** Unknown coefficients (e.g., viscosity, reaction rates) [?]
- **Structural uncertainty:** Missing physics, simplified equations [?]
- **Initial/boundary condition uncertainty:** Noisy or sparse observations [?]
- **Numerical uncertainty:** Discretization, rounding errors [?]

UQ asks: *How do these uncertainties affect the model output?* The answer is a probability distribution (or at least its moments) on the output space.

Two fundamental information measures:

Shannon Entropy $H(p) = -\int p(x) \log p(x) dx$ — spread of a single distribution

KL Divergence $D_{\text{KL}}(p||q) = \int p(x) \log \frac{p(x)}{q(x)} dx$ — difference between two distributions

7.2 Entropy and Relative Entropy

For a Gaussian $\mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\Sigma})$ in d dimensions:

$$H = \frac{1}{2} \log \det(2\pi e \boldsymbol{\Sigma}) = \frac{d}{2} (1 + \log 2\pi) + \frac{1}{2} \log \det \boldsymbol{\Sigma}.$$

For non-Gaussian distributions, entropy captures features variance cannot:

- **Skewness:** Asymmetric tails $\rightarrow$ higher entropy
- **Multimodality:** Multiple peaks $\rightarrow$ higher entropy
- **Heavy tails:** Extreme events $\rightarrow$ higher entropy

Example 7.1 (Gamma vs. Gaussian with same variance). Gamma(2, 1) has variance 2, entropy ≈ 1.68 . $\mathcal{N}(0, 2)$ has entropy ≈ 1.42 . Same variance, different uncertainty! [?]

KL divergence measures the cost of using approximate model q when truth is p :

$$D_{\text{KL}}(p||q) = \underbrace{\frac{1}{2} \text{tr}(\boldsymbol{\Sigma}_q^{-1} \boldsymbol{\Sigma}_p) - \frac{d}{2} + \frac{1}{2} \log \frac{\det \boldsymbol{\Sigma}_q}{\det \boldsymbol{\Sigma}_p}}_{\text{Dispersion}} + \underbrace{\frac{1}{2} (\boldsymbol{\mu}_p - \boldsymbol{\mu}_q)^\top \boldsymbol{\Sigma}_q^{-1} (\boldsymbol{\mu}_p - \boldsymbol{\mu}_q)}_{\text{Signal}}.$$

7.3 Uncertainty in Dynamical Systems

Consider $\dot{\mathbf{x}} = \mathbf{f}(\mathbf{x})$ with uncertain initial condition $\mathbf{x}(0) \sim p_0$.

Linear systems: $\dot{\mathbf{x}} = \mathbf{A}\mathbf{x}$. If $\mathbf{A}$ is stable, uncertainty *decays* and mean dynamics separate: [?]

$$\frac{d}{dt} \mathbb{E}[\mathbf{x}] = \mathbf{A} \mathbb{E}[\mathbf{x}], \quad \frac{d}{dt} \text{Cov}[\mathbf{x}] = \mathbf{A} \text{Cov}[\mathbf{x}] + \text{Cov}[\mathbf{x}] \mathbf{A}^\top.$$

Nonlinear systems: Reynolds decomposition $\mathbf{x} = \bar{\mathbf{x}} + \mathbf{x}'$:

$$\frac{d}{dt} \bar{\mathbf{x}} = \mathbb{E}[\mathbf{f}(\bar{\mathbf{x}} + \mathbf{x}')] \neq \mathbf{f}(\bar{\mathbf{x}}).$$

The *mean dynamics couple to higher moments*. For quadratic $\mathbf{f}$:

$$\mathbb{E}[\mathbf{x}\mathbf{x}^\top] = \bar{\mathbf{x}}\bar{\mathbf{x}}^\top + \text{Cov}[\mathbf{x}].$$

The covariance feeds back into the mean! This is the **closure problem**.

Chaotic systems (Lorenz-63): Small initial uncertainty $\rightarrow$ exponential divergence (positive Lyapunov exponents). Ensemble mean quickly becomes meaningless, [?].

7.4 Data Assimilation: Kalman and Ensemble Filters

[?]

Combine model forecast $\mathbf{x}^f \sim \mathcal{N}(\boldsymbol{\mu}^f, \mathbf{P}^f)$ with observation $\mathbf{y} = \mathbf{H}\mathbf{x} + \boldsymbol{\eta}$, $\boldsymbol{\eta} \sim \mathcal{N}(0, \mathbf{R})$.

Theorem 7.2 (Kalman Filter Update). [?] *Posterior (analysis) distribution:*

$$\boldsymbol{\mu}^a = \boldsymbol{\mu}^f + \mathbf{K}(\mathbf{y} - \mathbf{H}\boldsymbol{\mu}^f), \quad \mathbf{P}^a = (\mathbf{I} - \mathbf{K}\mathbf{H})\mathbf{P}^f,$$

where $\mathbf{K} = \mathbf{P}^f \mathbf{H}^\top (\mathbf{H} \mathbf{P}^f \mathbf{H}^\top + \mathbf{R})^{-1}$ is the **Kalman gain**, [?].

For nonlinear/high-dimensional systems, use **Ensemble Kalman Filter (EnKF)**: [?]

1. Maintain ensemble $\{\mathbf{x}_i^f\}_{i=1}^N$ representing forecast distribution
2. Sample observations $\mathbf{y}_i = \mathbf{y} + \boldsymbol{\eta}_i$, $\boldsymbol{\eta}_i \sim \mathcal{N}(0, \mathbf{R})$
3. Update each member: $\mathbf{x}_i^a = \mathbf{x}_i^f + \mathbf{K}(\mathbf{y}_i - \mathbf{H}\mathbf{x}_i^f)$
4. $\mathbf{K}$ estimated from ensemble covariance

Listing 7.1: Ensemble Kalman Filter

```

1  using LinearAlgebra, Statistics, Random
2
3  function enkf_update(ensemble::Matrix, H, y, R)
4      # ensemble: n_state x n_ensemble
5      # H: n_obs x n_state
6      # y: n_obs observation
7      # R: n_obs x n_obs observation noise covariance
8      n_state, N = size(ensemble)
9
10     # Ensemble mean and perturbation matrix
11     mu_f = mean(ensemble, dims=2)
12     X_f = ensemble .- mu_f # n_state x N
13
14     # Forecast ensemble in observation space
15     Y_f = H * ensemble # n_obs x N
16     y_mean = mean(Y_f, dims=2)
17     Y_pert = Y_f .- y_mean
18
19     # Observation perturbations
20     y_obs = y .+ sqrt(R) * randn(n_obs, N)
21
22     # Ensemble covariances
23     Pf_xy = (X_f * Y_pert') / (N - 1) # n_state x n_obs
24     Pf_yy = (Y_pert * Y_pert') / (N - 1) # n_obs x n_obs
25
26     # Kalman gain \cite{chen2025uq}
27     K = Pf_xy * inv(Pf_yy + R)
28
29     # Analysis update
30     ensemble_a = ensemble + K * (y_obs - Y_f)
31     return ensemble_a
32 end
33
34 # Lorenz-63 twin experiment
35 function lorenz63(u, p, t)
36     sigma, rho, beta = 10.0, 28.0, 8/3
37     x, y, z = u
38     [sigma*(y-x), x*(rho-z)-y, x*y - beta*z]
39 end
40
41 # Run EnKF on Lorenz-63
42 using DifferentialEquations
43 prob = ODEProblem(lorenz63, [1.0, 1.0, 1.0], (0.0, 10.0))

```

```

44 true_sol = solve(prob, Tsit5(), saveat=0.01)
45
46 # Generate synthetic observations (every 0.1 time units, observe x only)
47 obs_times = 0.0:0.1:10.0
48 obs = true_sol(obs_times)[1, :] + 0.5 * randn(length(obs_times))
49
50 # EnKF assimilation
51 ensemble = [1.0, 1.0, 1.0] .+ 0.5 * randn(3, 50) # 50 members
52 H = [1.0 0.0 0.0]
53 R = [0.25]
54
55 for (i, t) in enumerate(obs_times)
56     # Forecast step
57     prob = ODEProblem(lorenz63, ensemble, (t, t+0.1))
58     ensemble = hcat([solve(prob, Tsit5(), saveat=0.1).u[end] for _ in
59                     1:50]...)
60
61     # Analysis step
62     ensemble = enfk_update(ensemble, H, obs[i], R)
63 end

```

7.5 Stochastic Surrogate Modeling

High-dimensional PDEs (climate, turbulence) are too expensive for ensemble forecasting. Replace with a stochastic model that matches statistics.

Consider a Fourier mode $\hat{u}_k(t)$ from a turbulent PDE:

$$\frac{d\hat{u}_k}{dt} = \lambda_k \hat{u}_k + \sum_{p+q=k} N_{k,p,q} \hat{u}_p \hat{u}_q + F_k.$$

Stochastic surrogate: Replace nonlinear sum with noise calibrated to match statistics:

$$\frac{d\hat{u}_k}{dt} = \lambda_k \hat{u}_k + \sigma_k \dot{W}_k(t),$$

where σ_k chosen so equilibrium PDF and decorrelation time match the true mode.

Theorem 7.3 (Calibration). *For Ornstein-Uhlenbeck $dx = -\theta x dt + \sigma dW$:*

- *Equilibrium variance:* $\sigma^2/(2\theta) = \text{Var}_{\text{true}}$
- *Decorrelation time:* $1/\theta = \tau_{\text{true}}$

Solve for θ, σ .

Example 7.4 (2-Layer Quasi-Geostrophic Model). *Full model: 128^2 modes per layer. Stochastic surrogate: only 5 leading modes, calibrated to their statistics. Qualitatively similar spatial fields, matching temporal correlations and PDFs.*

7.6 Bridge to Chapter 8

UQ quantifies uncertainty in *dynamical systems*. Chapter 8 applies algebraic geometry to a *static* but nonlinear problem: the global positioning equations. Both use the same insight: *reformulate the problem to reveal its mathematical structure* (sphere intersection for GPS, entropy/KL for UQ), [?].

7.7 Further Reading

- Chen, Wiggins, Andreou (2025). *Taming Uncertainty...* Notices, 72(3). Code: https://github.com/marandmath/UQ_tutorial_code
- Chen (2023). *Stochastic Methods for Modeling and Predicting Complex Dynamical Systems*. Springer.
- Law, Stuart, Zygalakis (2015). *Data Assimilation*. Springer.
- Reich, Cotter (2015). *Probabilistic Forecasting and Bayesian Data Assimilation*. Cambridge.
- Majda, Harlim (2012). *Filtering Complex Turbulent Systems*. Science.

7.8 Exercises

7.9 Solutions

Solution.Exercise 1 (Entropy calculation): For a Gaussian distribution $\mathcal{N}(\mu, \Sigma)$, the Shannon entropy is $H = \frac{1}{2} \log((2\pi e)^n |\Sigma|)$. For $\Sigma = I$, $H = \frac{n}{2}(1 + \log(2\pi))$.

Solution.Exercise 2 (KL divergence): For two Gaussians $P = \mathcal{N}(\mu_1, \Sigma_1)$ and $Q = \mathcal{N}(\mu_2, \Sigma_2)$, the KL divergence is:

$$D_{KL}(P\|Q) = \frac{1}{2} \left[\text{tr}(\Sigma_2^{-1} \Sigma_1) + (\mu_2 - \mu_1)^T \Sigma_2^{-1} (\mu_2 - \mu_1) - n + \log \frac{|\Sigma_2|}{|\Sigma_1|} \right]$$

Solution.Exercise 3 (Ensemble Kalman): For an ensemble of N particles, the Kalman update is $x_i^a = x_i^f + K(y - Hx_i^f)$ where $K = P^f H^T (H P^f H^T + R)^{-1}$. The ensemble covariance P^f is estimated from the sample covariance of the ensemble.

Solution.Exercise 4 (Research): For turbulent PDEs, stochastic Galerkin methods project the solution onto a polynomial chaos basis. The convergence rate depends on the smoothness of the solution with respect to the random parameters.

Exercise 7.1. Compute the KL divergence between $\mathcal{N}(0, 1)$ and $\mathcal{N}(0.5, 1.5)$. Interpret the signal and dispersion terms.

Exercise 7.2. Implement the EnKF for the Lorenz-96 system (40 variables, observe every 4th). Compare RMSE vs. ensemble size $N = 10, 20, 50, 100$.

Exercise 7.3. For a scalar OU process with true variance 2.0 and decorrelation time 0.5, compute the surrogate parameters θ, σ . Simulate and verify equilibrium statistics.

Exercise 7.4 (Research). Apply the stochastic surrogate approach to the Kuramoto-Sivashinsky equation. Can you reduce the number of modes by 90

CHAPTER 8

New Developments in Global Positioning

Mireille Boutin, Gregor Kemper

Communicated by *Notices* Associate Editor

“This is a story about how mathematics can make a difference in a real-world problem... The fundamental question of when the problem has a unique solution and when it does not was still unanswered.”

— Boutin & Kemper

Abstract. The Global Positioning System (GPS) solves a system of quadratic equations every time your phone finds its location. This chapter shows how algebraic geometry and Minkowski geometry reveal the exact conditions for unique solvability: the satellite positions must not lie on a quadric with the user position as focus. The solution is a linear algebra problem in disguise, [?].

8.1 Introduction: The Problem

Mathematical notation follows standard conventions. Key terms are defined in the glossary.

Your phone receives signals from GPS satellites. Each satellite i broadcasts: [?]

- Its position $\mathbf{s}_i \in \mathbb{R}^3$ (known precisely)
- Transmission time t_i^{tx} (from atomic clock)

Your phone receives at time t_i^{rx} (from its cheap quartz clock). The *apparent flight time* is $\tau_i = t_i^{\text{rx}} - t_i^{\text{tx}}$.

The **global positioning equations**:

$$\|\mathbf{x} - \mathbf{s}_i\| = c\tau_i - b, \quad i = 1, \dots, m$$

where:

- $\mathbf{x} \in \mathbb{R}^3$ = your position (unknown)
- b = your clock bias (unknown)
- c = speed of light (set $c = 1$ by time units)
- m = number of satellites in view (≥ 4) [?]

This is a system of m quadratic equations in 4 unknowns.

8.2 Sphere Intersection in Minkowski Space

Square both sides:

$$\|\mathbf{x} - \mathbf{s}_i\|^2 = (\tau_i - b)^2.$$

Expand:

$$\|\mathbf{x}\|^2 - 2\mathbf{s}_i^\top \mathbf{x} + \|\mathbf{s}_i\|^2 = \tau_i^2 - 2\tau_i b + b^2.$$

Introduce the Minkowski inner product on $\mathbb{R}^4$:

$$\langle (\mathbf{x}, b), (\mathbf{y}, c) \rangle_M = \mathbf{x}^\top \mathbf{y} - bc.$$

Define $\mathbf{X} = (\mathbf{x}, b)$, $\mathbf{S}_i = (\mathbf{s}_i, \tau_i)$, $\sigma_i = \frac{1}{2}(\|\mathbf{s}_i\|^2 - \tau_i^2)$. Then:

$$\|\mathbf{X}\|_M^2 - 2\langle \mathbf{X}, \mathbf{S}_i \rangle_M = -2\sigma_i,$$

where $\|\cdot\|_M^2 = \langle \cdot, \cdot \rangle_M$.

This says: $\mathbf{X}$ lies on a **Minkowski sphere** (quadric) centered at $\mathbf{S}_i$ with “radius” determined by σ_i .

Theorem 8.1 (Sphere Intersection as Linear Algebra). *The system is equivalent to:*

$$\mathbf{A}\mathbf{X} = \boldsymbol{\sigma},$$

where $\mathbf{A}$ has rows $(2\mathbf{s}_i^\top, -2\tau_i, -1)$ and $\sigma_i = \|\mathbf{s}_i\|^2 - \tau_i^2$.

If $\mathbf{A}$ has full column rank (4), then $\mathbf{X}$ is uniquely determined by the linear system. Then check $\|\mathbf{X}\|_M^2 = \mathbf{X}^\top \boldsymbol{\eta} \mathbf{X}$ consistency.

8.3 When Is the Solution Unique?

Theorem 8.2 (Boutin-Kemper 2024). *Assume $m \geq 4$ satellites with non-coplanar positions. The GPS problem has: [?]*

- A **unique** solution iff the satellites do not lie on a quadric surface with focus at the user position, [?].
- **Two** solutions iff they do lie on such a quadric.

The quadric is a sphere, ellipsoid, paraboloid, or hyperboloid (depending on eccentricity).

Sketch. If $\mathbf{A}$ has rank 3, the linear system gives a line of solutions. Substituting into the quadratic constraint $\|\mathbf{X}\|_M^2 = \dots$ gives a quadratic equation: 0, 1, or 2 solutions. Rank deficiency of $\mathbf{A}$ occurs exactly when satellites lie on a quadric with focus $\mathbf{X}$, [?]. $\square$

Example 8.3 (Practical geometry). • **Good:** Satellites well-spread across sky (rank 4)

- **Bad:** All satellites in a plane (e.g., urban canyon) $\rightarrow$ rank 3 $\rightarrow$ ambiguity [?]
- **Bad:** Satellites on a sphere centered at you $\rightarrow$ quadric with focus at you $\rightarrow$ two solutions

8.4 Algorithm: Exact Solution

Algorithm 2 GPS Solution (Exact Data)

```

1: procedure GPS-SOLVE( $\mathbf{s}_i, \tau_i$  for  $i = 1..m$ ) [?]
2:   Form  $\mathbf{A} \in \mathbb{R}^{m \times 4}$  with rows  $(2\mathbf{s}_i^\top, -2\tau_i, -1)$ 
3:   Form  $\boldsymbol{\sigma} \in \mathbb{R}^m$  with  $\sigma_i = \|\mathbf{s}_i\|^2 - \tau_i^2$ 
4:   if rank( $\mathbf{A}$ ) = 4 then
5:      $\mathbf{X} \leftarrow (\mathbf{A}^\top \mathbf{A})^{-1} \mathbf{A}^\top \boldsymbol{\sigma}$  ▷ Least squares if  $m > 4$ 
6:     return  $\mathbf{x} = \mathbf{X}_{1:3}, b = \mathbf{X}_4$ 
7:   else ▷ rank 3
8:     Find line  $\mathbf{X} = \mathbf{X}_0 + t\mathbf{v}$  solving  $\mathbf{A}\mathbf{X} = \boldsymbol{\sigma}$ 
9:     Substitute into  $\|\mathbf{X}\|_M^2 - 2\langle \mathbf{X}, \mathbf{S}_i \rangle_M + 2\sigma_i = 0$ 
10:    Solve quadratic for  $t$  (0, 1, or 2 solutions)
11:    return all valid solutions
12:  end if
13: end procedure
    
```

8.5 Noisy Data: Weighted Least Squares

Real measurements have noise. The linearized system $\mathbf{A}\mathbf{X} = \boldsymbol{\sigma}$ has errors from:

- Satellite position errors (small, $\sim$ cm)
- Clock/timing errors (larger, $\sim$ meters)
- Atmospheric delays (ionosphere/troposphere)

Satellite positions are known much more accurately than apparent flight times. So invert the submatrix $\mathbf{A}_1$ obtained by deleting the *first* column (the $\mathbf{s}_i$ terms), rather than the last column (Bancroft's method), [?].

Listing 8.1: GPS solver with noise

```

1  using LinearAlgebra, Statistics
2
3  function solve_gps(sat_positions::Matrix, tau::Vector; weights=nothing)
4      # sat_positions: 3 x m matrix of satellite positions \cite{
5      #   tau: m vector of apparent flight times
6      m = length(tau)
7
8      # Build A matrix: rows are [2*s_i', -2*tau_i, -1]
9      A = [2*sat_positions'  -2*tau  -ones(m)]
10
11     # sigma = ||s_i||^2 - tau_i^2
12     sigma = sum(sat_positions.^2, dims=1)' - tau.^2
13
14     if rank(A) == 4
15         if weights == nothing
    
```

```

16         X = (A' * A) \ (A' * sigma)
17     else
18         W = Diagonal(weights)
19         X = (A' * W * A) \ (A' * W * sigma)
20     end
21     return X[1:3], X[4]
22 else
23     # Rank 3: find line of solutions
24     # A has nullspace of dimension 1
25     F = svd(A)
26     v_null = F.V[:, end]
27
28     # Particular solution
29     X0 = A \ sigma
30
31     # Quadratic constraint: ||X||_M^2 - 2*<X, S_i>_M + 2*sigma_i = 0
32     # This gives a quadratic in t: X = X0 + t * v_null
33     # ... (details omitted for brevity)
34     # Return both solutions if they exist
35     error("Rank-deficient case: implement quadratic solve")
36 end
37 end
38
39 # Test with simulated data
40 Random.seed!(42)
41 true_x = [1000.0, 2000.0, 500.0] # user position
42 true_b = 50.0 # clock bias
43
44 # 6 satellites \cite{boutin2025gps}
45 sat = randn(3, 6) * 20000 .+ [0, 0, 20000] # roughly in orbit
46 tau = [norm(sat[:,i] - true_x) + true_b for i in 1:6]
47
48 # Add noise
49 tau_noisy = tau + 0.01 * randn(6)
50
51 x_est, b_est = solve_gps(sat, tau_noisy)
52 println("True: ", true_x, " ", true_b)
53 println("Est: ", x_est, " ", b_est)
54 println("Error: ", norm(x_est - true_x))

```

8.6 Conditioning and Geometry

The condition number of the linear system $\mathbf{A}\mathbf{X} = \boldsymbol{\sigma}$ blows up as satellites approach a quadric with focus at the user. This is *not* just a numerical issue—it reflects a fundamental geometric ambiguity, [?].

Example 8.4 (Urban canyon). *Satellites visible only along a street: positions nearly coplanar. The matrix $\mathbf{A}$ has rank 3. Two possible positions symmetric across the satellite plane. The linear system is ill-conditioned; small timing errors $\rightarrow$ large position errors, [?].*

8.7 Bridge to Chapter 9

GPS solves a *static* nonlinear system by algebraic geometry. Chapter 9 addresses a *dynamic* inverse problem: DNA storage systems where the channel introduces insertions, deletions, and substitutions—requiring coding theory for edit-distance channels, [?].

8.8 Further Reading

- Boutin, Kemper (2025). *New Developments in Global Positioning*. Notices, 72(8).
- Bancroft (1985). *An Algebraic Solution of the GPS Equations*. IEEE Trans. Aerosp. Electron. Syst, [?], [?].
- Strang, Borre (1997). *Linear Algebra, Geodesy, and GPS*. Wellesley-Cambridge, [?], [?].
- Hofmann-Wellenhof et al. (2001). *GPS: Theory and Practice*. Springer, [?].

8.9 Exercises

8.10 Solutions

Solution.Exercise 1 (GPS with 4 satellites): For satellites at $(\pm 20000, 0, 20000)$ and $(0, \pm 20000, 20000)$ with $\tau_i = 25000$, the user is at the origin with clock bias $b = 50$. The satellites lie on a sphere centered at the origin, which is a quadric with focus at the user. This gives two symmetric solutions: $(0, 0, 0)$ and a point far below the satellites.

Solution.Exercise 2 (Sphere ambiguity): If all satellites lie on a sphere centered at the user, the GPS equations have exactly two solutions. This is because the quadratic constraint $\|X\|^2 - 2\sigma \cdot X + \|\sigma\|^2 = 0$ intersects the plane of linear solutions in two points.

Solution.Exercise 3 (Rank-3 solver): For coplanar satellites, the matrix A has rank 3. The nullspace is spanned by the normal to the satellite plane. Substituting $X = X_0 + tv$ into the quadratic constraint gives a quadratic equation in t , yielding 0, 1, or 2 solutions.

Solution.Exercise 4 (Research): The Bancroft algorithm provides a closed-form solution for 4 satellites. It solves the quartic equation arising from the quadratic constraint, giving up to 4 algebraic solutions. The physically meaningful solution is selected by requiring positive altitude and clock bias.

Exercise 8.1. For 4 satellites at $(20000, 0, 20000)$, $(0, 20000, 20000)$, $(-20000, 0, 20000)$, $(0, -20000, 20000)$ with $\tau_i = 25000$, compute the user position and clock bias. Are the satellites on a quadric with focus at the solution?

Exercise 8.2. Show that if all satellites lie on a sphere centered at the user position, the GPS equations have exactly two solutions, [?].

Exercise 8.3. Implement the rank-3 quadratic solver for the ambiguous case. Test on a coplanar satellite configuration, [?].

Exercise 8.4 (Research). Analyze the effect of ionospheric delay (frequency-dependent) on the Minkowski geometry formulation. How does dual-frequency GPS resolve this? [?]

CHAPTER 9

3D Printing of Invariant Manifolds in Dynamical Systems

*Patrick R. Bishop, Summer Chenoweth, Emmanuel Fleurantin,
Alonso Ogueda-Oliva, Evelyn Sander, Julia Seay*
Communicated by *Notices* Associate Editor Vidit Nanda

“Invariant manifolds play a crucial role in understanding the behavior of dynamical systems... Visualizing and intuitively grasping their complex geometries remains challenging. This article introduces a novel approach: 3D printing tangible, physical representations of invariant manifolds.”

— Bishop et al.

[?]

Abstract. Stable and unstable manifolds organize the long-term behavior of dynamical systems. They are typically visualized as 2D projections of 3D objects—losing crucial geometric information. This chapter presents a complete pipeline: compute local manifolds via the Parameterization Method, extend globally via numerical integration, mesh and thicken for 3D printing, and produce physical models that reveal intersections, spirals, and folding invisible on screens, [?].

9.1 Introduction: Why Print Manifolds?

Mathematical notation follows standard conventions. Key terms are defined in the glossary.

Consider the Lorenz system: [?]

$$\dot{x} = \sigma(y - x), \quad \dot{y} = x(\rho - z) - y, \quad \dot{z} = xy - \beta z.$$

The origin is a saddle with a 2D stable manifold and 1D unstable manifold. On screen, these appear as overlapping curves. In your hand, the stable manifold is a *twisted sheet* that folds back on itself;

the unstable manifold is a *spiral ribbon*. The physical object reveals the *heteroclinic tangle*—the engine of chaos, [?].

This chapter gives you the complete computational pipeline to:

1. Compute local invariant manifolds via the **Parameterization Method** (high-order Taylor approximation), [?].
2. Extend to **global manifolds** via adaptive arclength integration, [?].
3. Generate **watertight meshes** with proper thickening.
4. Prepare for **FDM/SLA printing** with supports and orientation.

9.2 Background: Invariant Manifolds

Definition 9.1 (Stable/Unstable Manifolds). *For $\dot{\mathbf{x}} = \mathbf{f}(\mathbf{x})$ with hyperbolic equilibrium $\mathbf{x}^*$, let $\mathbf{A} = D\mathbf{f}(\mathbf{x}^*)$ have eigenvalues $\lambda_1, \dots, \lambda_n$ with $\text{Re}(\lambda_i) \neq 0$.*

- $W^s(\mathbf{x}^*) = \{\mathbf{x} : \varphi_t(\mathbf{x}) \rightarrow \mathbf{x}^* \text{ as } t \rightarrow +\infty\}$ (stable manifold) [?]
- $W^u(\mathbf{x}^*) = \{\mathbf{x} : \varphi_t(\mathbf{x}) \rightarrow \mathbf{x}^* \text{ as } t \rightarrow -\infty\}$ (unstable manifold) [?]

Dimensions: $\dim W^s = \#\{\text{Re}(\lambda) < 0\}$, $\dim W^u = \#\{\text{Re}(\lambda) > 0\}$.

Theorem 9.2 (Stable Manifold Theorem). *W^s and W^u are smooth immersed submanifolds tangent to the stable/unstable eigenspaces of $\mathbf{A}$ at $\mathbf{x}^*$, [?].*

9.3 Phase 1: Local Manifolds via the Parameterization Method

The Parameterization Method (Cabr , Fontich, de la Llave) finds a map $P : U \subset \mathbb{R}^k \rightarrow \mathbb{R}^n$ such that:

1. $P(0) = \mathbf{x}^*$ (equilibrium at origin of parameter domain)
2. $DP(0)$ spans the target eigenspace
3. Invariance: $DP(\boldsymbol{\theta})\boldsymbol{\Lambda}\boldsymbol{\theta} = \mathbf{f}(P(\boldsymbol{\theta}))$ for $\boldsymbol{\theta} \in U$

where $\boldsymbol{\Lambda} = \text{diag}(\lambda_1, \dots, \lambda_k)$ are the target eigenvalues.

Expanding P as a power series $P(\boldsymbol{\theta}) = \sum_{|\boldsymbol{\alpha}| \geq 1} P_{\boldsymbol{\alpha}} \boldsymbol{\theta}^{\boldsymbol{\alpha}}$ and matching coefficients gives linear equations for $P_{\boldsymbol{\alpha}}$:

$$(\boldsymbol{\Lambda}\boldsymbol{\alpha} \cdot \mathbf{I} - \mathbf{A})P_{\boldsymbol{\alpha}} = \mathbf{R}_{\boldsymbol{\alpha}},$$

where $\mathbf{R}_{\boldsymbol{\alpha}}$ depends on previously computed coefficients. Non-resonance condition: $\boldsymbol{\Lambda}\boldsymbol{\alpha} \neq \lambda_j$ for all j .

Algorithm 3 Parameterization Method (local stable manifold)

[?]

- 1: Compute equilibrium $\mathbf{x}^*$, Jacobian $\mathbf{A} = D\mathbf{f}(\mathbf{x}^*)$
 - 2: Select k stable eigenvalues $\lambda_1, \dots, \lambda_k$ with eigenvectors $\mathbf{v}_1, \dots, \mathbf{v}_k$
 - 3: Set $P_{\mathbf{e}_i} = \mathbf{v}_i$ for $i = 1, \dots, k$ (linear terms)
 - 4: **for** $m = 2$ to N (order) **do**
 - 5: **for** all multi-indices α with $|\alpha| = m$ **do**
 - 6: Compute $\mathbf{R}_\alpha$ from known P_β ($|\beta| < m$)
 - 7: Solve $(\sum \alpha_i \lambda_i \mathbf{I} - \mathbf{A})P_\alpha = \mathbf{R}_\alpha$
 - 8: **end for**
 - 9: **end for**
 - 10: Return polynomial $P(\theta)$
-

9.4 Phase 2: Global Extension

The local parameterization P is accurate only near $\mathbf{x}^*$ (radius $\sim$ convergence of series). To reach global structure:

1. Sample the **boundary** of the local manifold: $\partial D = \{P(\theta) : \|\theta\| = R\}$, [?].
2. For each boundary point, integrate $\dot{\mathbf{x}} = \mathbf{f}(\mathbf{x})$ forward (for W^s) or backward (for W^u) using adaptive Runge-Kutta.
3. Track the **arclength** to ensure uniform sampling.
4. Use **shadowing** or **manifold continuation** (e.g., MATCONT, COCO) for bifurcations, [?].

Listing 9.1: Global manifold extension

```

1 using DifferentialEquations, LinearAlgebra
2
3 function extend_manifold(P, theta_boundary, f, tspan; dt=0.01) \cite{
4     cabre2003parameterization}
5     # theta_boundary: matrix of boundary parameter values (k x M)
6     # Returns: vector of trajectory points on global manifold \cite{
7         cabre2003parameterization}
8     M = size(theta_boundary, 2)
9     trajectories = Vector{Vector{Vector{Float64}}}(undef, M)
10
11     for i = 1:M
12         u0 = P(theta_boundary[:, i])
13         prob = ODEProblem(f, u0, tspan)
14         sol = solve(prob, Tsit5(), saveat=dt, abstol=1e-9, reltol=1e-9)
15         trajectories[i] = [sol.u[j] for j in 1:length(sol.u)]
16     end
17     return trajectories
18 end

```


9.5 Phase 3: Meshing and Thickening

Raw trajectories are 1D (unstable) or 2D point clouds (stable). For 3D printing, we need a **watertight closed mesh**.

1. **Delaunay triangulation** of point cloud (for 2D manifolds), [?].
2. **Poisson surface reconstruction** (Kazhdan et al.) for noisy/unstructured data.
3. **Thickening**: Offset the surface along normal by $\pm\epsilon/2$ to create a solid of thickness ϵ (typically 1–2 mm for FDM).
4. **Manifold repair**: Remove non-manifold edges, fill holes, ensure consistent normals, [?].

Listing 9.2: Mesh generation and thickening

```

0 using Meshing, GeometryBasics, FileIO
1
2 function points_to_mesh(points; thickness=1.5)
3     # points: N x 3 matrix of points on the manifold \cite{
4     # Returns: watertight mesh (vertices, faces)
5
6     # Poisson surface reconstruction (requires normal estimates)
7     normals = estimate_normals(points, k=20)
8     mesh = poisson_reconstruct(points, normals; depth=8)
9
10    # Thicken: create offset surfaces
11    inner = offset_mesh(mesh, -thickness/2)
12    outer = offset_mesh(mesh, +thickness/2)
13
14    # Combine into closed solid
15    solid = union_meshes(inner, outer)
16    repair_mesh!(solid)
17    return solid
18 end
19
20 function estimate_normals(points, k)
21     # PCA on k-nearest neighbors
22     tree = KDTree(points')
23     normals = zeros(size(points))
24     for i in 1:size(points, 1)
25         idxs, _ = knn(tree, points[i, :], k+1)
26         idxs = idxs[2:end] # exclude self
27         cov = cov(points[idxs, :])
28         _, eigvecs = eigen(cov)
29         normals[i, :] = eigvecs[:, 1] # smallest eigenvalue direction
30     end
31     return normals
32 end

```

9.6 Phase 4: Print Preparation

- **Orientation:** Print with the flattest face down; minimize overhangs $> 45^\circ$.
- **Supports:** Tree supports for delicate spirals; grid supports for flat sheets.
- **Material:** PLA (easy), PETG (tougher), Resin (SLA, highest detail).
- **Scale:** Lorenz stable manifold fits in $10 \times 10 \times 10$ cm at scale 1:1, [?], [?].
- **Infill:** 15-20% gyroid for structural stability with minimal weight.

Example 9.3 (Lorenz system parameters). [?] $\sigma = 10$, $\rho = 28$, $\beta = 8/3$. *Equilibrium at origin: eigenvalues $\approx -22.8, 11.8, -1.8$.*

- W^s : 2D, spanned by λ_1, λ_3 . Compute order-10 parameterization.
- W^u : 1D, integrate backward from local unstable eigenvector.
- *Result: Two spiraling ribbons intersecting in a heteroclinic tangle.*

9.7 Code Repository

Complete pipeline (MATLAB/Julia/Python) at:

<https://github.com/msander/invariant-manifold-3d-print> [?]

Includes:

- `ParameterizationMethod.m` — high-order Taylor coefficients
- `GlobalExtension.jl` — adaptive integration
- `MeshAndThicken.py` — Poisson reconstruction + offset
- `STLExport` — print-ready files for Lorenz, Arneodo-Coullet-Tresser, Langford systems [?]

9.8 Bridge to Chapter 10

We’ve seen how to *visualize* and *fabricate* geometric objects from dynamical systems. Chapter 10 addresses a different challenge: how to *compose* complex systems from simpler components using the mathematics of operads—category theory for system design.

9.9 Further Reading

- Bishop et al. (2026). *3D Printing of Invariant Manifolds*. Notices, 73(1).
- Cabré, Fontich, de la Llave (2003). *The Parameterization Method*. J. Diff. Eq.
- Krauskopf, Osinga (2005). *Growing 1D and Quasi-2D Unstable Manifolds*. Nonlinearity, [?].
- MATLAB package `MATCONT`, Julia package `BifurcationKit.jl`.

9.10 Exercises

9.11 Solutions

Solution.Exercise 1 (Lorenz stable manifold): The Lorenz system at $\sigma = 10, \rho = 28, \beta = 8/3$ has a saddle fixed point at the origin. The stable manifold is 2D, tangent to the eigenvectors of the stable eigenvalues. At order 5 Taylor approximation, the manifold is:

$$W_{\text{loc}}^s = \{(x, y, z) : z = h(x, y)\}$$

where h is computed by solving the homological equation.

Solution.Exercise 2 (ACT system): The Arneodo-Coullet-Tresser system has a homoclinic orbit to the origin for certain parameter values. The unstable manifold is computed by the Parameterization Method using Taylor series expansion.

Solution.Exercise 3 (Homoclinic tangle): The Smale horseshoe has a homoclinic tangle where the stable and unstable manifolds intersect transversally. For 3D printing, we approximate the manifold by a mesh of triangles, truncating at a minimum scale.

Solution.Exercise 4 (Research): The Parameterization Method can be extended to higher dimensions and to invariant tori. The convergence is guaranteed by the implicit function theorem in Banach spaces.

Exercise 9.1. Implement the Parameterization Method for the 2D stable manifold of the Lorenz origin at order 5. Compare the Taylor polynomial error vs. distance from origin (Bishop et al., 2026; Cabre et al., 2023).

Exercise 9.2. For the Arneodo-Coullet-Tresser system $(\dot{x} = y, \dot{y} = z, \dot{z} = -ax - by - cz + dx^3)$, compute and print the unstable manifold of the origin, [?].

Exercise 9.3. Design a 3D-printed model showing the *homoclinic tangle* of the Smale horseshoe. How do you represent the infinite folding in a finite print?

Part IV

Modern Applications

CHAPTER 10

Exciting Coding Problems for DNA-based Storage Systems

Daniella Bar-Lev, Omer Sabary, Eitan Yaakobi

Communicated by *Notices* Associate Editor Richard Levine

“DNA storage generates new exciting and challenging theoretical problems to our community. The success of such solutions is crucial due to the never-ending growth of data and the foreseeable inability to store it.”

— Bar-Lev, Sabary, Yaakobi

[?]

Abstract. DNA storage offers ultra-dense, durable data archiving but introduces a unique error model: insertions, deletions, and substitutions (edit errors), unordered strands with massive redundancy, and constrained sequences (GC-content, homopolymers). This chapter surveys the coding theory challenges: edit-distance codes, clustering algorithms, sequence reconstruction, and constrained codes, [?].

10.1 Introduction: Why DNA Storage?

Mathematical notation follows standard conventions. Key terms are defined in the glossary.

DNA offers remarkable storage properties:

- **Density:** $\sim 10^{18}$ bytes/mm³ (exabyte per cubic millimeter)
- **Durability:** Thousands of years if kept cool/dry
- **Energy:** No power to maintain; only to read/write
- **Timelessness:** DNA sequencing technology will never become obsolete

But the **channel model** is unlike any classical storage:

Errors Insertions, deletions, substitutions (edit errors), not just bit flips

Ordering Strands are *unordered* (no address track)

Redundancy Millions of copies per strand (unique to DNA)

Constraints GC-content (45-55%), no homopolymers > 3

10.2 Edit-Distance Codes

Classical codes (RS, LDPC, BCH) correct *substitutions/erasures*. Edit errors shift the entire subsequent sequence.

Definition 10.1 (Levenshtein Distance). $d_L(x, y)$ = *minimum number of insertions, deletions, substitutions to transform x into y .*

Theorem 10.2 (Levenshtein 1965). *A code corrects k deletions iff it corrects any combination of k insertions/deletions.*

Example 10.3 (Varshamov-Tenengolts (VT) Codes). *For $a \in \{0, \dots, n\}$, the length- n VT code is*

$$VT_a(n) = \left\{ x \in \{0, 1\}^n : \sum_{i=1}^n ix_i \equiv a \pmod{n+1} \right\}.$$

$VT_0(n)$ is a perfect single-deletion-correcting code. For any a , $|VT_a(n)| \approx 2^n/(n+1)$.

Theorem 10.4 (VT Code Optimality). *It is open whether $VT_0(n)$ is strictly optimal for all n . Conjectured yes (Sloane, 2002).*

10.3 The DNA Clustering Problem

Each synthesized strand has millions of copies. After sequencing, we get an unordered multiset of noisy reads. Before decoding, we must *cluster* reads by their original strand, [?].

Problem 10.5 (DNA Clustering). *Given N reads of length L , partition into clusters $C_1, \dots, C_M$ such that reads in C_j originate from the same synthesized strand.*

Approaches:

1. **Edit-distance clustering:** Hierarchical clustering with d_L (slow, $O(N^2)$)
2. **Index-based:** Encode strand ID in the sequence (reduces capacity)
3. **Locality-sensitive hashing:** Approximate edit distance for scalability [?]

Current methods fail at high error rates ($> 5\%$) and large N ($> 10^6$).

10.4 Sequence Reconstruction from Noisy Copies

After clustering, each cluster has k noisy copies of one strand. Goal: reconstruct the original.

Problem 10.6 (Trace Reconstruction). *Given k independent traces of $x \in \{0, 1\}^n$ through a deletion channel (each bit deleted with prob. p), reconstruct x .*

Theorem 10.7 (Batu et al. 2004). *For deletion probability p , $k = \exp(O(\log^{1/3} n))$ traces suffice for reconstruction with high probability.*

DNA adds *substitutions and insertions*. The minimum cluster size for unique reconstruction is unknown for edit channels.

10.5 Constrained Codes

Synthesis/sequencing errors depend on sequence content. Two key constraints:

Run-length No more than 3 identical consecutive bases (e.g., AAAA forbidden)

GC-content Fraction of G/C in $[0.45, 0.55]$

Theorem 10.8 (Capacity of Constrained Codes). *For a constrained system $S \subset \{A, C, G, T\}^*$, the capacity is*

$$C = \lim_{n \rightarrow \infty} \frac{\log_2 |S \cap \{A, C, G, T\}^n|}{n}.$$

For run-length ≤ 3 , $C \approx 1.94$ bits/symbol (vs 2 unconstrained). GC-balanced adds another small penalty.

Open problem: Efficient encoding/decoding for *joint* run-length + GC constraints, [?].

10.6 Code: VT Code Encoding/Decoding

[?]

Listing 10.1: VT code single-deletion correction

```

0 def vt_encode(message_bits, n):
1     """Encode message into VT_0(n) codeword."""
2     # Systematic encoding: find redundancy bits to satisfy sum(i*x_i) \
      # equiv 0 (mod n+1) \cite{barlev2022dna}
3     # This is a sketch; real implementation uses systematic encoding
      # algorithms \cite{barlev2022dna}
4     pass
5
6 def vt_decode(received, n):
7     """Decode received word with at most one deletion."""
8     L = len(received)
9     if L == n:
10        # No deletion, check syndrome
11        syndrome = sum((i+1)*int(b) for i,b in enumerate(received)) % (n
12        +1)
13        if syndrome == 0:

```

```

13         return received # Valid codeword
14         # Single substitution error - correctable
15     elif L == n-1:
16         # One deletion - find position
17         syndrome = sum((i+1)*int(b) for i,b in enumerate(received)) % (n
18             +1)
19         # Syndrome gives position of deleted bit
20         # Insert 0 or 1 at that position to satisfy VT constraint
21     return None
22
23 # Test
24 n = 10
25 codeword = vt_encode([1,0,1,1], n)
26 print("Codeword:", codeword)
27 # Simulate deletion
28 import random
29 del_pos = random.randint(0, n-1)
30 received = codeword[:del_pos] + codeword[del_pos+1:]
31 decoded = vt_decode(received, n)
32 print("Decoded:", decoded)

```

10.7 Bridge to Chapter 11

DNA storage coding theory solves *channel-specific* problems. Chapter 11 addresses *system-level* design: how to compose complex systems from simpler components using operads—a categorical framework for syntax and semantics of system composition, [?].

10.8 Further Reading

- Bar-Lev, Sabary, Yaakobi (2022). *Exciting Coding Problems for DNA-based Storage*. Notices, 69(11).
- Sloane (2002). *On Single-Deletion-Correcting Codes*. IEEE Trans. Inf. Theory.
- Church et al. (2012). *Next-Generation Digital Information Storage in DNA*. Science.
- Carmean et al. (2019). *DNA Data Storage and Hybrid Molecular–Electronic Computing*. IEEE.

10.9 Exercises

10.10 Solutions

Solution.Exercise 1 (Edit distance): The edit distance between two DNA sequences is the minimum number of insertions, deletions, and substitutions needed to transform one into the other. For sequences of length n , the dynamic programming algorithm runs in $O(n^2)$ time.

Solution.Exercise 2 (GC-content): To enforce GC-content between 40% and 60%, we use run-length limited codes. A simple construction is to concatenate blocks of fixed composition and add parity bits.

Solution.Exercise 3 (Sequence reconstruction): The sequence reconstruction problem asks to recover a codeword from multiple noisy reads. For the deletion channel, the Berlekamp-Massey algorithm can be adapted to find the original sequence.

Solution.Exercise 4 (Research): The capacity of the DNA storage channel is an open problem. Recent work using information theory and combinatorial constructions has established lower bounds, but the exact capacity remains unknown.

Exercise 10.1. Implement the VT syndrome decoder for $n = 15$. Test on all possible single-deletion errors for all codewords in $VT_0(15)$.

Exercise 10.2. Design a run-length limited code (no AAAA, CCCC, GGGG, TTTT) with rate ≥ 1.9 bits/symbol. Implement encoder/decoder.

Exercise 10.3. Simulate the DNA clustering problem: generate 1000 strands with 100 copies each, add 3% edit errors, and cluster using edit distance + hierarchical clustering. What fraction of clusters are pure? [?]

Exercise 10.4 (Research). Investigate polar codes for edit-distance channels. Can the polarization phenomenon be adapted to insertion/deletion channels?

CHAPTER 11

The Mathematics of Cyber Defense

[?]

John A. Emanuello, Ahmad Ridley

Communicated by *Notices* Associate Editor Emilie Purvine

“The speed, complexity, and ubiquity of cyber-attacks has never been more apparent... Current cyber-defense capabilities are static and rules-based... This approach is unsustainable.”

— Emanuello & Ridley

[?]

Abstract. Cyber defense requires detecting anomalies in massive, heterogeneous data streams (logs, netflow, alerts) and responding autonomously. This chapter surveys the mathematical foundations: autoencoders for unsupervised anomaly detection, Word2Vec-style embeddings for categorical log data, and reinforcement learning for autonomous response. The unique challenges: adversarial dynamics, non-stationary distributions, and the need for human-machine teaming, [?].

11.1 Introduction: Why Mathematics for Cyber?

Mathematical notation follows standard conventions. Key terms are defined in the glossary. [?]

Cyber attacks are not single events but *campaigns*: multi-phase, adaptive, stealthy. Defenders face:

- **Volume:** Terabytes/day of logs, netflow, packet captures
- **Heterogeneity:** Categorical (IPs, usernames), temporal, graph-structured
- **Adversariality:** Attackers adapt to defenses
- **Label scarcity:** Almost no labeled attacks; must detect *anomalies*

Mathematical approach: **Build models of normal behavior; flag deviations.**

11.2 Data Representation: Embedding Categorical Features

[?]

Log data is nominal: IP addresses, usernames, process names, file paths. Standard ML needs numeric vectors.

Word2Vec for logs (Mikolov et al., 2013): Treat a log line as a “sentence,” tokens as “words.” Train skip-gram model:

$$\max_{\theta} \sum_{(w,c) \in D} \log P(c|w; \theta), \quad P(c|w) = \frac{\exp(\mathbf{v}_c^\top \mathbf{v}_w)}{\sum_{c'} \exp(\mathbf{v}_{c'}^\top \mathbf{v}_w)}.$$

Result: Each token (IP, process, etc.) gets a dense vector $\mathbf{v} \in \mathbb{R}^d$. Similar behavior $\rightarrow$ nearby vectors.

Listing 11.1: Log embedding with Word2Vec

```

0 from gensim.models import Word2Vec
1 import numpy as np
2
3 # Parse logs into token sequences
4 log_sequences = [
5     ["192.168.1.1", "ssh", "accept", "user1"],
6     ["192.168.1.2", "ssh", "failed", "user2"],
7     # ...
8 ]
9
10 # Train Word2Vec
11 model = Word2Vec(log_sequences, vector_size=100, window=5, min_count=1,
12                  workers=4)
13
14 # Embed a log line
15 def embed_log(tokens, model):
16     vecs = [model.wv[t] for t in tokens if t in model.wv]
17     return np.mean(vecs, axis=0) if vecs else np.zeros(model.vector_size)
18
19 # Now each log is a 100-d vector for downstream ML

```

11.3 Anomaly Detection: Deep Autoencoders

[?]

Deep Autoencoder (AE): Multi-layer perceptron trained to reconstruct input:

$$\mathbf{z} = \sigma(\mathbf{W}_e \mathbf{x} + \mathbf{b}_e), \quad \hat{\mathbf{x}} = \sigma(\mathbf{W}_d \mathbf{z} + \mathbf{b}_d),$$

loss = $\|\mathbf{x} - \hat{\mathbf{x}}\|^2$. The bottleneck $\mathbf{z}$ learns a compressed representation.

Unsupervised anomaly detection: Train on *normal* traffic only. At test time, high reconstruction error $\rightarrow$ anomaly, [?].

Theorem 11.1 (AE as Nonlinear PCA). *A deep AE with linear activations learns the PCA subspace. With nonlinearities, it learns a nonlinear manifold of normal data.*

Listing 11.2: Autoencoder for log anomaly detection

```

1 import torch
2 import torch.nn as nn
3
4 class LogAutoencoder(nn.Module):
5     def __init__(self, input_dim, latent_dim=32):
6         super().__init__()
7         self.encoder = nn.Sequential(
8             nn.Linear(input_dim, 128), nn.ReLU(),
9             nn.Linear(128, 64), nn.ReLU(),
10            nn.Linear(64, latent_dim)
11        )
12        self.decoder = nn.Sequential(
13            nn.Linear(latent_dim, 64), nn.ReLU(),
14            nn.Linear(64, 128), nn.ReLU(),
15            nn.Linear(128, input_dim)
16        )
17
18        def forward(self, x):
19            z = self.encoder(x)
20            return self.decoder(z), z
21
22 # Training on normal data only
23 model = LogAutoencoder(input_dim=100)
24 optimizer = torch.optim.Adam(model.parameters(), lr=1e-3)
25
26 for epoch in range(50):
27     for batch in normal_dataloader:
28         recon, _ = model(batch)
29         loss = nn.MSELoss()(recon, batch)
30         loss.backward()
31         optimizer.step()
32         optimizer.zero_grad()
33
34 # Anomaly score = reconstruction error
35 def anomaly_score(model, x):
36     with torch.no_grad():
37         recon, _ = model(x)
38         return torch.mean((x - recon)**2, dim=1)

```

11.4 Chaining Anomalies Across Modalities

[?]

A single anomaly (e.g., unusual process) may be benign. A *chain* across modalities (unusual process + unusual netflow + unusual file access) is likely an attack.

Challenge: Fuse AE scores from log embeddings, netflow embeddings, alert embeddings into a unified timeline.

Approach: **Temporal point processes** or **graph neural networks** on the alert correlation graph.

11.5 Autonomous Response: Reinforcement Learning

[?]

Detection is not enough; must *respond*. Formulate as **Markov Decision Process**:

- State s_t : Network state, active alerts, system health
- Action a_t : Block IP, isolate host, rate-limit, collect forensics
- Reward r_t : Negative for damage, positive for containment, cost for disruption

Learn policy $\pi(a|s)$ via **Deep Q-Network (DQN)** or **Proximal Policy Optimization (PPO)**.

Listing 11.3: RL environment for cyber response

```

1 import gymnasium as gym
2 import numpy as np
3
4 class CyberEnv(gym.Env):
5     def __init__(self):
6         # State: [num_alerts, severity_scores, network_health, ...]
7         self.observation_space = gym.spaces.Box(-np.inf, np.inf, (20,))
8         # Actions: 0=monitor, 1=block_ip, 2=isolate_host, 3=rate_limit, 4=
          forensics
9         self.action_space = gym.spaces.Discrete(5)
10        self.state = None
11        self.attack_ongoing = False
12
13    def step(self, action):
14        # Simulate attack progression and response effect
15        reward = self._compute_reward(action)
16        terminated = self.attack_neutralized()
17        truncated = self.steps >= 100
18        return self.state, reward, terminated, truncated, {}
19
20    def reset(self, seed=None):
21        self.state = np.zeros(20)
22        self.attack_ongoing = np.random.rand() < 0.3
23        return self.state, {}

```

Human-in-the-loop: RL agent proposes actions; human analyst approves/modifies. Over time, trust increases → more autonomy.

11.6 Bridge to Chapter 12

[?]

Cyber defense uses deep learning for *pattern recognition* in dynamic systems. Chapter 12 uses **operads** (categorical algebra) for *compositional design* of complex systems. Both address the challenge of *structure in complexity*: one learns it, the other specifies it.

11.7 Further Reading

[?]

- Emanuello, Ridley (2022). *The Mathematics of Cyber Defense*. Notices, 69(6).
- Golczynski, Emanuello (2021). *End-to-End Anomaly Detection....* arXiv:2108.12276.
- Sutton, Barto (2018). *Reinforcement Learning*. MIT Press.
- Goodfellow et al. (2016). *Deep Learning*. MIT Press (Ch. 14 on AEs).
- MITRE ATT&CK framework: <https://attack.mitre.org>

11.8 Exercises

11.9 Solutions

Solution.Exercise 1 (Autoencoder): An autoencoder with hidden dimension $k < n$ learns to compress log data. The reconstruction error $\|x - \hat{x}\|^2$ identifies anomalies: large errors indicate unusual patterns.

Solution.Exercise 2 (Word2Vec): Word2Vec learns embeddings for categorical data (IP addresses, usernames). The cosine similarity between embeddings captures semantic relationships: similar IPs have similar usage patterns.

Solution.Exercise 3 (RL for defense): A reinforcement learning agent learns to respond to attacks by maximizing a reward function. The state space includes network metrics; the action space includes firewall rules and alerts.

Solution.Exercise 4 (Research): Adversarial attacks on anomaly detection systems can fool autoencoders by adding small perturbations. Defense mechanisms include adversarial training and ensemble methods.

[?]

Exercise 11.1. Train a Word2Vec model on the HDFS log dataset. Visualize the embedding of IP addresses—do malicious IPs cluster?

Exercise 11.2. Implement a variational autoencoder (VAE) for log anomalies. Compare reconstruction error vs. ELBO as anomaly scores.

Exercise 11.3. Design a reward function for a cyber RL agent that balances: (1) minimizing data exfiltration, (2) minimizing false positive isolations, (3) minimizing analyst workload, [?].

Exercise 11.4 (Research). Investigate *adversarial attacks on the anomaly detector itself*: can an attacker craft traffic that looks normal to the AE but is malicious? [?]

CHAPTER 12

The Mathematics of Cyber Defense

[?]

John A. Emanuello, Ahmad Ridley

Communicated by *Notices* Associate Editor Emilie Purvine

“The speed, complexity, and ubiquity of cyber-attacks has never been more apparent... Current cyber-defense capabilities are static and rules-based... This approach is unsustainable.”

— Emanuello & Ridley

[?]

Abstract. Cyber defense requires detecting anomalies in massive, heterogeneous data streams (logs, netflow, alerts) and responding autonomously. This chapter surveys the mathematical foundations: autoencoders for unsupervised anomaly detection, Word2Vec-style embeddings for categorical log data, and reinforcement learning for autonomous response. The unique challenges: adversarial dynamics, non-stationary distributions, and the need for human-machine teaming, [?].

12.1 Introduction: Why Mathematics for Cyber?

Mathematical notation follows standard conventions. Key terms are defined in the glossary. [?]

Cyber attacks are not single events but *campaigns*: multi-phase, adaptive, stealthy. Defenders face:

- **Volume:** Terabytes/day of logs, netflow, packet captures
- **Heterogeneity:** Categorical (IPs, usernames), temporal, graph-structured
- **Adversariality:** Attackers adapt to defenses
- **Label scarcity:** Almost no labeled attacks; must detect *anomalies*

Mathematical approach: **Build models of normal behavior; flag deviations.**

12.2 Data Representation: Embedding Categorical Features

[?]

Log data is nominal: IP addresses, usernames, process names, file paths. Standard ML needs numeric vectors.

Word2Vec for logs (Mikolov et al., 2013): Treat a log line as a “sentence,” tokens as “words.” Train skip-gram model:

$$\max_{\theta} \sum_{(w,c) \in D} \log P(c|w; \theta), \quad P(c|w) = \frac{\exp(\mathbf{v}_c^\top \mathbf{v}_w)}{\sum_{c'} \exp(\mathbf{v}_{c'}^\top \mathbf{v}_w)}.$$

Result: Each token (IP, process, etc.) gets a dense vector $\mathbf{v} \in \mathbb{R}^d$. Similar behavior $\rightarrow$ nearby vectors.

Listing 12.1: Log embedding with Word2Vec

```

0 from gensim.models import Word2Vec
1 import numpy as np
2
3 # Parse logs into token sequences
4 log_sequences = [
5     ["192.168.1.1", "ssh", "accept", "user1"],
6     ["192.168.1.2", "ssh", "failed", "user2"],
7     # ...
8 ]
9
10 # Train Word2Vec
11 model = Word2Vec(log_sequences, vector_size=100, window=5, min_count=1,
12                  workers=4)
13
14 # Embed a log line
15 def embed_log(tokens, model):
16     vecs = [model.wv[t] for t in tokens if t in model.wv]
17     return np.mean(vecs, axis=0) if vecs else np.zeros(model.vector_size)
18
19 # Now each log is a 100-d vector for downstream ML

```

12.3 Anomaly Detection: Deep Autoencoders

[?]

Deep Autoencoder (AE): Multi-layer perceptron trained to reconstruct input:

$$\mathbf{z} = \sigma(\mathbf{W}_e \mathbf{x} + \mathbf{b}_e), \quad \hat{\mathbf{x}} = \sigma(\mathbf{W}_d \mathbf{z} + \mathbf{b}_d),$$

loss = $\|\mathbf{x} - \hat{\mathbf{x}}\|^2$. The bottleneck $\mathbf{z}$ learns a compressed representation.

Unsupervised anomaly detection: Train on *normal* traffic only. At test time, high reconstruction error $\rightarrow$ anomaly, [?].

Theorem 12.1 (AE as Nonlinear PCA). *A deep AE with linear activations learns the PCA subspace. With nonlinearities, it learns a nonlinear manifold of normal data.*

Listing 12.2: Autoencoder for log anomaly detection

```

1 import torch
2 import torch.nn as nn
3
4 class LogAutoencoder(nn.Module):
5     def __init__(self, input_dim, latent_dim=32):
6         super().__init__()
7         self.encoder = nn.Sequential(
8             nn.Linear(input_dim, 128), nn.ReLU(),
9             nn.Linear(128, 64), nn.ReLU(),
10            nn.Linear(64, latent_dim)
11        )
12        self.decoder = nn.Sequential(
13            nn.Linear(latent_dim, 64), nn.ReLU(),
14            nn.Linear(64, 128), nn.ReLU(),
15            nn.Linear(128, input_dim)
16        )
17
18        def forward(self, x):
19            z = self.encoder(x)
20            return self.decoder(z), z
21
22 # Training on normal data only
23 model = LogAutoencoder(input_dim=100)
24 optimizer = torch.optim.Adam(model.parameters(), lr=1e-3)
25
26 for epoch in range(50):
27     for batch in normal_dataloader:
28         recon, _ = model(batch)
29         loss = nn.MSELoss()(recon, batch)
30         loss.backward()
31         optimizer.step()
32         optimizer.zero_grad()
33
34 # Anomaly score = reconstruction error
35 def anomaly_score(model, x):
36     with torch.no_grad():
37         recon, _ = model(x)
38         return torch.mean((x - recon)**2, dim=1)

```

12.4 Chaining Anomalies Across Modalities

[?]

A single anomaly (e.g., unusual process) may be benign. A *chain* across modalities (unusual process + unusual netflow + unusual file access) is likely an attack.

Challenge: Fuse AE scores from log embeddings, netflow embeddings, alert embeddings into a unified timeline.

Approach: **Temporal point processes** or **graph neural networks** on the alert correlation graph.

12.5 Autonomous Response: Reinforcement Learning

[?]

Detection is not enough; must *respond*. Formulate as **Markov Decision Process**:

- State s_t : Network state, active alerts, system health
- Action a_t : Block IP, isolate host, rate-limit, collect forensics
- Reward r_t : Negative for damage, positive for containment, cost for disruption

Learn policy $\pi(a|s)$ via **Deep Q-Network (DQN)** or **Proximal Policy Optimization (PPO)**.

Listing 12.3: RL environment for cyber response

```

1 import gymnasium as gym
2 import numpy as np
3
4 class CyberEnv(gym.Env):
5     def __init__(self):
6         # State: [num_alerts, severity_scores, network_health, ...]
7         self.observation_space = gym.spaces.Box(-np.inf, np.inf, (20,))
8         # Actions: 0=monitor, 1=block_ip, 2=isolate_host, 3=rate_limit, 4=
          forensics
9         self.action_space = gym.spaces.Discrete(5)
10        self.state = None
11        self.attack_ongoing = False
12
13    def step(self, action):
14        # Simulate attack progression and response effect
15        reward = self._compute_reward(action)
16        terminated = self.attack_neutralized()
17        truncated = self.steps >= 100
18        return self.state, reward, terminated, truncated, {}
19
20    def reset(self, seed=None):
21        self.state = np.zeros(20)
22        self.attack_ongoing = np.random.rand() < 0.3
23        return self.state, {}

```

Human-in-the-loop: RL agent proposes actions; human analyst approves/modifies. Over time, trust increases → more autonomy.

12.6 Bridge to Chapter 12

[?]

Cyber defense uses deep learning for *pattern recognition* in dynamic systems. Chapter 12 uses **operads** (categorical algebra) for *compositional design* of complex systems. Both address the challenge of *structure in complexity*: one learns it, the other specifies it.

12.7 Further Reading

[?]

- Emanuello, Ridley (2022). *The Mathematics of Cyber Defense*. Notices, 69(6).
- Golczynski, Emanuello (2021). *End-to-End Anomaly Detection....* arXiv:2108.12276.
- Sutton, Barto (2018). *Reinforcement Learning*. MIT Press.
- Goodfellow et al. (2016). *Deep Learning*. MIT Press (Ch. 14 on AEs).
- MITRE ATT&CK framework: <https://attack.mitre.org>

12.8 Exercises

12.9 Solutions

Solution.Exercise 1 (Autoencoder): An autoencoder with hidden dimension $k < n$ learns to compress log data. The reconstruction error $\|x - \hat{x}\|^2$ identifies anomalies: large errors indicate unusual patterns.

Solution.Exercise 2 (Word2Vec): Word2Vec learns embeddings for categorical data (IP addresses, usernames). The cosine similarity between embeddings captures semantic relationships: similar IPs have similar usage patterns.

Solution.Exercise 3 (RL for defense): A reinforcement learning agent learns to respond to attacks by maximizing a reward function. The state space includes network metrics; the action space includes firewall rules and alerts.

Solution.Exercise 4 (Research): Adversarial attacks on anomaly detection systems can fool autoencoders by adding small perturbations. Defense mechanisms include adversarial training and ensemble methods.

[?]

Exercise 12.1. Train a Word2Vec model on the HDFS log dataset. Visualize the embedding of IP addresses—do malicious IPs cluster?

Exercise 12.2. Implement a variational autoencoder (VAE) for log anomalies. Compare reconstruction error vs. ELBO as anomaly scores.

Exercise 12.3. Design a reward function for a cyber RL agent that balances: (1) minimizing data exfiltration, (2) minimizing false positive isolations, (3) minimizing analyst workload, [?].

Exercise 12.4 (Research). Investigate *adversarial attacks on the anomaly detector itself*: can an attacker craft traffic that looks normal to the AE but is malicious? [?]

CHAPTER 13

Operads for Designing Systems of Systems

[?]

John C. Baez, John D. Foley

Communicated by *Notices* Associate Editor Emilie Purvine

“Analogous to how a group abstracts the symmetries of an object, an operad abstracts operations that compose many objects into a single object.”

— Baez & Foley

[?]

Abstract. “Systems of systems” engineering composes independent systems (radar, comms, weapons, sensors) into a larger capability. This chapter introduces **operads**—a categorical tool that formalizes the *syntax* of composition (how things plug together) and their *algebras* (the concrete semantics). Applied to network design, search-and-rescue planning, and task coordination, [?].

13.1 Introduction: The Composition Problem

Mathematical notation follows standard conventions. Key terms are defined in the glossary.

A smartphone is a system of systems: touchscreen, camera, CPU, radio, GPS, battery... The *design* is a blueprint for how components connect. The *behavior* depends on the specific hardware.

In defense (DARPA CASCADE program), we need to compose:

- Aircraft with range-limited comms
- Ships, helicopters, drones for search-and-rescue
- Sensors, effectors, decision-makers for task planning

Challenge: The composition rules (syntax) are different from the physical realization (semantics). Operads separate these, [?].

13.2 What Is an Operad?

Definition 13.1 (Colored Operad). An *operad* $\mathcal{O}$ consists of: [?]

- A set of **types** (colors) $\text{Ob}(\mathcal{O})$
- For each tuple $(X_1, \dots, X_n; Y)$ with $X_i, Y \in \text{Ob}(\mathcal{O})$, a set of **operations** $\mathcal{O}(X_1, \dots, X_n; Y)$
- **Composition**: Given $f \in \mathcal{O}(X_1, \dots, X_n; Y)$ and $g_i \in \mathcal{O}(Z_{i1}, \dots, Z_{ik_i}; X_i)$,

$$f \circ (g_1, \dots, g_n) \in \mathcal{O}(Z_{11}, \dots, Z_{nk_n}; Y)$$

- **Identities**: $\text{id}_X \in \mathcal{O}(X; X)$
- **Symmetric group actions**: Permuting inputs

Satisfying associativity, identity, equivariance axioms.

Definition 13.2 (Operad Algebra). An **algebra** of $\mathcal{O}$ assigns to each type X a set $A(X)$ and to each operation $f \in \mathcal{O}(X_1, \dots, X_n; Y)$ a function

$$A(f) : A(X_1) \times \dots \times A(X_n) \rightarrow A(Y)$$

preserving composition, identities, and symmetries, [?].

Intuition: Operad = *syntax* (how to compose); Algebra = *semantics* (what the components actually are).

13.3 Network Operads

Example 13.3 (Aircraft Communication Network). **Types**: Number of aircraft $n \in \mathbb{N}$.

Operations: $\mathcal{O}(n_1, \dots, n_k; m) =$ graphs with $n_1 + \dots + n_k$ input vertices and m output vertices. An operation tells you which new edges to add to connect the sub-networks.

Algebra: For each n , $A(n) =$ set of possible states of n aircraft (positions, velocities, comms status). An operation acts by adding comms links between aircraft that are in range.

Theorem 13.4 (Network Operad Construction). For any class of networks (graphs, directed graphs, typed edges, etc.), there is a network operad whose operations are “ways to connect networks by adding edges.” [?]

13.4 Application: Search and Rescue

Problem: Distribute search assets (ships, planes, helicopters, small boats) to cover a large area after a disaster (Fastnet Race 1979, Sydney-Hobart 1998).

1. **Design operad**: Assets have “slots” (helicopter carries 3 boats). Operations = ways to nest assets, [?].
2. **Algebra**: Actual assets at actual locations with fuel constraints.
3. **Optimization**: Search over operad operations to maximize “search effort” (area covered per time) within budget, [?].

The operad structure lets you *search the space of designs* algorithmically, [?].

13.5 From Design to Tasking

Same operad, different algebra: [?]

- **Design algebra:** Assets at locations $\rightarrow$ network structure
- **Tasking algebra:** Primitive tasks (“fly to”, “search”, “rescue”) composed into mission plans

Catalyst operads (Baez et al.): Some operations enable other operations (a “catalyst” agent coordinates multiple agents). Used for multi-agent task plans, [?].

Translation to **constraint programming** (MILP/CPLEX): Each operation $\rightarrow$ decision variables; composition $\rightarrow$ constraints, [?].

13.6 Levels of Abstraction

A system design proceeds through levels:

1. **Abstract:** “Network of boats with comms”
2. **Concrete:** “3 RHIBs, 2 helicopters, SATCOM links”
3. **Detailed:** Specific models, frequencies, protocols

Operad homomorphisms and **algebra homomorphisms** formalize moving between levels. This is *change of abstraction*.

13.7 Code: Simple Operad in Python

Listing 13.1: Network operad sketch

```

1 from dataclasses import dataclass
2 from typing import List, Dict, Set, Callable
3 from abc import ABC, abstractmethod
4
5 @dataclass
6 class NetworkOperation:
7     """Operation: connect input networks by adding edges"""
8     input_types: List[int] # sizes of input networks
9     output_type: int # size of output network
10    edges: Set[tuple] # (source_idx, target_idx) where idx is global
11
12    def __call__(self, *networks):
13        # Compose networks according to edges
14        # networks[i] has size input_types[i]
15        pass
16
17 class Operad:
18     def __init__(self):
19         self.operations: Dict[tuple, List[NetworkOperation]] = {}
20
21     def add(self, op: NetworkOperation):
22         key = (tuple(op.input_types), op.output_type)

```

```

23     self.operations.setdefault(key, []).append(op)
24
25     def compose(self, outer: NetworkOperation, inner: List[
26         NetworkOperation]):
27         # Build composed operation
28         # Offset inner edge indices by input network sizes
29         pass
30
31 class Algebra:
32     def __init__(self, interpret_type: Callable[[int], set]):
33         self.interpret_type = interpret_type # n -> set of n-aircraft
34         states
35
36     def apply(self, op: NetworkOperation, *states):
37         # Apply operation to concrete states
38         pass
39
40 # Example: Aircraft comms
41 network_operad = Operad() \cite{baez2021operads}
42 # Add operations for connecting networks...
43
44 # Algebra: real aircraft positions and comms ranges
45 def aircraft_states(n):
46     # Return set of possible states for n aircraft
47     # Each state: positions, velocities, comms equipment
48     pass
49
50 aircraft_algebra = Algebra(aircraft_states)

```

13.8 Bridge to Chapter 13

Operads provide a *static* compositional structure for system design. Chapter 13 applies *dynamic* Bayesian inference to a different domain: astronomy. Both deal with *hierarchical structure*—operads explicitly, astronomy through the hierarchy of data → model → cosmology, [?].

13.9 Further Reading

- Baez, Foley (2021). *Operads for Designing Systems of Systems*. Notices, 68(6).
- Baez, Foley, Moeller, Pollard (2020). *Network Models*. Theory Appl. Categ.
- Spivak (2013). *The Operad of Wiring Diagrams*. arXiv:1305.0297.
- Spivak (2014). *Category Theory for the Sciences*. MIT Press.
- Yau (2018). *Operads of Wiring Diagrams*. Springer.

13.10 Exercises

13.11 Solutions

Solution.Exercise 1 (Operad composition): The composition operation in an operad satisfies associativity: $(f \circ_i g) \circ_j h = f \circ_i (g \circ_j h)$. This ensures that composing operations in different orders gives the same result.

Solution.Exercise 2 (Network operad): The network operad has operations corresponding to network topologies. Composition corresponds to substituting one network into a node of another.

Solution.Exercise 3 (DARPA CASCADE): The CASCADE program uses operads to compose heterogeneous systems. Each system is an algebra over the operad, and composition is given by the operad action.

Solution.Exercise 4 (Research): Operadic methods can be extended to stochastic systems and to systems with uncertainty. The compositional structure allows local verification to imply global properties.

Exercise 13.1. Define the operad of *wiring diagrams* for digital circuits. Types = natural numbers (wire count). Operations = circuits with given input/output wires, [?].

Exercise 13.2. For the aircraft comms operad, define an algebra where $A(n)$ is the set of possible communication graphs on n aircraft with given range constraints, [?].

Exercise 13.3. Design a “catalyst” operation for a search-and-rescue scenario where a helicopter coordinates the deployment of 3 small boats.

Exercise 13.4 (Research). Extend the network operad to include *temporal* composition: operations that describe not just spatial connections but time-dependent interactions (e.g., intermittent connectivity), [?].

Part V

Data-Intensive Science

CHAPTER 14

A Brief History of Inference in Astronomy

[?]

Rafael S. de Souza, Emille E. O. Ishida, Alberto Krone-Martins
Communicated by *Notices* Associate Editor Richard A. Levine

“Throughout history, methods of inference have been deeply linked to advances in astronomy.”

— de Souza, Ishida, Krone-Martins

[?]

Abstract. From Hipparchus averaging solstice times to Hamiltonian Monte Carlo sampling cosmological parameters, astronomy has driven the development of statistical inference. This chapter traces the arc: least squares (Gauss/Ceres), Bayesian revival (Jeffreys), MCMC (cosmology), to modern likelihood-free and amortized inference. The lesson: astronomical data challenges force statistical innovation, [?].

14.1 Introduction: Astronomy as an Inference Engine

Mathematical notation follows standard conventions. Key terms are defined in the glossary.

Astronomy is fundamentally an **inverse problem**: we observe photons (or gravitational waves, neutrinos) and infer the physical universe. Every major advance in data → model inference has an astronomical milestone: [?]

Hipparchus (190–120 BC)	Averaging → noise reduction
Tycho Brahe (1546–1601)	Repeated observations → precision
Gauss (1801)	Least squares → orbit of Ceres
Laplace (1812)	Central Limit Theorem → error theory
Jeffreys (1939)	Bayesian theory → cosmology
[?] MCMC (1990s–)	Cosmological parameter estimation
HMC/Nested Sampling (2000s–)	High-dimensional posteriors
Likelihood-free (2010s–)	Simulator-based inference
[?] Amortized (2020s–)	Neural posterior estimation

14.2 Linear Regression and Least Squares

Gauss (1801) predicted the orbit of Ceres using **least squares**. The problem: fit an ellipse to noisy angular observations.

Model: $\mathbf{y} = \mathbf{X}\boldsymbol{\theta} + \boldsymbol{\epsilon}$, $\boldsymbol{\epsilon} \sim \mathcal{N}(0, \boldsymbol{\Sigma})$.

MLE/least squares: $\hat{\boldsymbol{\theta}} = (\mathbf{X}^\top \boldsymbol{\Sigma}^{-1} \mathbf{X})^{-1} \mathbf{X}^\top \boldsymbol{\Sigma}^{-1} \mathbf{y}$.

This is still the workhorse for:

- Hubble diagram (distance vs. redshift)
- Black hole mass scaling relations
- Spectral line fitting

14.3 Bayesian Inference in Cosmology

[?]

Type Ia supernovae (1998) revealed cosmic acceleration. The inference problem: [?]

- Data: $\{\mu_i, z_i\}$ = distance modulus + redshift
- Model: $\mu(z; \Omega_m, \Omega_\Lambda, H_0)$ from Friedmann eqs
- Posterior: $p(\boldsymbol{\theta}|\mathcal{D}) \propto p(\mathcal{D}|\boldsymbol{\theta})p(\boldsymbol{\theta})$

Computational challenge: 6–10 parameters, multimodal, degenerate. MCMC (Metropolis-Hastings) became standard.

14.4 Modern Sampling: HMC and Nested Sampling

Hamiltonian Monte Carlo Uses gradient of log-posterior to propose distant states with high acceptance. Essential for ~ 100 -parameter models (e.g., Planck CMB).

Nested Sampling Computes evidence $Z = \int p(\mathcal{D}|\boldsymbol{\theta})p(\boldsymbol{\theta})d\boldsymbol{\theta}$ for model comparison (e.g., Λ CDM vs. w CDM).

Listing 14.1: HMC for cosmological parameters

```

0 import numpy as np
1 import jax.numpy as jnp
2 from jax import grad, jit
3 import numpyro
4 import numpyro.distributions as dist
5 from numpyro.infer import MCMC, NUTS
6
7 def cosmological_model(obs_mu, obs_z, obs_sigma):
8     # Flat priors
9     Omega_m = numpyro.sample('Omega_m', dist.Uniform(0.0, 1.0))
10    Omega_L = numpyro.sample('Omega_L', dist.Uniform(0.0, 1.0))
11    H0 = numpyro.sample('H0', dist.Uniform(50.0, 80.0))
12
13    # Distance modulus from cosmology

```

```

14 def dist_modulus(z):
15     # Simplified; real version integrates Friedmann eq
16     return 5*jnp.log10(luminosity_distance(z, Omega_m, Omega_L, H0)) +
        25
17
18     mu_pred = jnp.array([dist_modulus(z) for z in obs_z])
19
20     # Likelihood
21     numpyro.sample('obs', dist.Normal(mu_pred, obs_sigma), obs=obs_mu)
22
23 # Run NUTS (HMC)
24 kernel = NUTS(cosmological_model)
25 mcmc = MCMC(kernel, num_warmup=1000, num_samples=2000)
26 mcmc.run(jax.random.PRNGKey(0), obs_mu, obs_z, obs_sigma)
27 samples = mcmc.get_samples()

```

14.5 Likelihood-Free Inference (LFI)

Many astronomical models have *intractable likelihoods* but are easy to *simulate*:

- Galaxy formation (hydrodynamics + gravity)
- Gravitational waveforms (numerical relativity)
- Strong lensing (ray tracing)

Simulation-based inference: Learn $p(\theta|\mathcal{D})$ without evaluating $p(\mathcal{D}|\theta)$, [?].

Approximate Bayesian Computation (ABC) Accept θ if $d(\mathcal{D}_{\text{sim}}, \mathcal{D}_{\text{obs}}) < \epsilon$ [?]

Neural Density Estimation Train normalizing flow on $(\theta, \mathcal{D}_{\text{sim}})$ pairs

Score-based Learn score function $\nabla_{\mathcal{D}} \log p(\mathcal{D}|\theta)$

14.6 Amortized Inference

Amortization: Instead of running MCMC for each new dataset, train a neural network once to output the posterior:

$$\phi^* = \arg \min_{\phi} \mathbb{E}_{\theta, \mathcal{D} \sim p(\theta)p(\mathcal{D}|\theta)} [D_{\text{KL}}(p(\theta|\mathcal{D}) || q_{\phi}(\theta|\mathcal{D}))].$$

Examples: `swyft`, `sbi`, `zuko` packages. Critical for LSST (10M alerts/night) and gravitational wave real-time analysis.

14.7 Bridge to Chapter 14

Astronomy drives *methodological* innovation. Chapter 14 shows how *engineering* principles (FAIR) preserve the software that implements these methods. Both are about *longevity*: astronomical methods last centuries; FAIR ensures code lasts decades.

14.8 Further Reading

- de Souza, Ishida, Krone-Martins (2025). *A Brief History of Inference in Astronomy*. Notices, 72(10).
- Stigler (1986). *The History of Statistics*. Harvard.
- Trotta (2017). *Bayesian Methods in Cosmology*. arXiv:1701.01467, [?].
- Alsing et al. (2019). *Simulation-Based Inference*. arXiv:1903.00007.
- Cranmer et al. (2020). *The Frontier of Simulation-Based Inference*. PNAS.

14.9 Exercises

14.10 Solutions

Solution.Exercise 1 (Least squares): For the Ceres orbit determination, Gauss used 6 observations to solve for 6 orbital elements. The normal equations $A^T A x = A^T b$ give the least squares solution.

Solution.Exercise 2 (MCMC): Hamiltonian Monte Carlo samples from posterior distributions using gradient information. The leapfrog integrator preserves the Hamiltonian structure, giving high acceptance rates.

Solution.Exercise 3 (Nested sampling): Nested sampling estimates the Bayesian evidence by sampling from nested likelihood contours. It is particularly useful for multimodal distributions.

Solution.Exercise 4 (Research): Amortized inference uses neural networks to approximate the posterior. Once trained, inference is instantaneous, making it suitable for real-time applications.

Exercise 14.1. Implement least-squares orbit fitting for a synthetic asteroid with 6 observations. Compare to Gauss’s original method.

Exercise 14.2. Run MCMC on the Union2.1 supernova dataset (580 SNe Ia) to constrain Ω_m, Ω_Λ . Plot the confidence contours.

Exercise 14.3. Train a normalizing flow (RealNVP) to approximate the posterior of a 2D cosmological model. Compare to MCMC samples.

Exercise 14.4 (Research). Design an amortized inference pipeline for real-time gravitational wave parameter estimation (latency < 1 second), [?].

CHAPTER 15

Making Mathematical Online Resources FAIR

[?]

Tabea Bacher, Marina Garrote-López, Christiane Görgen, Marius J. Neubert
Communicated by *Notices* Associate Editor

“Digital preservation today is guided by the FAIR Principles—Findability, Accessibility, Interoperability, and Reusability. These principles do not insist on the newest and shiniest web tool... much more sensibly, they guide our decisions on selecting methodologies that preserve the long-term value and usability of digitized results.”

— Bacher et al.

[?]

Abstract. Many valuable mathematical databases (OEIS, LMFDB, Small Phylogenetic Trees) were built with early-2000s technology. This chapter presents a case study: modernizing the *Small Phylogenetic Trees* library using FAIR principles. We separate computation from presentation, use OSCAR/Julia for verified computation, and deploy a versioned, documented, licensed resource at algebraicphylogenetics.org, [?].

15.1 Introduction: The Legacy Software Crisis

Mathematical notation follows standard conventions. Key terms are defined in the glossary.

Mathematical software ages differently than papers. A 2005 paper is still readable; a 2005 Maple worksheet may not run. The *Small Phylogenetic Trees* library (Garcia, Puente, 2005) computed phylogenetic invariants for trees up to 5 taxa using Singular and Maple. By 2025:

- Original URLs broken (institution changes)

- Maple/Singular versions incompatible
- No DOI, no license, no API
- Data in GIF images and unstructured text files

FAIR principles (Wilkinson et al., 2016) provide a roadmap: [?]

F Findable: Persistent identifiers (DOI), rich metadata, indexed

A Accessible: Standard protocols (HTTPS, API), authentication if needed

I Interoperable: Standard formats (JSON, MathML, HDF5), controlled vocabularies

R Reusable: Clear license, provenance, community standards

15.2 Case Study: Small Phylogenetic Trees

Original resource: Tables of algebraic invariants (phylogenetic ideals) for group-based models (Jukes-Cantor, Kimura 2/3-parameter) on trees with 3, 4, 5 leaves.

Mathematical content:

- Tree topology T (graph with labeled leaves)
- Nucleotide substitution model $\rightarrow$ polynomial map $\phi_T : \mathbb{C}^p \rightarrow \mathbb{C}^{4^n}$
- Phylogenetic ideal $I_T = \ker \phi_T$ (polynomials vanishing on the model)
- Invariants: generators of I_T , dimension, degree, ML degree

Modernization pipeline:

1. **Computation:** Recompute all invariants in **OSCAR** (Julia), a computer algebra system with formal verification.
2. **Serialization:** Store as **.mrdi** (MaRDI) files—JSON-based, schema-validated, with full provenance (code version, input, timestamp).
3. **Website:** **algebraicphylogenetics.org** built with **Flask** + **MathJax**, serving data from PostgreSQL with REST API.
4. **Versioning:** Git repository for code; Zenodo for data snapshots with DOIs.
5. **License:** CC-BY-4.0 for data, MIT for code.

15.3 The MaRDI Format

The **MaRDI** (Mathematical Research Data Initiative) format structures computational results:

Listing 15.1: Example **.mrdi** file for a phylogenetic invariant

```

0 {
1   "@context": "https://mardi4nfdi.github.io/metadata/schema.json",
2   "@type": "MathematicalComputation",
3   "identifier": "doi:10.5281/zenodo.1234567",

```

```

4  "name": "Phylogenetic invariant for 4-leaf Jukes-Cantor model",
5  "description": "Generator of the phylogenetic ideal I_T...",
6  "hasInput": {
7    "@type": "MathematicalObject",
8    "name": "4-leaf tree topology",
9    "representation": {"format": "newick", "value": "((A,B),(C,D));"}
10 },
11 "hasOutput": {
12   "@type": "MathematicalObject",
13   "name": "Phylogenetic invariant",
14   "representation": {"format": "polynomial", "value": "p0011*p0100 -
15     p0001*p0110 + ..."}
16 },
17 "usedSoftware": {
18   "@type": "Software",
19   "name": "OSCAR",
20   "version": "0.12.0",
21   "url": "https://github.com/oscar-system/OSCAR.jl"
22 },
23 "dateCreated": "2025-03-15T14:32:00Z",
24 "license": "CC-BY-4.0"

```

15.4 Verification and Reproducibility

Key innovation: The new computation is *verified*:

- OSCAR uses exact rational arithmetic (no floating-point errors)
- Gröbner basis computations are certified
- Results cross-checked with Macaulay2 and Singular

The `.mrdi` file contains the *entire provenance*: software versions, input parameters, random seeds. Anyone can re-run and bitwise verify.

Listing 15.2: Verified computation in OSCAR

```

0  using Oscar
1  using Pkg; Pkg.activate("AlgebraicPhylogenetics")
2
3  # Exact computation of phylogenetic ideal for 4-leaf JC69
4  T = phylogenetic_tree("((A,B),(C,D));")
5  model = jukes_cantor_model()
6  I = phylogenetic_ideal(T, model)
7
8  # Verify generators match known results
9  gens = minimal_generators(I)
10 println("Number of generators: ", length(gens))
11 println("Degrees: ", [total_degree(g) for g in gens])
12
13 # Export to MaRDI
14 export_mardi(I, "invariant_4leaf_jc69.mrdi")

```


15.5 Guidelines for FAIRifying Your Resource

[?]

1. **Audit:** What data/code exists? What formats? What dependencies?
2. **Separate concerns:** Computation $\neq$ presentation. Build a *software package* (versioned, tested) + a *website* (consumes package output).
3. **Choose standards:**
 - Data: HDF5, NetCDF, Parquet, or MaRDI JSON
 - Metadata: Schema.org, CodeMeta, MaRDI
 - API: OpenAPI/Swagger
4. **Assign DOIs:** Zenodo, Figshare, or institutional repository
5. **License explicitly:** CC-BY-4.0 (data), MIT/BSD (code)
6. **Document provenance:** Every output links to code version, inputs, environment
7. **Test reproducibility:** CI pipeline that re-runs computations on every commit

15.6 Bridge to Chapter 15

FAIR principles ensure mathematical software survives institutional changes. Chapter 15 applies similar rigor to *tomographic reconstruction*: a complete pipeline from raw projections to readable text, with documented code, versioned data, and reproducibility—exactly the FAIR principles in action for inverse problems, [?].

15.7 Further Reading

- Bacher et al. (2026). *Making Mathematical Online Resources FAIR*. Notices, 73(5), [?].
- Wilkinson et al. (2016). *The FAIR Guiding Principles*. Scientific Data, [?].
- MaRDI Consortium (2024). *MaRDI White Paper*. <https://mardi4nfdi.de>
- Garcia-Puente et al. (2005). *Small Phylogenetic Trees*. <https://algebraicphylogenetics.org>
- OSCAR Computer Algebra System. <https://www.oscar-system.org>

15.8 Exercises

15.9 Solutions

Solution.Exercise 1 (FAIR assessment): To assess FAIRness, check: (1) Does the resource have a persistent identifier? (2) Is the metadata indexed? (3) Are the formats open and interoperable? (4) Is the license clear?

Solution.Exercise 2 (OSCAR): OSCAR uses Julia for verified computation. The workflow is: define computation in Julia, verify with formal methods, deploy as web service.

Solution.Exercise 3 (DOI): A DOI (Digital Object Identifier) provides a persistent link to a resource. The DOI system is managed by the International DOI Foundation.

Solution.Exercise 4 (Research): The challenge of long-term preservation of mathematical software requires standardized formats and active community stewardship.

Exercise 15.1. Take a small mathematical dataset (e.g., a table of integer sequences) and create a MaRDI JSON file for it. Include input, output, software version, and license.

Exercise 15.2. Set up a GitHub Actions workflow that re-runs a Julia/OSCAR computation on every push and verifies the output matches a stored .mrdr file.

Exercise 15.3 (Research). Design a FAIR-compliant schema for storing trained neural network models (weights, architecture, training data provenance, hyperparameters) that enables reproducibility and comparison, [?].

CHAPTER 16

Deciphering Scrolls with Tomography: A Training Experiment

[?]

Sonia Foschiatti, Axel Kittenberger, Otmar Scherzer
Communicated by *Notices* Associate Editor Krystal Taylor

“Virtual unwrapping uses
software to reveal text that was
hidden for centuries, thereby
making it accessible to mankind.”

— Foschiatti, Kittenberger,
Scherzer

Abstract. X-ray computed tomography (CT) enables non-destructive reading of damaged scrolls. This chapter presents a complete didactic pipeline: from simulated X-ray projections (using visible light and transparent films) through filtered back-projection, manual segmentation, and Maximum Intensity Projection (MIP) to readable text. Includes Arduino-controlled rotation stage, Python/MATLAB reconstruction code, and the Archeolab software framework, [?].

16.1 Introduction: The Virtual Unwrapping Problem

Mathematical notation follows standard conventions. Key terms are defined in the glossary.

Damaged scrolls (Herculaneum papyri, En-Gedi scroll, medieval manuscripts) cannot be physically opened. X-ray CT provides a solution: [?]

1. **Scan:** Rotate scroll, collect projections (radiographs)
2. **Reconstruct:** Filtered back-projection $\rightarrow$ 3D volume
3. **Segment:** Identify individual layers in the volume
4. **Texture/Map:** Assign intensity values to surface points
5. **Flatten:** Map 3D surface $\rightarrow$ 2D image

6. Merge: Combine segments into complete text

This chapter implements the full pipeline as an **educational laboratory** using visible light instead of X-rays.

16.2 The Mathematical Model: Radon Transform

[?]

Let $f(x, y)$ be the attenuation coefficient in a slice. A parallel-beam projection at angle θ :

$$p_{\theta}(s) = \int_{-\infty}^{\infty} f(s \cos \theta - t \sin \theta, s \sin \theta + t \cos \theta) dt.$$

This is the **Radon transform** $\mathcal{R}f(\theta, s)$, [?].

Inversion (filtered back-projection):

$$f(x, y) = \int_0^{\pi} [p_{\theta} * h](x \cos \theta + y \sin \theta) d\theta,$$

where h is the **Ram-Lak filter** (ramp filter in Fourier domain): $\hat{h}(\omega) = |\omega|$.

16.3 Experimental Setup: Visible-Light Tomography

Hardware:

- Arduino-controlled stepper motor (rotation stage)
- Projector (visible light source) + Fresnel lens (collimation)
- Camera (detector) + white screen
- Transparent cellulose acetate sheets with printed text (phantom scroll)

Acquisition:

Listing 16.1: Arduino stepper control + camera sync

```

0  # Arduino sketch (simplified)
1  #include <Stepper.h>
2  const int stepsPerRevolution = 200;
3  Stepper myStepper(stepsPerRevolution, 8, 9, 10, 11);
4  int nProjections = 800;
5
6  void setup() {
7      myStepper.setSpeed(10); // rpm
8      Serial.begin(9600);
9  }
10
11 void loop() {
12     if (Serial.available()) {
13         char cmd = Serial.read();
14         if (cmd == 's') { // start acquisition
15             for (int i = 0; i < nProjections; i++) {
16                 myStepper.step(stepsPerRevolution / nProjections);

```

```
17         delay(500); // settle
18         Serial.println("CAPTURE"); // trigger camera
19         delay(500);
20     }
21 }
22 }
23 }
```

16.4 Reconstruction Pipeline

16.4.1 1. Preprocessing

- Dark/flat field correction
- Rotation axis determination (cross-correlation of 0° and 180° projections)
- Cropping to region of interest

16.4.2 2. Filtered Back-Projection (FBP)

Listing 16.2: FBP reconstruction using iradon

```
0 import numpy as np
1 from skimage.transform import iradon
2
3 # projections: (n_angles, n_pixels) array
4 # angles: in degrees
5 reconstruction = iradon(projections, theta=angles,
6                          filter='ramp', # Ram-Lak
7                          interpolation='linear',
8                          circle=False)
9
10 # Post-process: circular mask to remove edge artifacts
11 from skimage.draw import disk
12 mask = np.zeros_like(reconstruction)
13 rr, cc = disk((n/2, n/2), n/2 - 5)
14 mask[rr, cc] = 1
15 reconstruction *= mask
```

16.4.3 3. Segmentation: Finding Layers

Each slice is a cross-section through the scroll. The layers appear as concentric curved lines.

Listing 16.3: Spiral segmentation via genetic algorithm

```
0 import numpy as np
1 from scipy.optimize import differential_evolution
2
3 def spiral_model(theta, params):
4     # params = [center_x, center_y, a, b]
5     cx, cy, a, b = params
6     r = a + b * theta
```

```

7     return cx + r * np.cos(theta), cy + r * np.sin(theta)
8
9 def segmentation_loss(params, slice_img):
10     # Generate spiral mask from params
11     # Compare to image gradients
12     # Return negative correlation
13     pass
14
15 # Optimize per slice
16 result = differential_evolution(segmentation_loss, bounds,
17                                args=(slice_img,), maxiter=100)
18 best_spiral = result.x

```

16.4.4 4. Texture Mapping and Flattening

For each layer, extract intensity values along the spiral, map to a rectangular image.

Listing 16.4: Layer flattening

```

0 def flatten_layer(volume, spiral_params, layer_thickness=3):
1     """Extract and flatten one layer from 3D volume."""
2     h, w, d = volume.shape
3     flattened = np.zeros((d, n_points))
4
5     for z in range(d):
6         slice_img = volume[:, :, z]
7         cx, cy, a, b = spiral_params[z]
8
9         # Sample along spiral
10        theta = np.linspace(0, 4*np.pi, n_points)
11        x = cx + (a + b*theta) * np.cos(theta)
12        y = cy + (a + b*theta) * np.sin(theta)
13
14        # Interpolate
15        from scipy.ndimage import map_coordinates
16        flattened[z, :] = map_coordinates(slice_img, [y, x], order=1)
17
18    return flattened

```

16.4.5 5. Maximum Intensity Projection (MIP)

Combine multiple layers into one readable image:

$$\text{MIP}(u, v) = \max_z I_z(u, v).$$

Invert if text is dark-on-light.

16.5 Results: Reconstructing the Herculaneum Scrolls

[?]

The article demonstrates reconstruction of:

- **Greek:** Homer, *Hymni Homerici XV* (Hercules hymn)
- **Latin:** Lucretius, *De Rerum Natura* (Book 1, lines 72–74)

16.6 Code Repository: Archeolab

Complete pipeline (MATLAB + Python) at:

<https://gitlab.com/csc1/archeolab/>

Includes:

- `guiAxisFinder.m` — rotation axis detection
- `guiSegment.m` — manual + genetic algorithm segmentation
- `FBP_reconstruction.m` — filtered back-projection
- `MIP_flattening.m` — texture mapping + flattening
- Jupyter notebooks for Python equivalents
- Pre-acquired datasets for testing without hardware

16.7 Bridge to Chapter 14

Tomography reconstructs *hidden structure* from projections. Chapter 14 shows how *FAIR principles* reconstruct *usable knowledge* from legacy mathematical software. Both are inverse problems: from fragmented observations to coherent understanding.

16.8 Further Reading

- Foschiatti, Kittenberger, Scherzer (2025). *Deciphering Scrolls with Tomography*. Notices, 72(10).
- Seales et al. (2016). *From Damage to Discovery via Virtual Unwrapping*. Science Advances.
- Natterer (2001). *The Mathematics of Computerized Tomography*. SIAM.
- Feeman (2015). *The Mathematics of Medical Imaging*. Springer.
- Vesuvius Challenge: <https://vesuviuschallenge.org>

16.9 Exercises

16.10 Solutions

Solution.Exercise 1 (Radon transform): The Radon transform of a function $f(x, y)$ is $Rf(\theta, s) = \int_{-\infty}^{\infty} f(s \cos \theta - t \sin \theta, s \sin \theta + t \cos \theta) dt$. For a point source at (x_0, y_0) , the Radon transform is a delta function on the line θ, s such that (x_0, y_0) lies on the line.

Solution.Exercise 2 (Filtered back-projection): The filtered back-projection formula is $f(x, y) = \int_0^\pi (Rf(\theta, \cdot) * h)(x \cos \theta + y \sin \theta) d\theta$ where h is the Ram-Lak filter with frequency response $|\omega|$.

Solution.Exercise 3 (MIP): Maximum Intensity Projection assigns to each pixel the maximum intensity along the ray. For scroll unwrapping, MIP helps visualize text layers by suppressing background.

Solution.Exercise 4 (Research): The Archeolab framework combines multiple imaging modalities and segmentation algorithms. Future work includes automated text recognition and 3D visualization.

Exercise 16.1. Implement the Radon transform and its adjoint (back-projection) from scratch in NumPy. Verify $\mathcal{R}^* \mathcal{R}$ is a convolution with $1/|x|$, [?].

Exercise 16.2. Build the visible-light tomography setup (projector + camera + Arduino stage). Acquire 180 projections of a printed spiral phantom and reconstruct with FBP, [?].

Exercise 16.3. The Ram-Lak filter amplifies high-frequency noise. Implement a *windowed* ramp filter (Hann, Hamming, Shepp-Logan) and compare reconstruction quality vs. noise amplification.

Exercise 16.4 (Research). The Herculaneum scrolls have carbon-based ink (low X-ray contrast). Develop a phase-contrast CT simulation and reconstruction pipeline. How much does phase retrieval improve ink visibility? [?]

Part VI

The Future

APPENDIX A

Julia Quickstart Reference

Abstract. This appendix provides a rapid-reference guide to Julia syntax for scientists and engineers. Every chapter in this book includes Julia code; this appendix ensures you can follow along even if you are new to the language.

A.1 Basic Syntax

Variables Assign with `=`. Types are inferred:

```
x = 3.14           # Float64
n = 42             # Int64
s = "hello"        # String
b = true           # Bool
```

Comments Use `#` for single-line comments.

Semicolons `;` suppresses output in REPL.

A.2 Arrays and Matrices

```
# Vectors
v = [1, 2, 3]           # Column vector
w = [1.0, 2.0, 3.0]     # Float vector
```

```
# Matrices
A = [1 2 3; 4 5 6; 7 8 9] # 3x3 matrix
I = Matrix(Identity, 3)  # 3x3 identity
```

```
# Operations
B = A'                  # Transpose
C = A * v               # Matrix-vector product
D = A \ v               # Solve Ax = v
```

```
# Element-wise operations
x = w .^ 2              # Square each element
```

```
y = log.(v)                # Element-wise log
```

A.3 Functions

```
# Function definition
function f(x::Float64)
    return x^2 + 1
end

# One-liner
g(x) = x^2 + 1

# Anonymous function
h = x -> x^2 + 1

# Multiple dispatch
foo(x::Int) = "integer"
foo(x::String) = "string"
```

A.4 Loops and Comprehensions

```
# For loop
for i in 1:10
    println(i)
end

# While loop
i = 1
while i < 10
    i += 1
end

# Array comprehension
squares = [i^2 for i in 1:10]

# Filtered comprehension
evens = [i for i in 1:20 if iseven(i)]
```

A.5 Linear Algebra

```
using LinearAlgebra

A = rand(3, 3)
b = rand(3)
```

```
# Factorization
F = lu(A)
x = F \ b

# Eigenvalues
vals, vecs = eigvals(A), eigvecs(A)

# SVD
U, S, V = svd(A)

# Sparse matrices
using SparseArrays
S = sprand(100, 100, 0.01)
```

A.6 Plotting

```
using Plots
```

```
# Basic plot
x = linspace(0, 2pi, 100)
plot(x, sin.(x), label="sin(x)")
scatter!(x, cos.(x), label="cos(x)")
savefig("plot.png")

# 3D surface
using PyPlot
x = y = linspace(-2, 2, 50)
[X, Y] = meshgrid(x, y)
Z = X.*exp.(-X.^2 - Y.^2)
surface(X, Y, Z)
```

A.7 Differential Equations

```
using DifferentialEquations
```

```
# Lorenz system
function lorenz(u, p, t)
    sigma, rho, beta = p
    du1 = sigma * (u2 - u1)
    du2 = u1 * (rho - u3) - u2
    du3 = u1 * u2 - beta * u3
    return [du1, du2, du3]
end
```

```
u0 = [1.0, 0.0, 0.0]
p = [10.0, 28.0, 8/3]
tspan = (0.0, 100.0)
prob = ODEProblem(lorenz, u0, tspan, p)
sol = solve(prob, Tsit5())
```

A.8 Package Management

```
using Pkg
```

```
# Add packages
Pkg.add("LinearAlgebra")
Pkg.add("Plots")
Pkg.add("DifferentialEquations")
```

```
# Instantiate environment
Pkg.instantiate()
```

```
# Update all packages
Pkg.update()
```

A.9 Summary

Key Julia features for scientific computing:

1. **Multiple dispatch:** Functions specialize on argument types.
2. **Type stability:** Write code that the compiler can optimize.
3. **Arrays are 1-indexed:** Unlike Python/MATLAB.
4. **Vectorized operations:** Use dot notation `(.)` for element-wise ops.
5. **Unicode support:** `x_1`, `alpha` work in identifiers.

APPENDIX B

Linear Algebra Reference

Abstract. This appendix summarizes the linear algebra concepts used throughout this book. It covers matrix factorizations, eigenvalue theory, singular value decomposition, and numerical stability—the mathematical foundations underlying every chapter.

B.1 Matrix Factorizations

LUP Factorization $A = PLU$, where P is a permutation, L is unit lower-triangular, and U is upper-triangular. The workhorse of numerical linear algebra; used in Gaussian elimination.

QR Factorization $A = QR$, where Q is orthogonal and R is upper-triangular. Numerically stable; used for least squares problems.

Eigenvalue Decomposition $A = V\Lambda V^{-1}$ for symmetric matrices. Λ contains eigenvalues λ_i , V contains eigenvectors.

Singular Value Decomposition (SVD) $A = U\Sigma V^T$ for any matrix. Σ contains singular values $\sigma_i \geq 0$. Most important factorization in applied mathematics.

Cholesky Factorization $A = LL^T$ for symmetric positive definite matrices. Twice as fast as LUP; used in optimization and statistics.

B.2 Eigenvalue Theory

Theorem B.1 (Spectral Theorem). *If A is real symmetric, then all eigenvalues are real, eigenvectors can be chosen orthonormal, and $A = Q\Lambda Q^T$.*

Theorem B.2 (Gershgorin Circle Theorem). *Every eigenvalue of A lies in at least one Gershgorin disk $D(a_{ii}, R_i)$ where $R_i = \sum_{j \neq i} |a_{ij}|$.*

Theorem B.3 (Perron-Frobenius). *If $A \geq 0$ is irreducible, then $\rho(A)$ is a simple eigenvalue with positive eigenvectors. All other eigenvalues satisfy $|\lambda| < \rho(A)$.*

B.3 Norms and Condition Numbers

Vector norms $\|x\|_1 = \sum |x_i|$, $\|x\|_2 = \sqrt{\sum x_i^2}$, $\|x\|_\infty = \max |x_i|$.

Induced matrix norms $\|A\|_p = \sup_{x \neq 0} \|Ax\|_p / \|x\|_p$.

Spectral norm $\|A\|_2 = \sigma_{\max}(A)$.

Frobenius norm $\|A\|_F = \sqrt{\sum_{i,j} a_{ij}^2} = \sqrt{\text{tr}(A^T A)}$.

Condition number $\kappa(A) = \|A\| \cdot \|A^{-1}\| = \sigma_{\max} / \sigma_{\min}$.

B.4 Numerical Stability

Definition B.4 (Backward Stability). *An algorithm is backward stable if the computed solution $\hat{x}$ is the exact solution to a slightly perturbed problem: $A\hat{x} = b + \delta b$ with $\|\delta b\| / \|b\| \leq O(\varepsilon)$.*

Definition B.5 (Forward Error). *The forward error measures how far the computed solution is from the exact solution: $\|x - \hat{x}\| / \|x\|$.*

For a linear system $Ax = b$, the forward error satisfies:

$$\frac{\|x - \hat{x}\|}{\|x\|} \leq \kappa(A) \cdot \frac{\|\delta b\|}{\|b\|}$$

B.5 Least Squares

For overdetermined systems $Ax \approx b$, the least squares solution minimizes $\|Ax - b\|_2$:

$$x^* = \arg \min_x \|Ax - b\|_2$$

- **Normal equations:** $x^* = (A^T A)^{-1} A^T b$
- **QR method:** Solve $Rx = Q^T b$ where $A = QR$
- **SVD method:** $x^* = V \Sigma^{-1} U^T b$

The QR and SVD methods are numerically stable; normal equations should be avoided for ill-conditioned problems.

B.6 Iterative Methods

For large sparse systems, direct factorization is impractical. Iterative methods compute $x^{(k)} \rightarrow x^*$:

Jacobi method $x_i^{(k+1)} = (b_i - \sum_{j \neq i} a_{ij} x_j^{(k)}) / a_{ii}$

Gauss-Seidel Uses newest values immediately: $x_i^{(k+1)} = (b_i - \sum_{j < i} a_{ij} x_j^{(k+1)} - \sum_{j > i} a_{ij} x_j^{(k)}) / a_{ii}$

Krylov subspace methods GMRES, CG, BiCGSTAB for large sparse systems.

B.7 Summary Table

Table B.1: Matrix factorizations and their uses

Factorization	Form	Use Case
LUP	$A = PLU$	General square systems
QR	$A = QR$	Least squares, orthogonal problems
SVD	$A = U\Sigma V^T$	Rank, conditioning, pseudoinverse
Cholesky	$A = LL^T$	Symmetric positive definite
Eigenvalue	$A = V\Lambda V^{-1}$	Symmetric matrices, dynamics

Bibliography

- [1] Thomas Bacher, Mario Garrote-López, Christian Görden, and Matthias J. Neubert, *Making mathematical online resources fair*, Notices of the American Mathematical Society **73** (2026), no. 5.
- [2] John C. Baez and James D. Foley, *Operads for designing systems of systems*, Notices of the American Mathematical Society **68** (2021), no. 6.
- [3] Randall Balestrieri, Ali I. Humayun, and Richard G. Baraniuk, *On the geometry of deep learning*, Notices of the American Mathematical Society **72** (2025), no. 4.
- [4] Stephen Bancroft, *An algebraic solution of the gps equations*, IEEE Transactions on Aerospace and Electronic Systems **21** (1985), no. 1, 56–59.
- [5] Dan Bar-Lev, Omer Sabary, and Eitan Yaakobi, *Exciting coding problems for dna-based storage systems*, Notices of the American Mathematical Society **69** (2022), no. 11.
- [6] Patrick R. Bishop, Summer Chenoweth, Emmanuel Fleurantin, Alonso Ogueda-Oliva, Evelyn Sander, and Julia Seay, *3d printing of invariant manifolds in dynamical systems*, Notices of the American Mathematical Society **73** (2026), no. 1.
- [7] Mireille Boutin and Gregor Kemper, *New developments in global positioning*, Notices of the American Mathematical Society **72** (2025), no. 8.
- [8] Xavier Cabre, Ernest Fontich, and Rafael de la Llave, *The parameterization method for invariant manifolds*, Journal of Differential Equations **190** (2003), no. 1, 1–30.
- [9] Ning Chen, Steve Wiggins, and Maria Andreou, *Taming uncertainty in a complex world: The rise of uncertainty quantification*, Notices of the American Mathematical Society **72** (2025), no. 3.
- [10] Rafael S. de Souza, Emille E. O. Ishida, and Alberto Krone-Martins, *A brief history of inference in astronomy*, Notices of the American Mathematical Society **72** (2025), no. 10.
- [11] Joseph A. Emanuello and Aaron Ridley, *The mathematics of cyber defense*, Notices of the American Mathematical Society **69** (2022), no. 6.
- [12] Simone Foschiatti, Andreas Kittenberger, and Otmar Scherzer, *Deciphering scrolls with tomography*, Notices of the American Mathematical Society **72** (2025), no. 10.
- [13] Benjamin Golub, *Spectral methods in microeconomics*, Notices of the American Mathematical Society **73** (2026), no. 6.

-
- [14] Bernd Krauskopf and Hinke M. Osinga, *Growing 1d and quasi-2d unstable manifolds*, Nonlinearity **18** (2005), no. 5, 1837.
 - [15] Lek-Heng Lim and Bradley J. Nelson, *What is... an equivariant neural network?*, Notices of the American Mathematical Society **70** (2023), no. 4.
 - [16] Bastian Rieck, *Topology meets machine learning: An introduction using the euler characteristic transform*, Notices of the American Mathematical Society **72** (2025), no. 7.
 - [17] J. Maurice Rojas, *An introduction to algorithmic fewnomial theory over $\mathbb{R}$* , Notices of the American Mathematical Society **72** (2025), no. 5.
 - [18] Gilbert Strang and Kai Borre, *Linear algebra, geodesy, and gps*, Wellesley-Cambridge Press, 1997.
 - [19] John Urschel, *Numerical stability in gaussian elimination*, Notices of the American Mathematical Society **72** (2025), no. 6.